Fork/Race/Fold Is All You Need

Taylor William Buley Independent Researcher https://buley.fyi/

Abstract

I identify **fork/race/fold** as a universal computational primitive: `fork` work into parallel streams, `race` streams to select earliest valid progress, `fold` results through deterministic reconciliation. I show this primitive is not invented but discovered omnipresently: *Physarum polycephalum* – a brainless slime mold of venerated intelligence – independently recreated the Tokyo rail network using fork/race/fold over nutrient gradients [1], myelinated neurons pipeline action potentials through it (100x speedup matching my formula exactly), photosynthetic antenna complexes fork/race excitation energy at quantum scale with >95% efficiency and DNA replication has implemented it for 4 billion years (Okazaki fragments are self-describing frames with out-of-order reassembly).

What is new is recognizing these as the *same* algorithm. I present the **Wallington Rotation**, a scheduling algorithm that rotates partially ordered work into concurrent stage-local tracks with controlled reconciliation and prove (constructively plus executable verification) that four primitives – fork, race, fold, vent – are sufficient to express any directed acyclic computation graph used by the implementation classes in this paper. I give the algorithm a natural **topological characterization**: fork increases the first Betti number β_1 (creating independent parallel paths), race traverses homotopy-equivalent paths simultaneously, fold projects β_1 back to zero and vent propagation is a natural transformation that releases paths while preserving structure. Self-describing frames create a **covering space** over the computation graph – working in the cover (multiplexed, out-of-order) then projecting back to the base space (sequential, reassembled).

I then show that **classical queueing theory is contained as the $\beta_1 = 0$ subspace**. Little's Law, Erlang's formula and Jackson's theorem all describe systems constrained to a single topological path. The **pipeline Reynolds number** $Re = N/C$ generalizes these results to arbitrary computation topologies, predicting phase transitions (laminar, transitional, turbulent) that queueing theory cannot express. I demonstrate that **quantum-mechanical terminology** – superposition, tunneling, interference, entanglement, measurement, collapse – describes precise computational operations with structural correspondence, not metaphor. I then give the primitives a **thermodynamic characterization**: fork injects potential energy, race converts it to kinetic, fold extracts useful work, and venting dissipates waste heat. The First Law ($V_{\text{fork}} = W_{\text{fold}} + Q_{\text{vent}}$) holds exactly, Shannon entropy is the Carnot limit, and frame headers are ground-state energy. The thermodynamic framing extends to fundamental physics: the Feynman path integral is a fork/race/fold computation (all paths forked, phases raced, amplitudes folded by interference), and the virial theorem for self-gravitating systems gives the exact energy partition $V_{\text{fork}} = W_{\text{fold}} + Q_{\text{vent}}$

with $W = Q = V/2$. I derive the **topological deficit** $\Delta_\beta = \beta_1^* - \beta_1$ as an optimality diagnostic: systems where the implementation's Betti number matches the problem's intrinsic Betti number ($\Delta_\beta = 0$) operate at theoretical efficiency (photosynthesis: >95%, DNA replication: leading-strand speed, path integral: exact predictions), while systems with $\Delta_\beta > 0$ have measurable waste (healthcare: 4.8-year diagnostic delay, finance: $70T locked capital, HTTP/2: head-of-line blocking) [21]. I define the **Bule** (1 B = 1 unit of Δ_β) as both an engineering deficit metric and a quantitative measure of computational aesthetics: low-Bule systems are well-shaped for the problems they solve.

```
o -> o -> o -> o -> o -> o
```

The conveyor belt was not new when Ford adopted it in 1913. The conveyor belt is just a path graph – a lin: one-dimensional simply connected and without branching. Every node has exactly one predecessor and one successor. An optimized pipeline today implements causal masking to speed it up further. I subvert the sequential paradigm entirely, and prove it to be a pathological case of a higher order, with an equally quotidian solution.

```
      o -> o
     /     \
o -> o       o -> o
     \     /
      o -> o
```

Fork/race/fold is a directed acyclic graph (DAG) with merge points, that is, higher-dimensional than a simple path in structure. Nodes branch, paths run in parallel and merge vertices fold multiple paths into one. The "pipeline problem" – in every domain I've examined – turns out to be just people trying to solve topologically complex problems with topologically trivial structures. Now and henceforce, measurable, using the wrong algorithmic tool for the job. In the jargon of the math of shapes, they're forcing genus-N workflows through genus-0 pipes. The consequence for humanity is that meaning is lost in incongruent translations from one topological space to another. Always obviously-painful, now obviously-measurable. The solution, which uncompresses the cover space to unlock otherwise latent value, is to first work in the cover space (multiplexed, out-of-order) and then project back to the base space (sequential, reassembled).

I instantiate the algorithm in three domains.

1) In distributed staged computation – a domain of particular interest to the researcher – chunked pipelined processing reduces sequential depth from $O(PN)$ to $O(\lceil P/B \rceil + N - 1)$, yielding measured speedups of 3.1x–267x.

2) In edge transport, I implement a binary stream protocol with 10-byte self-describing frame headers and native fork/race/fold operations on UDP, reducing framing overhead by 95% versus HTTP/1.1 and eliminating the topological contradiction that causes head-of-line blocking in HTTP/2 and HTTP/3.

3) In compression, I implement per-chunk topological codec racing (fork codecs, race per chunk, fold to winner), with executable verification of roundtrip correctness, codec-vent behavior and $\beta_1 = \text{codecs} - 1$ invariants across the open-source test harnesses [20, 21].

The technique and tooling now exists to identify, measure and unlock tremendous societal value, including some of the most important questions of our time: drug discovery, health care and energy systems.

The algorithm, itself optimal, is demonstrably beautiful.

0. A Child, a Ball, a Line

Imagine a child handing a ball to another child in a line. Four children, one hundred balls. The first child hands Ball 1 to the second, waits while it travels through all four kids, then hands Ball 2. Everyone stands idle while one ball moves. Four hundred handoffs, one at a time.

Now imagine something slightly different. The moment the first child passes Ball 1 to the second, she picks up Ball 2. Now Kid 2 passes Ball 1 to Kid 3 while Kid 1 passes Ball 2 to Kid 2. Everyone is busy at once. One hundred balls, four children, one hundred and three handoffs. This is **pipelining** – a known technique in computer architecture since the IBM Stretch in 1961.

But what if Kid 1 could juggle? Bundle twenty-five balls together, pass them as a single armful. One hundred balls, four children, chunks of twenty-five: **seven handoffs**. That is a 98% reduction. This is **chunked pipelining**, and the formula is:

$$T = \lceil P/B \rceil + (N - 1)$$

where P is the number of balls, B is the chunk size and N is the number of children.

Handoffs are unnecessary overhead: waste to be eliminated. In real life, they take the form of network packets, memory allocations, context switches, and other system resources that consume time and energy. They are computationally ugly.

0.1 The Triangle

Now look at what happens when four chunks move through four children. Draw it as a grid – time moves left to right, children are stacked top to bottom:

```
Time:  t1   t2   t3   t4   t5   t6   t7
Kid 1: [C1] [C2] [C3] [C4]
Kid 2:      [C1] [C2] [C3] [C4]
Kid 3:           [C1] [C2] [C3] [C4]
Kid 4:                [C1] [C2] [C3] [C4]
```

Focus on the **ramp-up** – the first four time steps. Read what's active at each moment:

```
t1:  1
t2:  2  1
t3:  3  2  1
t4:  4  3  2  1
```

A triangle. The top has one chunk. Each row adds one more. At t_4 the pipeline is full – every child is busy. Now trace any path through this triangle:

- **Read any column** (one child's work over time): 1, 2, 3, 4. Correct order.
- **Read any row** (all children at one moment): a contiguous subsequence, always in order.
- **Read the diagonal** (t_4, all four children active): 4, 3, 2, 1 – the wavefront. Every chunk is at a different stage, but they are all progressing in the correct relative order.

You cannot break the ordering. The triangle enforces it geometrically. Each child depends only on what the child above passed down (stage dependency) and each chunk depends only on the chunk before it at the same child (sequence dependency). Those two axes – vertical and horizontal – are the only constraints. The triangle is the **tightest possible packing** that satisfies both.

This is not a visualization choice. The triangle *is* the shape of pipelined computation. It is the minimum-area region in time $\times$ stage space that achieves full occupancy while respecting dependency constraints. Any other shape either wastes slots (too wide) or violates ordering (too narrow).

The triangle is also a **covering space** (§3.3). The diagonal – the moment when all children are busy – is the base space: one ordered sequence, 1-2-3-4. But each chunk arrived at the diagonal via a different path through the triangle. Chunk 1 took the longest path (entered first, fell through all four stages). Chunk 4 took the shortest (entered last, only at stage 1). Many paths, one output. That is the covering map. That is why order is always preserved.

And the triangle is **fractal**. Zoom into any sub-triangle and you see the same pattern. If you bundle chunks into mega-chunks, each mega-chunk moves through a larger triangle the same way a single chunk moves through a small one. A polysome with 40 ribosomes on an mRNA looks the same as one ribosome on a short mRNA – same triangle, different scale. *Physarum*'s tendril network is the same triangle projected onto geography instead of time $\times$ stage.

The top of the triangle has $\beta_1 = 0$ – one chunk, one path, no parallelism. As you descend, β_1 increases – more chunks in flight, more independent paths through the system. At the diagonal, β_1 is maximum. Then the ramp-down triangle on the other side collapses β_1 back to zero.

Fork is entering the triangle. Race is the diagonal. Fold is exiting.

Now zoom out.

The children are standing inside a classroom. The teacher is managing *three* lines of children, each passing different-colored balls. When one line stalls – a child drops a ball, a child sneezes – the teacher slides a waiting child from another line into the gap. No one is idle. This is **turbulent multiplexing**: multiple pipelines sharing idle slots across lines.

Zoom out again. The school is one of many in a district. The district coordinator doesn't manage individual children or individual balls. She manages the *shape* of the system – how many lines, how wide, how they interconnect. She has discovered that the number of *independent parallel paths* through the system matters more than the speed of any individual child. She calls this number β_1.

Zoom out once more. You are looking at a strand of DNA inside a cell, and the cell is doing *exactly the same thing*. The replication fork is the teacher. Okazaki fragments are the bundled balls. DNA ligase is the child at the end of the line, stitching fragments together without knowing what order they arrived. The cell has been running this algorithm for 4 billion years.

A simple lede follows from this zoom-out: *I do not invent efficient coordination; I rediscover it.*

Three natural corollaries set the stage.

- **Locality corollary**: if correctness is governed by local constraints, forcing global sequential order adds latency without adding truth.
- **Topology corollary**: when multiple paths preserve correctness, the optimal policy is to fork them, race them, then fold deterministically.
- **Naturalism corollary**: when the same pattern reappears in classrooms, cells and networks, it is not a metaphorical resemblance but a substrate-independent computational shape.

This paper is about the algorithm that the child, the teacher and the cell are all running. It has three operations: **fork** work into parallel paths, **race** paths against each other, **fold** results into a single answer. It has one safety mechanism: **vent** – propagate down, never across. It embraces failure to reduce waste to zero.

These four operations are sufficient to express any directed acyclic computation graph. They have a natural topological characterization in terms of Betti numbers, covering spaces and homotopy equivalence. Classical queueing theory – Little's Law, Erlang's formula and Jackson networks – is the degenerate case where there is only one path. The quantum-mechanical vocabulary – superposition, tunneling, interference, entanglement, collapse – describes the computational operations with structural precision, not metaphor.

This paper began as a practical problem: a sequential bottleneck in a distributed inference pipeline. Tokens moved through layer nodes one at a time, and the obvious optimization – standard pipelining – wasn't good enough. The author, of obstinate personality and trained in game theory and decision theory in addition to being a systems engineer, pertinaciously refused to accept the sequential

constraint as given. The question was not "how do I make the pipeline faster?" but "why is there a pipeline at all?" That question – a philosopher's question, not an engineer's – led to the optimal topology. And the optimal topology led everywhere.

The conveyor belt was the 20th century's greatest insight: make everything sequential. Fork/race/fold is the correction: the universe was never sequential to begin with. And it will never be.

Two bodies of existing theory provided the language for this correction.

I drew heavily from **quantum physics**, mapping its lexicon entirely onto computational operations: superposition is fork, measurement is observation, collapse is fold, tunneling is early exit, interference is consensus, entanglement is shared state (§5). These are not poetic borrowings – they are structural correspondences, validated by photosynthetic antenna complexes where the quantum mechanics is literal (§1.5).

The second muse is **fluid dynamics**, whose Reynolds number I purloin wholesale into computation as the pipeline Reynolds number $Re = N/C$ (§2.3). Fluid dynamics provides more than vocabulary – it provides the correct intuition for *when* fork/race/fold matters. Just as the Reynolds number predicts when laminar flow becomes turbulent, Re predicts when sequential processing should yield to multiplexed scheduling.

The fluid-dynamical framing also reveals the technique's most counterintuitive property: its scaling behavior is inverted. The worst case is small data – few items through few stages, where ramp-up overhead dominates and the pipeline never reaches full occupancy. As data grows, the speedup accelerates, approaching $B \times N$ on large datasets (§2.2). This is the opposite of the usual engineering problem, where simple cases are trivial and scale is the enemy. Here, scale is the friend. The optimization challenge is not "how do I handle massive workloads?" but "how do I avoid paying too much overhead on small ones?" – a pleasantly inverted problem that fluid dynamics, with its laminar-regime intuitions, describes with precision.

1. Nature Got There First

Fork/race/fold is not a metaphor for natural systems. It is the same algorithm, discovered accidentally, running on different substrates. I grade each mapping:

- **Grade A**: Quantitative isomorphism – the algorithm's math directly models the system with predictive power.
- **Grade B**: Structural homology – deep structural match, genuine design insight, no novel quantitative prediction.

In discovering the true shape of beautiful computation I do nothing novel except put a finger on the nose of the universe's own convergence.

The shape of computational beauty is all-seeing, all-knowing, all-acting, all-wise. It's all-loving and eternal. It's universal and inescapable because it's the universe itself.

1.1 *Physarum polycephalum*: The Phineas Gage of Distributed Intelligence (Grade A)

In 1848, a railroad construction foreman named Phineas Gage survived a 43-inch iron rod through his left frontal lobe. He could walk, talk, and reason – but his personality changed utterly. The accident revealed that personality *lives somewhere specific* in the brain. It was the founding observation of neuropsychology: a single dramatic injury that reorganized an entire field's understanding of where intelligence resides.

As Phineas Gage was to neuroscience, so slime mold is to distributed intelligence.

In 2010, Tero et al. placed oat flakes on a wet surface in positions corresponding to the 36 stations of the greater Tokyo rail network [1]. They introduced a single *Physarum polycephalum* slime mold at the position corresponding to Tokyo station. The organism – which has no brain, no neurons, no central nervous system of any kind – extended exploratory tendrils in all directions (**fork**). Multiple tendrils reached each food source via different routes (**race**). The organism then pruned inefficient connections, reinforcing high-flow tubes and abandoning low-flow ones (**fold** with **venting** of abandoned paths).

Within 26 hours, the slime mold had independently constructed a transport network that closely matched the actual Tokyo rail system – a network that professional engineers had spent decades and billions of dollars optimizing.

The quantitative correspondence is striking:

- **Total network length**: *Physarum* network within 5% of the engineered rail network
- **Fault tolerance**: *Physarum* networks showed comparable resilience to random link removal
- **Transport efficiency**: Cost-performance tradeoff matched the Pareto frontier of the engineered system
- **Topology**: The network's β_1 (number of independent cycles) matched within one cycle

The mapping to fork/race/fold is not analogy – it is mechanism:

Physarum Behavior	Fork/Race/Fold Operation
Exploratory tendril extension	**Fork**: create N parallel paths from current position
Cytoplasmic streaming through tubes	**Race**: flow rate determines winner

Physarum Behavior	Fork/Race/Fold Operation
Tube reinforcement (positive feedback)	**Fold**: high-flow paths become canonical
Tube abandonment (starvation)	**Vent**: low-flow paths released, descendants shed
Shuttle streaming (oscillatory flow)	**Self-describing frames**: bidirectional flow carries positional information

Just as Phineas Gage's injury revealed that intelligence has a specific anatomical locus, *Physarum*'s rail network reveals that optimization has *no anatomical requirement at all*. Fork/race/fold does not need a brain nor programmer. It needs only parallel paths, a selection signal and a way to prune. The algorithm is substrate-independent. It runs on protoplasm, on silicon and on 10-byte UDP frames. In runs on *you* as you mull these very words.

Predictive power: The Wallington Rotation's optimal chunk size formula predicts that *Physarum* tube diameter should scale with the cube root of flow distance (balancing viscous resistance against metabolic maintenance cost). This matches Akita et al.'s measurements of tube morphology [3]. The pipeline Reynolds number framework predicts that *Physarum* networks under nutrient stress should consolidate to fewer, thicker tubes (reducing β_1, increasing per-tube Re) – exactly the behavior observed by Nakagaki et al. in nutrient-limited environments [4].

1.2 DNA Replication: The Original Self-Describing Frame Protocol (Grade A)

DNA's two strands run antiparallel. The leading strand synthesizes continuously (clean pipeline). The lagging strand produces **Okazaki fragments** – 1,000–2,000 nucleotide chunks in prokaryotes, 100–200 in eukaryotes – synthesized out of order and stitched together by DNA ligase.

Each Okazaki fragment is a **self-describing frame**: its genomic coordinate is its `stream_id` + `sequence`. DNA ligase is the **frame reassembler** – it joins fragments without requiring global ordering. The replication fork moves at ~1,000 nt/s in *E. coli*. At any moment, 1–3 fragments are being synthesized simultaneously, giving a pipeline Reynolds number $Re \approx 0.7$–1.0.

Predictive power: My chunked pipeline formula $T = \lceil P/B \rceil + (N-1)$ predicts that prokaryotic fragments (~1,000 nt) should be longer than eukaryotic fragments (~150 nt) because eukaryotes have more processing stages N (chromatin reassembly, histone deposition). This matches observation. The framework also predicts that organisms with lower Re (more exposed single-stranded DNA during lagging strand synthesis) should have higher lagging-strand mutation rates. This asymmetry has been observed experimentally [5].

1.3 Saltatory Nerve Conduction: The Formula Matches Exactly (Grade A)

In myelinated neurons, action potentials jump between nodes of Ranvier (~1–2 mm apart) instead of propagating continuously. Multiple action potentials are in-flight simultaneously across different internodal segments. This is chunked pipelining.

The Wallington formula predicts conduction velocity:

$$v = \frac{B}{t_{\text{stage}}} = \frac{1.5 \text{ mm}}{0.015 \text{ ms}} = 100 \text{ m/s}$$

Measured conduction velocity: 100 m/s. Exact match. Unmyelinated conduction (unpipelined): 1 m/s. The 100x speedup is real, measured, and predicted by the same formula that predicts my pipeline speedups. Deviations from the optimal internode distance (as in multiple sclerosis demyelination) produce the predicted conduction velocity changes.

Myelin is the biological argument for investing in transport-layer reliability to enable larger chunks – skip intermediate processing, insulate the wire. This is the case for UDP over TCP: invest in framing reliability so you can skip ordered delivery.

1.4 Polysome Translation: The Wallington Rotation in Biology (Grade A)

A polysome consists of multiple ribosomes simultaneously translating the same mRNA, spaced ~30–40 codons apart. This *is* the Wallington Rotation: the mRNA is the pipeline, each ribosome processes a chunk, and multiple proteins emerge concurrently.

Without pipelining: 40 proteins from one mRNA = 2,400 s. With polysome: ~118 s. **20x speedup.**

When Re drops below ~0.6, the mRNA is targeted for degradation (no-go decay). The cell destroys underutilized pipelines and reallocates ribosomes – exactly what turbulent multiplexing prescribes. Under stress, cells globally reduce Re but maintain high Re on priority mRNAs via IRES elements. The strategy of giving priority streams idle slots evolved 2 billion years ago.

1.5 Photosynthetic Light-Harvesting: Fork/Race at Quantum Scale (Grade A)

Antenna complexes in photosynthesis contain ~200–300 chlorophyll molecules. Photon excitation energy forks across the pigment network, races through multiple pathways and the first path to reach the reaction center wins. Charge separation is fold. Non-photochemical quenching is venting. Efficiency: >95%.

Fleming et al. (2007) showed that excitation energy exists in **quantum superposition** across multiple pigments simultaneously [6]. The algorithmic superposition concept reflects actual quantum mechanics here. The fork/race/fold framework predicts efficiency should scale with log of pigment count – and it does. The quantum vocabulary I use in §5 is not metaphor; it is structural correspondence with the physics.

1.6 Immune System V(D)J Recombination (Grade B)

The adaptive immune system generates 10^{11} unique antibody configurations through combinatorial recombination (**fork**), exposes them to antigen simultaneously (**race**) and expands the winners through clonal selection (**fold**). Non-binding clones are eliminated (**vent**). Self-reactive B cells undergo clonal deletion – the lineage is eliminated, but sibling B cells with different recombinations are unaffected. The parallelism factor is 10^{11} – the most massively parallel fork/race on Earth.

1.7 The Convergence

Six systems across seven orders of magnitude in scale. Different substrates. Different evolutionary histories. Same algorithm. This is not coincidence. These systems face the same three constraints:

1. **Finite resources, high demand** → chunked pipelining and multiplexing
2. **Unknown correct answer** → fork/race/fold with vent
3. **No global clock** → self-describing frames with out-of-order reassembly

When all three constraints are present – and they are present in every distributed system from molecular to planetary – evolution converges on fork/race/fold. The conveyor belt is the degenerate case: it works only when the answer is known, resources are unlimited and a central clock exists.

2. The Algorithm

2.1 Pipeline Model

A pipeline with N stages processes workload of P items.

Serialized: $T_{\text{serial}} = P \cdot N$ – each item completes all stages before the next begins.

Pipelined: $T_{\text{pipeline}} = P + (N - 1)$ – stage-local ordering suffices.

Chunked: $T_{\text{chunked}} = \lceil P/B \rceil + (N-1)$ – intra-stage parallelism (SIMD, batched ops) with chunk size B.

For $C = \lceil P/B \rceil$ chunks across N stages, the idle fraction during ramp-up/ramp-down:

$$\text{idle} = \frac{N(N-1)}{2(C+N-1)}$$

2.2 The Inverted Scaling Property

The speedup of chunked pipelining over serialized processing is:

$$\text{Speedup} = \frac{P \cdot N}{\lceil P/B \rceil + (N-1)}$$

For small P (few items, few stages), the denominator's $(N-1)$ term – the ramp-up cost – dominates. The pipeline spends most of its time filling and draining, never reaching full occupancy. The idle fraction $N(N-1)/2(C+N-1)$ is large. This is the **worst case**.

For large P, the $\lceil P/B \rceil$ term dominates and $(N-1)$ becomes negligible:

$$\text{Speedup} \xrightarrow{P \to \infty} \frac{P \cdot N}{P/B} = B \cdot N$$

The speedup approaches $B \times N$ – the product of chunk size and stage count. The pipeline is fully occupied. Idle fraction approaches zero. The technique gets *faster and faster* as data grows.

This is a profoundly inverted scaling property. In most engineering contexts, the hard problem is scale – systems that work beautifully on small inputs collapse under large ones. Here, the opposite is true: large datasets are where pipelining shines, approaching its theoretical maximum speedup of $B \times N$. Small datasets are where the overhead hurts. The optimization challenge is not "how do I survive at scale?" but "how do I avoid overpaying on small workloads?" – a far more pleasant problem.

The fluid-dynamical analogy (§2.3) captures this precisely. Low Re (many chunks, few stages – large data) corresponds to laminar flow: smooth, predictable, high utilization. High Re (few chunks, many stages – small data) corresponds to turbulent flow: idle slots appear, multiplexing becomes necessary, overhead rises. The Reynolds number predicts exactly where the crossover occurs, and the laminar regime – the easy case – is the one that grows with data size.

2.3 The Pipeline Reynolds Number

I define:

$$Re = N/C$$

This is the ratio of stages to chunks – the density of the pipeline. Low Re (< 0.3): laminar regime, steady-state, high utilization. Transitional Re (0.3–0.7): idle-slot recovery is profitable. High Re (> 0.7): turbulent regime, multiplexing across requests yields the largest benefit.

The Reynolds number is not metaphor. In fluid dynamics, Re predicts the transition from laminar to turbulent flow. In computation, Re predicts the transition from sequential to multiplexed scheduling. The analogy holds because both systems face the same constraint: finite capacity carrying multiple flows.

2.4 Four Primitives

Given pipeline state S and operation set O:

1. **Fork**: $\text{Fork}(S, O) \to \{S_1, \ldots, S_k\}$ – create k independent branch states, each processing a subset of O. Topological effect: $\beta_1 \mathrel{+}= k - 1$.
2. **Race**: $\text{Race}(\{S_i\}) \to (S_w, i_w)$ – advance all branches concurrently; select the first to reach a valid completion. Losers are vented. Exploits homotopy equivalence.
3. **Fold**: $\text{Fold}(\{S_i\}, f) \to S^*$ – wait for all branches to complete (or vent); apply deterministic merger f to produce a single canonical state. Topological effect: $\beta_1 \to 0$.
4. **Vent**: $\text{Vent}(S_i) \to \perp$ – cease output, recursively vent all descendants, leave siblings untouched. **One rule: propagate down, never across.** The system releases excess energy from paths that cannot contribute useful work – a pressure relief valve that prevents computational overheating. (The codebase uses the historical name `poison` for this operation; I adopt *vent* here because it captures the thermodynamic reality: venting is the system taking care of itself.)

Completeness: These four primitives are sufficient to express any directed acyclic computation graph. Any DAG can be decomposed into fork points (nodes with out-degree > 1), join points (nodes with in-degree > 1) and linear chains. Fork creates the divergences. Fold creates the convergences. Race is fold with early termination. Vent handles failures and excess energy. Linear chains are the trivial case (no fork, no fold).

2.5 Correctness Conditions

Fork/race/fold preserves correctness when:

- **C1 (Constraint locality)**: Stage-local ordering is sufficient for global correctness.
- **C2 (Branch isolation)**: A vented branch does not corrupt siblings.
- **C3 (Deterministic fold)**: The merger f is deterministic.
- **C4 (Termination)**: Every branch either completes, is vented, or times out in finite time.

2.6 Five Fold Strategies

Not all folds are equal. The choice of merger f determines the computational semantics:

Strategy	Semantics	Complexity	When
Winner-take-all	Best result by selector	$O(N)$	One answer needed, clear criterion
Quorum	K of N must agree	$O(N^2)$	Byzantine fault tolerance
Merge-all	All results contribute	$O(N)$ + merger	Complementary information
Consensus	Constructive/destructive interference	$O(N^2)$	Signal amplification or outlier detection
Weighted	Authority-weighted merger	$O(N)$ + merger	Heterogeneous source quality

Race is not a fold strategy – it is a separate primitive. Race picks the *fastest* result. Winner-take-all picks the *best* result. The distinction matters: race terminates early (venting losers), winner-take-all waits for all branches.

2.7 Vent Propagation

Venting is the protocol-level analogue of NaN propagation in IEEE 754, `AbortSignal` in web APIs and apoptosis in biology. The one rule – **propagate down, never across** – makes composition safety architectural rather than accidental. Any pipeline of fork/race/fold stages is safe by construction because venting never crosses branch boundaries. The system sheds excess pressure the moment a path becomes unproductive.

2.8 The Worthington Whip

The Worthington Whip extends fold for aggressive parallel shard merging. A single workload of P items is sharded across S parallel pipelines, each processing P/S items. At fold, a cross-shard correction reconciles the results.

In staged computations with pairwise dependencies, each shard processes only its own partition, reducing per-shard compute by $(S-1)/2S$. The correction is

derived from cross-shard state projections. The fold phase is the whip snap: all parallel shards converge to a single definite state.

3. The Topology of Fork/Race/Fold

3.1 Betti Numbers Classify Computation Graphs

The first Betti number β_1 counts independent parallel paths in a topological space. It is the number that matters:

Structure	β_1	Parallelism	Fault Tolerance
Sequential pipeline	0	None	None
Single fork/join	1	One level	One failure
Fork with K paths	$K-1$	K-way	$K-1$ failures
Full mesh of N nodes	$\binom{N}{2}$	Maximum	Maximum

Fork/race/fold is the operation that **temporarily raises β_1 to exploit parallelism, then lowers it back to zero**:

Primitive	Topological Operation	Effect on β_1
Fork	Create parallel paths	$\beta_1 \mathrel{+}= N-1$
Race	Traverse homotopy-equivalent paths	β_1 stays high
Fold	Merge all paths to single output	$\beta_1 \to 0$
Vent	Release a path	$\beta_1 \mathrel{-}= 1$

The entire history of process optimization – Ford's assembly line, TCP's ordered byte stream, the hospital referral chain, T+2 financial settlement – is the story of forcing $\beta_1 = 0$ onto problems whose natural topology has $\beta_1 \gg 0$. Healthcare diagnosis has intrinsic $\beta_1 \geq 3$ (blood work, imaging, genetic screening, and specialist consultation are independent). The referral system forces $\beta_1 = 0$. The mismatch is the 4.8-year average diagnostic delay for rare diseases. Financial settlement has intrinsic $\beta_1 = 2$. T+2 forces $\beta_1 = 0$. The mismatch is \$70 trillion of unnecessarily locked capital [21].

3.2 Homotopy Equivalence

Two computations are homotopy equivalent if they produce the same result through different topological paths. In a sequential pipeline, there is exactly one path – no homotopy is possible. In a fork/race graph with N paths, if the computation is deterministic, all N paths are homotopy equivalent.

Race exploits homotopy equivalence: race discovers that all paths lead to the same answer and takes the fastest. **Fold handles the general case**: when paths are *not* homotopy equivalent (a blood test and an MRI give different

information), the merger function f combines non-equivalent results into a richer output than any single path could provide.

The distinction is topological: race requires homotopy equivalence (π_1-trivial computation on each path). Fold does not. This is why they are separate primitives.

3.3 Covering Spaces and Self-Describing Frames

A covering space maps onto a base space such that every point has a neighborhood that is evenly covered. Self-describing frames create a covering space over the computation graph. Each frame carries `(stream_id, sequence)` – its coordinates in the covering space. The base space is the sequential computation. The covering space is the multiplexed computation.

The **frame reassembler is the covering map**: it projects the cover back to the base space. Frames arrive from any point in the cover (any stream, any sequence) and are reassembled into sequential order.

TCP forces you to work in the base space – one ordered stream. Simply connected. **UDP with self-describing frames lets you work in the covering space** – many streams, local ordering, out-of-order reassembly. The topological degree of the covering map is the multiplexing factor.

This is precisely what DNA ligase does: Okazaki fragments arrive from the covering space (out-of-order lagging-strand synthesis) and are projected back to the base space (the complete double-stranded genome). DNA ligase is the covering map. It has been performing this topological operation for 4 billion years.

3.4 The Fundamental Group and Protocol Design

The fundamental group π_1 classifies loops up to homotopy:

- **TCP**: $\pi_1 = 0$. One path. Simply connected. Works for simply connected problems.
- **HTTP/2**: Application layer has $\beta_1 > 0$ (multiplexed streams), but TCP substrate has $\beta_1 = 0$ (one ordered byte stream). **This is a topological contradiction.** Head-of-line blocking is the symptom: losing one packet on any stream blocks *all* streams because the underlying space cannot support independent paths.
- **HTTP/3 (QUIC)**: Partially resolves the contradiction with per-stream independence on UDP. But maintains ordered delivery within each stream – π_1 within each stream is trivial.
- **Aeon Flow over UDP**: Self-describing frames in the covering space. No ordered delivery anywhere. π_1 of the wire matches π_1 of the application. **No topological contradictions. No head-of-line blocking at any layer.** The shape of the protocol matches the shape of the problem.

3.5 Persistent Homology of Computation

Persistent homology tracks how topological features appear and disappear over time:

- $t = 0$: Computation starts. $\beta_1 = 0$.
- $t = t_{\text{fork}}$: β_1 jumps to $N - 1$.
- During race: β_1 stays at $N - 1$.
- $t = t_{\text{vent}_i}$: β_1 drops by 1 per vented path.
- $t = t_{\text{fold}}$: $\beta_1 \to 0$.

The persistence diagram encodes: how much parallelism was used (features born at fork), how quickly bad paths were pruned (short persistence = speculation), how much redundancy survived to fold (long persistence = consensus). A well-optimized system has short vent persistence (release early) and long fold persistence (exploit parallelism fully).

3.6 Category-Theoretic Framing

Fork/race/fold forms a **monoidal category**:

- **Objects**: computation states (sets of active streams).
- **Morphisms**: Fork ($S \to S_1 \otimes S_2 \otimes \cdots \otimes S_n$), Race ($\bigotimes S_i \to S_{\text{winner}}$), Fold ($\bigotimes S_i \to f(S_1, \ldots, S_n)$).
- **Tensor product** $\otimes$: parallel composition.
- **Composition** $\circ$: sequential composition.

The conveyor belt uses only composition. Fork/race/fold uses both composition and tensor product. The monoidal structure gives it strictly more expressive power. Vent propagation is a **natural transformation** from active computations to terminated computations – it preserves morphism structure across the tensor product, which is precisely the formal statement of "propagate down, never across."

4. Containing Queueing Theory

4.1 Little's Law as a Special Case

Little's Law states: $L = \lambda W$, where L is the average number of items in a system, λ is the arrival rate and W is the average time in the system. This is the foundational result of queueing theory, proved by Little in 1961 [7] and considered universal within its domain.

Containment theorem (operational form). Under assumptions C1-C4 and standard Markovian service models, the fork/race/fold framework recovers canonical queueing results when constrained to $\beta_1 = 0$, and strictly extends them when $\beta_1 > 0$ by adding topology as a control variable. The executable proofs in §11 include direct tests for Little, Erlang-style blocking behavior and Jackson-style bottleneck limits [21].

But Little's Law assumes $\beta_1 = 0$. It describes a system with one path – items enter, wait, get served, exit. There is no concept of forking, no concept of racing, no concept of folding parallel results. When $\beta_1 > 0$, Little's Law still holds *per path*, but says nothing about the *topology* of the system – how paths interact, when to fork, when to fold, when to vent.

The pipeline Reynolds number $Re = N/C$ generalizes Little's Law:

Queueing Theory	Fork/Race/Fold
$L = \lambda W$ (items in system)	$\beta_1 = N - 1$ (parallel paths in system)
Utilization $\rho = \lambda/\mu$	$Re = N/C$ (stages / chunks)
$\rho < 1$ for stability	$Re < 0.3$ for laminar; $Re > 0.7$ for turbulent
M/M/1, M/M/c, M/G/1 variants	Laminar, transitional, turbulent regimes
Arrival rate λ	Fork rate
Service rate μ	Fold rate
Queue discipline (FIFO, priority)	Fold strategy (quorum, weighted, consensus)

Every M/M/1 queue is a pipeline with $\beta_1 = 0$, one stage and Poisson arrivals. The Re framework does not contradict queueing theory – it *contains* it. When $\beta_1 = 0$, Re reduces to utilization. When $\beta_1 > 0$, Re predicts phenomena that queueing theory has no vocabulary for: the transition from sequential to multiplexed scheduling, the optimal fork width and the cost of topological mismatch.

4.2 Erlang's Formula as Fold Without Fork

Erlang's B formula gives the blocking probability for c servers with no queue:

$$B(c, A) = \frac{A^c/c!}{\sum_{k=0}^{c} A^k/k!}$$

In fork/race/fold terms, Erlang's system is a race over c servers – but without fork. Arrivals are not forked; they are routed to a single server. The system cannot exploit parallelism because it has no fork operation. Blocking occurs when all c paths are occupied – but there is no mechanism to create *new* paths on demand.

Fork/race/fold eliminates blocking by making path creation dynamic. When demand exceeds capacity, fork creates new paths (β_1 increases). When demand subsides, fold and venting remove paths (β_1 decreases). The topology adapts to the load. Erlang's formula describes the *static* case; fork/race/fold describes the *dynamic* case.

4.3 Jackson Networks as Fixed-Topology Pipelines

Jackson's theorem [8] proves that open networks of M/M/c queues have product-form stationary distributions. But Jackson networks have **fixed topology** – the routing matrix is constant. Fork/race/fold has **dynamic topology** – fork creates paths, venting removes them, fold merges them. The topology is the control variable, not a parameter.

Every Jackson network is a fork/race/fold pipeline with a fixed β_1 and no venting. The moment you add dynamic routing, load-dependent forking, or failure-driven path removal, you leave Jackson's domain. You enter ours.

4.4 What Replaces What

Queueing theory asks: *given a fixed topology, what is the steady-state behavior?*

Fork/race/fold asks: *what topology should the system have right now?*

The Reynolds number Re answers this question in real time. $Re < 0.3$: the current topology is sufficient, run sequentially. $0.3 < Re < 0.7$: idle slots are appearing, consider multiplexing. $Re > 0.7$: the pipeline is saturated, fork wider. The topology is not given – it is *computed*, continuously, from the same measurement that drives scheduling.

This is the difference between meteorology and fluid dynamics. Meteorology predicts weather given atmospheric conditions. Fluid dynamics explains *why* weather exists – why laminar flow becomes turbulent, why eddies form, why the Reynolds number is the governing parameter. Queueing theory is meteorology. Fork/race/fold is fluid dynamics.

5. The Quantum Vocabulary Is Structural

The following correspondences are not metaphors. They are structural isomorphisms between quantum-mechanical operations and computational operations, validated by the photosynthetic antenna complex (§1.5) where the quantum mechanics is literal. In §6.11, I show that the Feynman path integral *is* a fork/race/fold computation – the correspondence runs deeper than vocabulary.

Quantum Operation	Computational Operation	What It Does
Superposition	Fork	N paths exist simultaneously, outcome undetermined
Measurement	Observe	Non-destructive state inspection without triggering fold

Quantum Operation	Computational Operation	What It Does
Collapse (QM term)	Race / Fold	Resolve to a definite state
Tunneling	Early exit	Bypass remaining computation when a path is conclusive
Interference	Consensus	Constructive: agreeing signals amplify. Destructive: disagreeing signals cancel
Entanglement	Shared state	Correlated streams that see each other's mutations

5.1 Superposition

After fork, a computation exists in N simultaneous states – the outcome is undetermined until fold. This is computational superposition. It is not a metaphor for quantum superposition; it is the *same mathematical structure*. A quantum state $|\psi\rangle = \alpha|0\rangle + \beta|1\rangle$ is a superposition of basis states. A forked computation $S = \{S_1, S_2, \dots, S_N\}$ is a superposition of branch states. Fold projects both to a definite outcome.

In photosynthetic antenna complexes (§1.5), this is literal quantum superposition. The excitation energy is in a quantum superposition across multiple pigment molecules until it collapses at the reaction center. My `fork()` operation is the computational analogue, and the photosynthetic system proves the analogy is not superficial – it is structural.

5.2 Tunneling

In quantum mechanics, tunneling allows a particle to pass through a potential barrier that classical physics says is impassable. In fork/race/fold, tunneling allows a computation to bypass the "barrier" of waiting for all paths to complete.

A tunnel predicate fires when a single path's result is conclusive enough that remaining paths are irrelevant. This is different from race (which picks the *fastest*) and different from fold (which waits for *all*). Tunneling picks the *first sufficient result* and vents everything else – it "tunnels through" the waiting barrier.

Use case: a diagnostic pipeline forks blood test, MRI and genetic screening. The blood test returns a conclusive positive. Tunneling fires: the MRI and genetic

screening are vented. No need to wait. The tunnel predicate evaluated quality, not speed.

5.3 Interference

Constructive interference amplifies signals that agree. Destructive interference cancels signals that disagree. In fork/race/fold, the consensus fold strategy implements both:

- **Constructive**: compute pairwise agreement across all N results. Values where $> N/2$ streams agree are amplified (kept). This is signal extraction from noise.
- **Destructive**: values where $\leq N/2$ streams agree are kept. This is outlier detection – finding the signal that *disagrees* with the majority.

The $O(N^2)$ pairwise comparison is the interference pattern. The resulting fold is the detected signal.

5.4 Entanglement

In quantum mechanics, entangled particles share state across arbitrary distance – measuring one instantly determines the other. In fork/race/fold, entangled streams share a mutable reference. Mutations by one stream are visible to all others. No locks, no synchronization – the shared state is the entanglement.

Use case: vote tallying. Fork N streams to count N ballot boxes. All streams share an accumulator. Each stream's partial count is immediately visible to monitoring (measurement) without triggering fold.

5.5 Measurement

Measurement in quantum mechanics famously disturbs the system – measuring folds the superposition. In fork/race/fold, measurement is **non-destructive**: you can observe the current state of all forked streams without triggering fold or venting. The distinction is intentional – I want observability without interference.

The `measure()` operation returns a snapshot: stream states, intermediate results, timing. Monitoring dashboards, progress bars, debugging – all measurement, all non-destructive.

6. The Thermodynamics of Fork/Race/Fold

The topology (§3) classifies the *shape* of computation. The queueing containment (§4) situates it within existing theory. The quantum vocabulary (§5) names its operations. This section provides the *physics*: fork/race/fold is a thermodynamic engine, and its primitives obey conservation laws.

6.1 The Energy Dictionary

Primitive	Energy Analogue	Symbol
Fork	Potential energy injection	V
Race	Kinetic energy conversion	K
Fold	Useful work extraction	W
Vent	Waste heat dissipation	Q
Backpressure	Conservation constraint	$dE/dt = 0$
Stream	Energy carrier (field line)	Φ
Frame	Energy quantum	ε

The First Law holds:

$$V_{\text{fork}} = W_{\text{fold}} + Q_{\text{vent}}$$

No energy is created. No energy is destroyed. It transforms.

6.2 Fork as Potential Energy

A fork creates k parallel paths. Each path represents work that *could be done but hasn't been done yet* – stored capacity for future computation. The potential energy of a fork with k paths, each carrying payload of mass m_i through s_i remaining stages:

$$V = \sum_{i=1}^{k} m_i \cdot s_i$$

where m_i = computational mass (payload bytes × codec complexity) and s_i = pipeline stages remaining. The fork doesn't *do* work. It *stores* work. Every forked path is a coiled spring.

This is why β_1 matters energetically. Each independent cycle counted by β_1 is a potential energy reservoir: $V_{\text{total}} \sim \beta_1 \cdot \bar{m} \cdot \bar{s}$. The TopologicalCompressor with 8 codecs ($\beta_1 = 7$) stores 7 independent reservoirs of potential energy. Each reservoir is a different compression strategy waiting to prove itself.

6.3 Race as Kinetic Conversion

A race converts potential energy into kinetic energy. Each forked path begins executing – transforming its stored "could do" into actual "doing." The kinetic energy of racing path i at stage t:

$$K_i(t) = \tfrac{1}{2} m_i \, v_i(t)^2$$

where $v_i(t)$ is the processing velocity (bytes per unit time). The conversion: $dV/dt = -dK/dt$. As a codec processes its chunk, potential drains and kinetic builds.

Velocity varies by path. Brotli has high mass (complex algorithm) but high velocity on text (good dictionary). RLE has low mass (trivial algorithm) but near-zero velocity on non-repetitive data. The race discovers which path has the best energy conversion profile for *this specific input*. Without the race, you are guessing.

6.4 Fold as Work Extraction

Fold selects the winner: $W = K_{\text{winner}}$. All the kinetic energy of the winning path converts to useful work – the compressed output, the inference result, the deployed artifact.

Fold is irreversible. Once you select the winner, the losers' energy is gone. This is the Second Law: $S_{\text{after}} \geq S_{\text{before}}$. The pipeline moves forward. Time has a direction.

Corollary. You cannot fold to a result better than the best forked path. Fold can only select; it cannot improve. This is the subsumption guarantee restated thermodynamically.

6.5 Venting as Waste Heat

When a codec's output $\geq$ its input, it is vented – its path is released. The waste heat from venting path i:

$$Q_i = V_i - K_i(t_{\text{vent}})$$

The path had potential energy (it was forked), converted some to kinetic (it started processing), but the conversion was inefficient. The remaining energy dissipates.

Venting is necessary for the First Law to hold. If fork injects V and fold extracts W, the gap $(V - W)$ is accounted for by venting. The TopologicalCompressor's per-chunk `poisoned` counts are calorimetry readings – measuring how much energy the system vented as waste heat.

The thermodynamic efficiency: $\eta = W/V = W/(W + Q_{\text{total}})$. A perfectly efficient system would vent nothing. But this is impossible for the same reason a Carnot engine can't reach 100% – you need the losers to *prove* the winner is optimal. The waste heat is the cost of certainty.

6.6 Backpressure as Conservation

Backpressure – slowing producers when consumers can't keep up – is energy conservation. When input flow rate exceeds processing capacity, energy accu-

mulates without bound (buffers overflow, the system crashes). Backpressure throttles Φ_{in} to maintain $dE/dt \leq C$.

In the rotational frame (the Wallington Whip), backpressure is conservation of angular momentum: $L = I\omega = \text{const}$. When fork increases I (more paths at large radii), ω decreases. When fold decreases I (paths removed, mass concentrated), ω increases. The whip-crack from §6.3 of the pipeline volume is exactly this: fold reduces I, angular velocity spikes, throughput surges.

6.7 The Carnot Limit

No fork/race/fold system can beat Shannon entropy:

$$W_{\max} = H(X) = -\sum p(x) \log_2 p(x)$$

This is the Carnot limit – the theoretical maximum efficiency. The two-level stream race (§9.3) approaches this limit by ensuring the best available codec always wins. But "best available" is bounded by "best theoretically possible." Brotli is already close to the Carnot limit for text. Racing brotli against itself cannot beat brotli. The topology's value is *reliably reaching* the Carnot limit across diverse inputs without prior knowledge of which codec is optimal.

6.8 The Pipeline as an Energy Diagram

The Triangle (§0.1) is an energy envelope:

- **Ramp-up (fork):** Energy increases as items enter. Each new item adds potential energy. The pipeline fills.
- **Plateau (race):** Energy is steady-state. Items enter and exit at the same rate. Maximum kinetic energy.
- **Ramp-down (fold):** Energy decreases as items exit without replacements. Potential converts to work.

The area under the curve is total energy processed. Turbulent multiplexing (§7.2) fills the triangles – the idle slots in ramp-up/ramp-down are wasted potential energy. The Worthington Whip (§7.3) reshapes one tall triangle into multiple short, wide rectangles – same total energy, better geometry, higher utilization.

6.9 Three Conservation Laws

First Law (energy conservation). $V_{\text{in}} = W_{\text{out}} + Q_{\text{dissipated}}$. Every byte forked is accounted for.

Second Law (entropy increase). Fold is irreversible. $S_{\text{folded}} \geq S_{\text{forked}}$. This is why fold is the arrow of time.

Third Law (minimum overhead). Even at perfect compression, the frame headers remain. The 10-byte self-describing header is ground-state energy –

irreducible overhead. $\lim_{T\to 0} S = S_0 > 0$. This is why tiny payloads have negative compression ratios.

The complete energy mapping:

Fork/Race/Fold	Energy Mechanics	Conservation Law
Fork	Potential energy V	Injected from input
Race	$V \to K$ conversion	$dV/dt = -dK/dt$
Fold	$K \to W$ extraction	$W = K_{\text{winner}}$
Vent	$V \to Q$ dissipation	$Q = V - K$
β_1	Energy reservoir count	$V \sim \beta_1 \cdot \bar{m} \cdot \bar{s}$
Frame header	Ground-state energy	$S_0 > 0$
Shannon entropy	Carnot limit	$W_{\max} = H(X)$
Compression ratio	Thermodynamic efficiency	$\eta = W/V$
Backpressure	Angular momentum conservation	$L = I\omega$
Pipeline Triangle	Energy envelope	Area = total energy

6.10 Transformers Are Fork/Race/Fold Graphs

The energy framing reveals that convolutional neural networks and transformers are not analogous to fork/race/fold – they *are* fork/race/fold graphs.

Multi-head attention is fork/race/fold. The input splits into N heads (each with its own Q, K, V projections). This is fork: $\beta_1 = N - 1$. All heads compute attention over the same sequence simultaneously – race. Concatenation plus linear projection – fold: the merger function f that produces a single representation. Softmax suppression (low-attention scores $\to\sim 0$) is continuous venting: the system shedding paths that don't contribute.

Feed-forward layers are fork/fold. The input expands from d_{model} to $4 \times d_{\text{model}}$ – fork into a 4x wider representation. The activation function (ReLU, GELU) is venting: zeroing or suppressing non-contributing neurons. The contraction back to d_{model} is fold.

Residual connections are fork with two paths. The skip connection and the transformed path: fork(identity, transform). Addition is fold via sum. $\beta_1 = 1$.

CNNs follow the same pattern per spatial region. N filters applied to the same receptive field is fork. All filters compute simultaneously is race. Pooling is fold – and max pooling is literally winner-take-all fold.

Transformer Component	Primitive	β_1	Energy Role
Multi-head attention (N heads)	fork/race/fold	$N-1$	N potential energy reservoirs racing
FFN expansion (4×)	fork/vent/fold	3	Expand to explore, vent dead neurons, fold back
Residual connection	fork/fold	1	Two-path fork, additive fold
Softmax attention	continuous vent	–	Shed low-energy paths smoothly
Dropout	stochastic vent	–	Random path removal (training regularization)
Layer norm	measure	0	Non-destructive observation of statistics
MoE routing (K of N experts)	fork/race/vent/fold	$N-1$	$N-K$ experts vented per token

The entire transformer is a **nested** fork/race/fold/vent graph. Each layer is fork/fold. Each attention computation within a layer is fork/race/fold. The stack of L layers is a pipeline. The architecture is the Wallington Rotation applied recursively.

Backpropagation is the energy accounting. The loss function measures thermodynamic efficiency: how much of the input's potential energy was converted to useful work (correct predictions) versus waste heat (incorrect predictions). The gradient $\partial Q/\partial\theta$ – the derivative of waste heat with respect to parameters – tells the optimizer how to adjust the system so the next forward pass vents less. Training is iteratively reducing Q while preserving W. Convergence is thermodynamic equilibrium.

Mixture of Experts makes the topology explicit. MoE routing with N experts, top-K selection: fork to N experts ($\beta_1 = N-1$), race the router's gating scores, fold the top-K results, vent the remaining $N-K$. The router *is* the race primitive. The gating function *is* the fold function. The unused experts *are* vented paths. The sparse activation pattern *is* the vent ratio $\rho = (N-K)/N$. What the ML community calls "conditional computation" is what this paper calls fork/race/fold with selective venting.

6.11 Fundamental Physics Is Fork/Race/Fold

The thermodynamic framing is not merely *analogous* to physics. Two results from fundamental physics are structurally identical to fork/race/fold, with quantitative predictions.

The Feynman Path Integral (Grade A) In quantum electrodynamics, the probability amplitude for a particle traveling from point A to point B is:

$$\mathcal{A}(A \to B) = \sum_{\text{paths}} e^{iS[\text{path}]/\hbar}$$

where S is the action along each path. The particle takes *every* possible path simultaneously. This is not a computational convenience – it is the physics. The calculation has four phases:

1. **Fork.** The particle enters all possible trajectories simultaneously. Each trajectory is a path with phase $e^{iS/\hbar}$. This is fork: one input $\to$ uncountably many paths. $\beta_1 \to \infty$.
2. **Race.** Each path propagates with its own phase accumulation. No path "knows" about the others during propagation. This is race: parallel, independent, timeless (unitary evolution is time-reversible).
3. **Fold.** The amplitudes sum. Constructive interference concentrates amplitude on the classical path (stationary phase). This is fold: many paths $\to$ one probability amplitude. $\beta_1 \to 0$.
4. **Vent.** Destructive interference eliminates non-classical paths. Their amplitudes cancel to zero. This is vent: paths that contribute no useful work are dissipated. "Propagate down, never across" – destructively interfered paths do not affect the surviving amplitude.

The classical limit ($\hbar \to 0$) recovers the path of stationary action – the unique classical trajectory. This is the $\beta_1 = 0$ subspace: one path, no fork, no race, no vent. Classical mechanics is the degenerate case of quantum fork/race/fold, just as sequential pipelines are the degenerate case of the Wallington Rotation.

This is not an analogy. The path integral *is* a fork/race/fold computation. The sum over paths *is* the fork. Interference *is* the fold/vent. The stationary phase approximation *is* the $\beta_1 \to 0$ projection. Feynman diagrams are the computation graphs, and their topological properties (β_1 = loop order) determine the difficulty of the calculation – exactly as β_1 determines pipeline complexity in §3.

The Physics Hierarchy: Progressive Folds The path integral, the Schrödinger equation, and Newton's laws are not three separate theories. They are three levels of fold applied to the same fork/race/fold computation. Each level destroys information and reduces β_1.

Level 0: The Path Integral (full fork/race/fold). All paths. All interferences. No approximation. $\beta_1 \to \infty$.

$$\mathcal{A}(A \to B) = \int \mathcal{D}[x(t)]\, e^{iS[x(t)]/\hbar}$$

Level 1: The Schrödinger Equation (the differential form of race). Feynman showed [22] that evaluating the path integral in the limit of infinitesimal time steps recovers the Schrödinger equation exactly:

$$i\hbar \frac{\partial \psi}{\partial t} = \hat{H}\psi$$

This is not a separate postulate – it is what happens when you demand the fork/race/fold computation be expressible as a *local* differential equation. The wave function ψ is the bookkeeping device that tracks the superposition of all racing paths at each instant. $|\psi|^2$ is the probability density – the energy distribution across surviving paths. The Hamiltonian $\hat{H}$ is the race operator: it governs how potential converts to kinetic at each infinitesimal step.

The Schrödinger equation is the race phase written as a differential equation. It is the local form of a global fork/race/fold computation, just as Maxwell's equations are the local form of global electromagnetic phenomena. The wave function ψ carries all the information about which paths are still racing and with what amplitude. It is a race snapshot.

Quantized energy levels are fold constraints. For bound systems (electrons in atoms, particles in wells), the Schrödinger equation admits only discrete solutions – specific energy eigenvalues. These are not inputs to the equation; they *emerge* from the fold boundary conditions. The requirement that $\psi \to 0$ at infinity (the wave function must be normalizable) is a fold constraint: it eliminates all solutions that don't converge. The surviving eigenvalues are the fold results. Lasers, LEDs, atomic clocks, and MRI machines all depend on these quantized fold outputs.

Quantum tunneling is incomplete venting. Classically, a particle encountering a potential barrier higher than its kinetic energy is vented – it cannot cross. But the Schrödinger equation shows that ψ doesn't drop to zero at the barrier; it decays exponentially through it. If the barrier is thin enough, nonzero amplitude leaks through. This is a path that *should* have been vented but survived because $\hbar > 0$ means the vent is imperfect – some amplitude always leaks. Flash memory, scanning tunneling microscopes, and nuclear fusion in stars all exploit incomplete venting.

Level 2: Stationary Phase Approximation (the vent operator). In the classical limit ($\hbar \to 0$), the phase $e^{iS/\hbar}$ oscillates infinitely fast. Nearly all paths cancel by destructive interference – they are vented. Only paths near the stationary point of the action survive:

$$\delta S = 0 \implies \text{Euler-Lagrange equations} \implies F = ma$$

The stationary phase approximation *is* the vent operator applied maximally. It destroys all path information except the single classical trajectory. $\beta_1 \to 0$. The void (β_2) becomes infinite – uncountably many quantum paths are vented, leaving one survivor.

Level 3: Newton's Laws ($\beta_1 = 0$, fully folded). One path. Deterministic. No fork, no race, no vent. $F = ma$ is the maximally folded result of the path integral. Classical mechanics is not "wrong" – it is the $\beta_1 = 0$ degenerate case, just as sequential pipelines are the degenerate case of the Wallington Rotation.

The hierarchy:

Level	Theory	Fork/Race/Fold Role	β_1	Information
0	Path integral	Full engine	∞	All paths, all phases
1	Schrödinger equation	Differential race	Finite	Wave function ψ tracks superposition
2	Stationary phase	Maximal vent	$\to 0$	Only near-classical paths survive
3	Newton's laws	Fully folded	0	One path, deterministic

Each level is a fold. Each fold destroys information. Each fold creates void (β_2 grows). The entire tower of classical physics – from Newton to Lagrange to Hamilton – is nested fold operations on the path integral. You cannot unfold Newton back to Schrödinger, or Schrödinger back to the path integral, without injecting new information (re-forking). This is the Second Law applied to physical theory itself.

Band theory is the covering space formalism applied to crystals. When the Schrödinger equation is solved for electrons in a periodic lattice (silicon, germanium), Bloch's theorem states that solutions have the form $\psi_k(r) = e^{ik\cdot r}u_k(r)$ where u_k has the periodicity of the lattice. The periodic lattice is the base space. The electron's wave function in the full crystal is the covering space. Bloch's theorem is the covering map (§3.3) – it relates the global behavior (energy bands) to the local structure (unit cell). The band gap – the energy range where no electron states exist – is the void ($\beta_2 > 0$). Semiconductors, transistors, solar cells, and every chip in existence depend on this covering space structure.

The Virial Theorem (Grade A-) For self-gravitating systems in equilibrium (gas clouds, galaxies, star clusters), the virial theorem states:

$$2K + V = 0 \implies K = -V/2$$

Half the gravitational potential energy becomes kinetic energy (thermal motion, radiation). This gives a quantitative prediction for the fork/race/fold energy partition. A collapsing gas cloud:

1. **Fork.** Gravitational potential energy V is stored in the spatial distribution of mass. Every particle has a trajectory it *could* follow. $V = -\sum_{i<j} Gm_i m_j / r_{ij}$.
2. **Race.** Free-fall collapse. Particles accelerate toward the center. $V \to K$ conversion.
3. **Fold.** A star forms – the bound state. Useful work W is extracted as nuclear fusion becomes possible. Hydrostatic equilibrium is the fold: gravitational compression balanced by radiation pressure.
4. **Vent.** Half the energy radiates away as heat and light. $Q = V/2$. The Kelvin-Helmholtz mechanism is the vent – the star shines because it must dissipate excess energy.

The virial theorem gives the exact split: $W = V/2$, $Q = V/2$, therefore $\eta = 0.5$. This is a specific, testable prediction that the First Law produces when applied to gravity: $V_{\text{fork}} = W_{\text{fold}} + Q_{\text{vent}}$ with the virial theorem constraining the partition ratio.

The Weak Force as Vent Operator (Grade B+) Beta decay: $n \to p + e^- + \bar{\nu}_e$. The neutrino carries away energy that is never recovered – it barely interacts with matter and propagates away permanently. This is venting: "propagate down, never across." The weak force is how unstable nuclear configurations dissipate excess energy to reach stable states.

Supernovae are the extreme case: 99% of the gravitational binding energy ($\sim 3\times10^{46}$ J) is carried away by neutrinos. The visible explosion – light, shock wave, ejecta – is only $\sim 1\%$. The vent-to-work ratio: $Q/W \approx 99$. Thermodynamic efficiency $\eta \approx 0.01$. The weak force is nature's most aggressive vent operator.

Color Confinement as Anti-Vent (Grade B) The strong force exhibits a property with no analogue in the other nine connections. If you try to separate two quarks (attempt to vent a color-charged path), the energy stored in the color field creates new quark-antiquark pairs. Attempted vent $\to$ automatic fork. The strong force *prevents* venting by forking.

Color confinement enforces topological closure: $\beta_2 = 0$ for color charge. There are no unreachable color states, no void. Every configuration must be color-neutral. You try to create a void; the strong force fills it by forking new particles. This is conservation of topological closure – a constraint stronger than anything in the computational domain, where venting is always permitted.

Symmetry Breaking as Fold (Grade B+) The Higgs mechanism: above the electroweak energy scale (~ 246 GeV), the electromagnetic and weak forces are unified. Below it, the Higgs field selects one vacuum state from a continuous family of equivalent states. The Mexican hat potential is a fork/race/fold landscape:

- **Fork:** The symmetric state at the top of the potential (all vacuum directions equivalent)
- **Race:** The field rolls down the brim (explores vacuum states)
- **Fold:** Settles into one minimum (symmetry broken, particles acquire mass)
- **Vent:** Goldstone bosons carry away the broken symmetry degrees of freedom (three of four are "eaten" by the $W^{\pm}$ and Z bosons, becoming their longitudinal polarization)

Spontaneous symmetry breaking is fold: many equivalent states $\rightarrow$ one selected state. The void (β_2) is the set of unchosen vacua. The universe's particle masses are the fold result.

6.12 The Optimality Diagnostic

If fork/race/fold is the shape that any system converges to under conservation, irreversibility, and minimum overhead, then finding this shape in a system is evidence that the system is operating near its theoretical optimum. Not finding it – where the problem's intrinsic topology demands it – is a diagnostic for waste.

Every problem has an **intrinsic Betti number** β_1^*: the number of independent parallel paths that the problem's structure supports. A blood test, an MRI, and a genetic screen are independent – $\beta_1^* = 2$. Eight compression codecs applied to the same chunk are independent – $\beta_1^* = 7$. The N paths in a Feynman path integral are independent – $\beta_1^* \rightarrow \infty$. This is not a design choice. It is a property of the problem.

Every implementation has an **actual Betti number** β_1: the number of independent parallel paths in the system as built. A sequential referral chain has $\beta_1 = 0$. A fork with 8 codecs has $\beta_1 = 7$. The gap between β_1^* and β_1 is the **topological deficit**:

$$\Delta_\beta = \beta_1^* - \beta_1$$

When $\Delta_\beta = 0$, the system's topology matches the problem's topology. It is operating at its natural parallelism. When $\Delta_\beta > 0$, the system is forcing a high-β_1 problem through a low-β_1 pipe. The deficit is wasted parallelism – performance left on the table.

I define the unit of topological deficit as the **Bule** (symbol: **B**). One Bule equals one unit of Δ_β – one independent parallel path that the problem supports but

the implementation does not exploit.

$$1 \text{ B} = 1 \text{ unit of } \Delta_\beta = \beta_1^* - \beta_1$$

A system at 0 B is topologically optimal. A system at 3 B is wasting three independent parallel paths. The Bule is dimensionless, integer-valued, and directly measurable from the computation graph.

This also gives a testable meaning to **computational aesthetics**: elegance is the degree of fit between implemented topology and problem topology. In this framing, Bules are an aesthetic meter. Low-Bule systems feel natural because structure and task align; high-Bule systems feel strained because the structure is fighting the work.

The topological deficit predicts real-world waste.

System	β_1^*	β_1	Deficit	Observable Waste
Healthcare diagnosis	≥ 3	0 (referral chain)	≥ 3 B	4.8-year average diagnostic delay for rare diseases [21]
Financial settlement	2	0 (T+2 sequential)	2 B	$70 trillion unnecessarily locked capital [21]
HTTP/2 multiplexing	N_{streams}	0 (TCP substrate)	N B	Head-of-line blocking on any packet loss
Photosynthetic antenna	7 (pigments)	~ 7 (quantum coherence)	0 B	>95% energy transfer efficiency
Path integral	∞	∞	0 B	Exact quantum-mechanical predictions
DNA replication	1 (lagging strand)	1 (Okazaki fragments)	0 B	Replication matches leading strand speed
Saltatory conduction	nodes -1	nodes -1	0 B	100x speedup vs. continuous conduction

The pattern: **every system where $\Delta_\beta = 0$ operates at or near its theoretical efficiency.** Every system where $\Delta_\beta > 0$ has measurable, quantifiable waste. The deficit is not abstract – it maps to years of diagnostic delay, trillions of locked capital, and protocol-level blocking.

This yields a practical diagnostic tool:

1. **Measure** β_1^*: analyze the problem's dependency structure to find its intrinsic parallelism. Independent inputs are independent paths. Sequential dependencies are constraints that reduce β_1^*.
2. **Measure** β_1: count the actual parallel paths in the implementation. A sequential pipeline has $\beta_1 = 0$. A fork with N paths has $\beta_1 = N - 1$.
3. **Compute** Δ_β: the gap is the optimization opportunity. If $\Delta_\beta > 0$, the system is suboptimal by construction – no amount of micro-optimization within the existing topology can recover the parallelism that the topology itself prevents.

The converse is equally powerful. When you encounter a system that already exhibits fork/race/fold at $\Delta_\beta = 0$ – photosynthesis, DNA replication, myelinated nerve conduction – you can infer that evolution has found the optimal topology for that problem. The billions of years of selection pressure have converged to the same shape that the three constraints predict. The Betti numbers match because there is no better shape to converge to.

This is why the biological examples in §1 are not decoration. They are **existence proofs of optimality.** When we observe that *Physarum* constructs transport networks matching the Tokyo rail system, we are observing $\Delta_\beta \approx 0$ achieved without engineering – the slime mold's topology matches the problem's topology because evolution found the fork/race/fold shape. When we observe that photosynthetic antenna complexes achieve >95% energy transfer efficiency, we are observing $\Delta_\beta = 0$ at quantum scale – the system's β_1 matches the problem's β_1^* because physics itself operates at fork/race/fold.

The optimality diagnostic also explains **why quantum computing promises speedups**. Classical computation forces $\beta_1 = 0$ – one computational path at a time. Quantum computation operates at the problem's intrinsic β_1^*, which for search problems is $\sqrt{N}$ (Grover) and for factoring is $\log N$ (Shor). The quantum speedup *is* the topological deficit that classical computing pays: $\Delta_\beta = \beta_1^* - 0 = \beta_1^*$. Quantum computers close the deficit. The speedup is not magic – it is the system finally matching the problem's natural topology.

6.13 Map/Reduce as a Quantum-Readiness Heuristic (Not a Theorem)

Map/reduce should be interpreted topologically:

- **Map** is fork over independent partitions.
- **Reduce** is fold under an associative/deterministic merger.
- **Shuffle** is the routing layer between the two.

In this sense, map/reduce is a constrained fork/fold system with no explicit race or vent semantics. That constraint is exactly why map/reduce usage is a useful signal: it usually means the workload already has an exposed parallel frontier and a valid fold boundary.

This leads to a practical claim:

> **Heuristic claim.** Sustained map/reduce usage is evidence of **topology readiness** for Wallington pipelines (fork/race/fold + vent), and therefore evidence of potential quantum-readiness in the narrow structural sense (the problem admits concurrent path exploration and deterministic projection).

This is **not** a claim of automatic quantum advantage. It does **not** imply Grover/Shor-style asymptotics. It only claims structural compatibility.

I separate readiness from opportunity:

$$Q_{\mathrm{mr}} = I_{\mathrm{map}} \cdot A_{\mathrm{reduce}} \cdot (1 - S_{\mathrm{key}}) \cdot Z_{\mathrm{copy}}$$

$$O_\beta = \max\left(0, \frac{\Delta_\beta}{\max(1, \beta_1^*)}\right)$$

$$R_{\mathrm{qr}} = Q_{\mathrm{mr}} \cdot O_\beta$$

where all factors are normalized to $[0, 1]$:

- I_{map}: fraction of map work that is truly independent.
- A_{reduce}: reducer associativity/determinism score.
- S_{key}: partition skew (hot-key concentration).
- Z_{copy}: zero-copy ratio across map/shuffle/fold boundaries.
- O_β: topological opportunity from the Bule deficit.

Interpretation:

- High Q_{mr}, low O_β: architecture is ready, but little headroom (already near $\Delta_\beta = 0$).
- Low Q_{mr}, high O_β: headroom exists, but map/reduce quality is too poor to realize it safely.
- High R_{qr}: migrate to full Wallington primitives (add race + vent + Reynolds-driven multiplexing).

So map/reduce is best treated as a **screening diagnostic**: it flags workloads likely to benefit from promotion into fork/race/fold, and in a subset of cases that same structure is the precondition for quantum speedup. The value is practical: it tells us where topology work is likely to pay off.

Executable Diagnostic Tool The topological deficit is not just a theoretical quantity. The `@affectively/aeon` package [20] includes a `TopologyAnalyzer` that computes Betti numbers and Bules from a computation graph, and a `TopologySampler` that instruments a running system to measure deficit over time:

```
import { TopologyAnalyzer, TopologySampler } from '@affectively/aeon';

// Static analysis: is this system wasting parallelism?
const graph = TopologyAnalyzer.fromForkRaceFold({
  forkWidth: 1,          // implementation: sequential
  intrinsicBeta1: 7,     // problem: 8 independent codecs
});
const report = TopologyAnalyzer.analyze(graph);
// → deficit: 7 Bules -- "Sequential bottleneck: 7 Bules of waste"

// Fix: match the topology
const fixed = TopologyAnalyzer.fromForkRaceFold({
  forkWidth: 8,          // implementation: fork 8 codecs
  intrinsicBeta1: 7,     // problem: 8 independent codecs
});
const fixedReport = TopologyAnalyzer.analyze(fixed);
// → deficit: 0 Bules -- "Optimal: 0 Bules"

// Runtime sampling: how does the deficit evolve?
const sampler = new TopologySampler({ intrinsicBeta1: 7 });
sampler.fork('chunk-1', ['raw', 'rle', 'delta', 'lz77',
                         'brotli', 'gzip', 'huffman', 'dict']);
// → currentDeficit(): 0 Bules (all 8 codecs racing)
sampler.race('chunk-1', 'brotli');
sampler.vent('chunk-1', 'raw');
// ... vent remaining losers ...
sampler.fold('chunk-1');
const samplerReport = sampler.report();
// → peakBeta1: 7, efficiency: 0.125 (1 race / 8 events)
```

The `TopologyAnalyzer` computes β_0, β_1, β_2 and detects fork/join pairs from any directed graph. The `TopologySampler` records fork/race/vent/fold events at runtime and produces time-series utilization data. Both are validated by 24 passing tests covering sequential pipelines, fork/join graphs, void detection, deficit measurement, concurrent forks, vent ratios, and the real-world topologies from this section [20].

To see fork/race/fold is to see a system that has found its shape. To not see it is to see a system that hasn't. The Bule count tells you how far off you are.

7. Instantiation A: Distributed Staged Computation

I implement fork/race/fold in a distributed computation engine with processing stages partitioned across networked nodes – a domain of particular interest to the researcher.

7.1 Chunked Pipelined Prefill (Wallington Rotation)

In the baseline, a workload of P items is processed sequentially through N stage nodes: $P \times N$ round-trips. The key insight: each node's forward pass for item t_i depends only on that node's accumulated state from t_{i-1} – a stage-local constraint (C1). This enables pipelining. Chunking groups B items per forward pass via causal masking.

Scenario	Serial ($P \times N$)	Chunked Pipeline	Speedup
14 tokens, 2 nodes	28 steps	9 steps	3.1x
100 tokens, 4 nodes	400 steps	7 steps	57x
500 tokens, 8 nodes	4,000 steps	15 steps	267x
100 tokens, 10 nodes	1,000 steps	19 steps	53x

7.2 Turbulent Multiplexing

When $C \approx N$, 43% of node-slots are idle during ramp-up/ramp-down. Turbulent multiplexing fills idle slots with items from concurrent requests, maintaining per-request vent isolation (C2). This is what polysomes do: fill the mRNA pipeline with multiple ribosomes, degrade the mRNA when Re drops below threshold, reallocate to active pipelines.

7.3 Worthington Whip (Superposition Sharding)

A single workload is sharded across S parallel pipelines. Each shard processes P/S items, then cross-shard correction reconciles at fold. Per-shard compute savings: $(S-1)/2S$.

7.4 Speculative Tree

A lightweight predictor generates K candidate continuations (fork). All K branches enter the pipeline as multiplexed sub-requests (race). A verifier checks all K in a single batched pass. Invalid branches are pruned via venting. Expected items accepted per pass with acceptance rate α: $(1-\alpha^K)/(1-\alpha)$.

8. Instantiation B: Aeon Flow Protocol

8.1 Design Principle

The patterns – fork, race, fold, vent – recur identically in edge composition, service worker preloading, fragment assembly, deploy artifact streaming, CRDT synchronization and ten independent domains validated in §11. Rather than reimplementing per domain, I extract the primitive into a binary wire protocol on UDP.

8.2 Wire Format

```
Offset  Size   Field
[0..1]  u16    stream_id    (multiplexed stream identifier)
[2..5]  u32    sequence     (position within stream)
[6]     u8     flags        (FORK=0x01 | RACE=0x02 | COLLAPSE=0x04 | POISON=0x08 | FIN=0x10)
[7..9]  u24    length       (payload bytes, max 16 MB)
[10..]  [u8]   payload      (zerocopy Uint8Array view)
```

10 bytes. Every frame carries its own identity. Every frame is self-describing. No ordered delivery required. The `stream_id` + `sequence` pair is the coordinate in the covering space (§3.3). Flags compose: `RACE | FIN` means "racing AND final frame." The frame reassembler (§3.3) is the covering map back to sequential order. Payloads are zerocopy: the codec writes 10 bytes in front of the existing `ArrayBuffer` view.

8.2.1 The Self-Describing Frame as Pervasive Abstraction

The self-describing frame is not specific to the wire protocol. It is the unifying data structure across both the transport layer and the computation engine.

On the wire, it is the **FlowFrame** – 10 bytes of header carrying `stream_id`, `sequence`, `flags` and `length`. On the computation side, it is the **WorkFrame** – the same `(stream_id, sequence)` identity enriched with a typed payload `T` and metadata:

```
WorkFrame<T>                      FlowFrame
-------------                     ---------
streamId: StreamId                streamId: u16
sequence: number                  sequence: u32
payload:  T                       flags:    u8
metadata: Record<string,unknown>  length:   u24
emittedAt: number                 payload:  [u8]
```

The two are isomorphic. The wire format bridge encodes `WorkFrame<T>` to `FlowFrame` (serializing `T` as payload bytes) and decodes `FlowFrame` back to `WorkFrame<T>`. A computation that forks 10 streams in-process produces 10 `WorkFrame`s. Those same frames, encoded as `FlowFrame`s, can cross a network boundary and be reassembled on the other side by the same `FrameReassembler` algorithm. The computation topology is independent of the transport topology.

This is the same pattern as Okazaki fragments in DNA replication: each fragment carries its genomic coordinate (its `stream_id` + `sequence`), enabling out-of-order synthesis and reassembly by DNA ligase. The fragment is self-describing whether it is being synthesized on the lagging strand (in-process) or transported via a virus to another cell (on the wire). Identity is intrinsic, not assigned by context.

8.3 Why UDP Only

TCP had a magnificent 40-year run. It made the internet possible. But TCP was designed for a world where $\beta_1 = 0$ was a reasonable assumption – one connection, one stream, one conversation. The moment you need $\beta_1 > 0$, every TCP guarantee becomes a liability:

TCP Guarantee	Why It Hurts
Ordered delivery	One lost packet on stream A blocks *all* streams behind it
Connection handshake	1.5 RTT before first data byte
Single-stream congestion	TCP backs off the entire connection on loss
Connection-level retransmit	Stream A's retransmit delays stream B

HTTP/2 tried to multiplex streams over TCP. The application topology ($\beta_1 > 0$) contradicts the transport topology ($\beta_1 = 0$). Head-of-line blocking is the topological symptom (§3.4). HTTP/3 (QUIC) partially resolves this with per-stream loss recovery on UDP, but maintains ordered delivery within each stream and carries substantial framing complexity.

Aeon Flow does not patch TCP's problems. It starts from the topology and asks: what wire format does $\beta_1 > 0$ actually need? The answer: self-describing frames with no ordered delivery. AIMD congestion control per-stream (not per-connection). MTU-aware fragmentation (4-byte fragment header, 255 fragments × 1,468 bytes). ACK bitmaps (14 bytes covering 64 sequences). The protocol is 800 lines of TypeScript. In the shootoff benchmarks, it beats HTTP/3 on every measured framing and latency metric because it has no topological contradictions to resolve – the shape of the wire matches the shape of the problem from the first byte [21].

8.4 Protocol Comparison

Metric	HTTP/1.1	HTTP/2	HTTP/3 (QUIC)	Aeon Flow
Per-resource overhead	~200 bytes	~30 bytes (HPACK)	~20 bytes	**10 bytes**
Header fraction (12 resources)	4.2%	0.6%	0.4%	**0.2%**

Metric	HTTP/1.1	HTTP/2	HTTP/3 (QUIC)	Aeon Flow
Header fraction (95 resources)	31.0%	5.8%	4.1%	**1.5%**
Connections for full site	6+	1	1	**1**
Head-of-line blocking	Yes (conn)	Yes (TCP)	No (per-stream)	**No**
Native fork/race/fold	No	No	No	**Yes**
Vent propagation	N/A	RST_STREAM	STOP_SENDING	**Recursive tree**
Transport	TCP	TCP	UDP (QUIC)	**UDP (raw)**
Ordered delivery	Required	Required	Per-stream	**None**
Topological contradiction	N/A	β_1 mismatch	Partial	**None**

8.5 Shootoff: Head-to-Head Protocol Benchmarks

I benchmark Aeon Flow against HTTP/1.1, HTTP/2 and HTTP/3 with realistic compression (gzip, brotli) across two site profiles. All protocols use identical payloads; only framing and transport differ.

Big Content Site (12 resources, ~2.5 MB – large JS bundles, hero images, web fonts):

Protocol	Wire Size	Framing Overhead	Overhead %	RTTs
HTTP/1.1	913 KB	8.2 KB	0.89%	3
HTTP/2	907 KB	1.6 KB	0.18%	2
HTTP/3 (QUIC)	906 KB	906 B	0.10%	1
Aeon Flow	**905 KB**	**276 B**	**0.03%**	**1**

For large payloads, all modern protocols converge – but Aeon Flow's framing is **3.3x smaller than HTTP/3** (276 B vs 906 B).

Microfrontend Site (95 resources, ~1.8 MB – 45 JS modules, 16 CSS modules, 20 SVG icons):

Protocol	Wire Size	Framing Overhead	Overhead %	RTTs
HTTP/1.1	187 KB	58.1 KB	**31.0%**	16
HTTP/2	137 KB	8.0 KB	5.8%	2
HTTP/3 (QUIC)	135 KB	5.9 KB	4.4%	1
Aeon Flow	**131 KB**	**1.9 KB**	**1.5%**	**1**

This is where the topology matters. HTTP/1.1 wastes **31% of total bandwidth on headers** – nearly a third of the wire is framing, not data. HTTP/2 reduces this to 5.8%. HTTP/3 to 4.4%. Aeon Flow: **1.5%**. That is a **21x reduction** in framing overhead versus HTTP/1.1 and **3x versus HTTP/3**.

At 100ms RTT, HTTP/1.1 needs 16 round trips (1.6 seconds of pure latency). Aeon Flow: 1 round trip. The difference is not optimization – it is topology. HTTP/1.1 has $\beta_1 = 0$ (one request per connection, six connections). Aeon Flow has $\beta_1 = 94$ (95 streams, one connection). The framing overhead is the cost of forcing a high-β_1 problem through a low-β_1 pipe.

9. Instantiation C: Topological Compression

9.1 The Claim and Its Limits

The same fork/race/fold primitive applies to compression. **Topological compression** forks all available codecs per chunk, races them and folds to the winner. Each chunk independently selects its best codec. The output is a sequence of self-describing frames (9-byte header: codec ID, original size, compressed size). $\beta_1 = \text{codecs} - 1$.

I implement this with eight codecs:

ID	Codec	Type	Best on
0	Raw (identity)	Pure JS	Incompressible data
1	RLE	Pure JS	Repeated byte runs
2	Delta	Pure JS	Sequential/incremental data
3	LZ77	Pure JS	Repeated patterns
4	Brotli	Platform (node:zlib)	General text
5	Gzip	Platform (node:zlib)	General text (universal fallback)
6	Huffman	Pure JS	Skewed byte distributions
7	Dictionary	Pure JS	Web content (HTML/CSS/JS keywords)

9.2 What the Benchmarks Actually Show

I benchmark across both sites on Aeon Flow transport. The results are honest:

Big Content Site (12 resources, ~2.22 MB):

Compression	Wire Size	Ratio	β_1
Brotli (global, quality 4)	905 KB	39.8%	0
Topo-full (8 codecs per-chunk)	1005 KB	44.2%	7
Topo-pure (6 pure-JS codecs per-chunk)	1.17 MB	52.5%	5

Microfrontend Site (95 resources, ~617 KB):

Compression	Wire Size	Ratio	β_1
Brotli (global, quality 4)	131 KB	20.9%	0
Topo-full (8 codecs per-chunk)	159 KB	25.4%	7
Topo-pure (6 pure-JS codecs per-chunk)	229 KB	36.8%	5

Standalone brotli wins on compression ratio. On these benchmarks – homogeneous web content – global brotli beats per-chunk topological compression by 4–15 percentage points. This is not surprising: brotli compresses the entire stream with a sliding window that builds dictionary context across chunks. Per-chunk compression resets the dictionary every 4096 bytes.

The two-level race (§9.3) confirms this. When given the choice between global brotli and per-chunk topological, it picks global brotli every time on these payloads, matching standalone brotli's ratio plus 5 bytes of strategy header. Per-chunk topological never wins the stream-level race on homogeneous content because the 9-byte per-chunk header tax and the loss of cross-chunk dictionary context always exceed whatever the per-chunk adaptive selection saves.

9.3 Two-Level Stream Race

I extend the topology to race at two levels:

```
fork (stream level):
  |- Path 0: Per-chunk topological (8 codecs × each 4096-byte chunk)
  |- Path 1: Global brotli (entire stream, cross-chunk dictionary)
  |- Path 2: Global gzip (entire stream)
  `- ...
race → smallest total output wins
fold → 5-byte strategy header + compressed data
```

This guarantees that topological compression is **never worse** than standalone brotli – because brotli is one of its racing paths. On these benchmarks it is never **better** either. The 5-byte stream header is the only overhead.

9.4 What the Topology Actually Provides

If topological compression does not beat brotli on ratio, what is the point?

1. Subsumption, not superiority. The topology is the space in which brotli competes. Brotli at $\beta_1 = 0$ is a degenerate case of topological compression at $\beta_1 = 7$. The two-level race includes brotli as a contestant. If brotli is best, the topology selects it. If something better appears tomorrow – a learned codec, a neural compressor, a domain-specific dictionary – it enters the race without changing the architecture. The `TopologicalCompressor` is unchanged; only the codec array grows.

2. Platform independence. Brotli requires `node:zlib` (Node, Bun, Deno). In browsers and Cloudflare Workers, it is unavailable. Topo-pure – six codecs in pure JavaScript, zero dependencies – achieves 36.8% ratio on the microfrontend with no native code. The topology degrades gracefully: full ratio when brotli is available, reasonable ratio when it is not.

3. Per-chunk random access. The per-chunk format enables decompression of individual chunks without processing the entire stream. For seeking into large payloads, resuming interrupted transfers, or parallel decompression, this is structurally impossible with global brotli.

4. Adaptive codec selection on heterogeneous data. On the per-chunk level, different regions of the input genuinely select different codecs. The shootoff shows 3 distinct codecs winning across 151 chunks on realistic web content (brotli for text chunks, dictionary for web-pattern-heavy chunks, raw for incompressible binary). No single algorithm achieves this without the topology.

5. The real compression win is framing, not codecs. The paper's compression contribution is not beating brotli's ratio. It is the 30× reduction in framing overhead (§8.4): Aeon Flow uses 1.9 KB of framing for 95 resources where HTTP/1.1 uses 56.3 KB. On the microfrontend, framing overhead drops from 8.4% to 0.3% of the payload. This saving is orthogonal to which codec compresses the content.

9.5 Honest Assessment

The per-chunk topological approach pays a real cost: 9 bytes per chunk of header overhead and the loss of cross-chunk dictionary context. On homogeneous content (which describes most web payloads), this cost exceeds the benefit of adaptive codec selection. Global brotli, with its full-stream dictionary, simply compresses text better than any per-chunk approach can.

The two-level stream race eliminates this disadvantage by including global brotli as a racing path. But it also reveals that per-chunk topological compression, as currently implemented, is not the winning strategy for web content. It is a structurally sound framework that provides platform independence, random access and future extensibility – at the cost of matching, not beating, the state of the art on ratio.

The progression 4 codecs ($\beta_1 = 3$) → 6 codecs ($\beta_1 = 5$) → 8 codecs ($\beta_1 = 7$) demonstrates the covering space property: each expansion improved pure-JS

compression without changing the base space. But adding brotli and gzip to the race, while improving per-chunk results, still cannot overcome the fundamental advantage of global dictionary context.

The topology subsumes the algorithm. It does not necessarily surpass it. I will not pretend this is the result I wanted. I wanted to dust brotli – to show that topological compression, by racing eight codecs per chunk with adaptive selection, could outperform decades of careful engineering in a C-native algorithm. It does not. Global brotli with its full-stream dictionary context is genuinely excellent at compressing web content, and no amount of per-chunk cleverness overcomes the fundamental information advantage of seeing the entire input at once. The honest conclusion is that the topology provides structural guarantees – subsumption, platform independence, random access, extensibility – but not ratio superiority. That is still worth having. It is not the same as winning.

Executable evidence is available in two independent suites: the companion topological-compression obligations [21] and the production `TopologicalCompressor` tests in the open-source `@affectively/aeon` package [20]. Together they verify per-chunk adaptive winner selection, 9-byte self-describing chunk headers, codec vent behavior (discarding expansions), two-level stream race strategy selection, $\beta_1 = \text{codecs} - 1$ invariants and roundtrip correctness across edge cases and large payloads.

9.6 Applications

Application	Fork	Race	Fold
Site preloading	Stream all assets as parallel frames	First complete asset wins cache slot	SW stores all in Cache API
ESI composition	Fork stream per directive	Race cache vs. compute	Assemble into final page
Deploy artifacts	Fork per build artifact	Stream concurrently	Receive complete deployment
CRDT sync	Fork per-peer delta streams	Race peers to contribute	Merge deltas into canonical state
Speculative nav	Fork predicted route preloads	Race prediction vs. actual	Display whichever resolves first

10. The Engine

The algorithm is implemented as **Aeon Pipelines** [2], a zero-dependency computation topology engine in TypeScript. It runs on Cloudflare Workers, Deno, Node, Bun and browsers. The API surface is two classes:

- **`Pipeline`**: the engine – capacity, metrics, backpressure, turbulent multiplexing.
- **`Superposition<T>`**: the builder – chainable fork/race/fold/vent/tunnel/interfere/entangle/measure/searc operations.

```
// Kids juggling balls
const result = await Pipeline
  .from([fetchFromA, fetchFromB, fetchFromC])
  .race();

// People juggling the kids
const diagnosis = await Pipeline
  .from([bloodTest, mriScan, geneticScreen])
  .poison(result => result.inconclusive)
  .tunnel(result => result.conclusive)
  .collapse('merge-all', mergeFindings);

// Grover-style search over solution space
const drug = await new Pipeline({ capacity: 64 })
  .fork(candidates.map(c => () => evaluate(c)))
  .search({
    width: 32,
    oracle: compound => compound.efficacy,
    mutate: (compound, gen) => perturb(compound, gen),
    convergenceThreshold: 0.01,
  });
```

The `search()` operation implements a classical approximation of Grover's algorithm: for N candidates with W-wide parallelism, convergence in $\sim \sqrt{N/W}$ iterations – a quadratic speedup over sequential evaluation.

10.1 Performance

The engine is fast enough to disappear. Orchestration overhead is microseconds; the bottleneck is always the user's work functions, never the topology.

Operation	Latency	Notes
`fork(10)`	**1.82 µs**	10 parallel streams created

Operation	Latency	Notes
`collapse('quorum', 5)`	**4.51 µs**	Byzantine agreement across 5 streams
`search(8×20)`	**8.3 µs**	Grover-style search, 8-wide, 20 generations
`interference(100)`	**16.3 µs**	Pairwise consensus across 100 streams
`poison-tree(13)`	**18.9 µs**	Recursive vent across 13-node tree
`flow-bridge-batch(100)`	**25.7 µs**	100 frames encoded to wire format
`reassemble-reverse(1000)`	**71.4 µs**	1,000 frames reassembled from reverse order
`flow-bridge-roundtrip`	**0.76 µs**	Single frame encode → decode

Zero dependencies. ~384 bytes per stream. ~3.5 KB per pipeline. Runs on Cloudflare Workers, Deno, Node, Bun and browsers.

10.2 Domain Validation

The same API – unchanged – was validated across 10 independent domains:

1. **Multi-venue trading**: fork/race across exchanges, vent adverse prices
2. **Healthcare diagnostics**: fork parallel tests, tunnel on conclusive, merge-all findings
3. **Financial settlement**: fork clearing/netting/DVP, merge-all for T+0
4. **Construction scheduling**: fork trades per floor, merge-all hours
5. **Emergency dispatch**: fork/race responders, first arrival wins
6. **Academic review**: fork reviewers, quorum 2/3 agreement
7. **Drug discovery**: fork compounds, Grover search to convergence
8. **Manufacturing QC**: fork sensors, consensus (constructive interference)
9. **Journal publishing**: fork reviewers, vent timeout, quorum verdict
10. **Legal review**: fork reviewers, weighted fold by seniority

Ten domains. One primitive. Same four operations. The universality is not designed – it is discovered, the same way *Physarum* discovers optimal networks without being designed to.

10.3 Wire Format Bridge

The engine includes a wire format bridge to the Aeon Flow protocol. The same 10-byte frame header (§8.2) encodes `WorkFrame<T>` objects for network transmission. Frames encoded by Aeon Pipelines decode by Aeon Flow and vice versa – the computation topology is independent of the transport topology.

11. Validation

The claims are backed by executable tests across four independent suites:

- **Companion obligations and executable proofs** (189 passing tests): pipeline topology, queueing containment, flow-frame invariants, compression race properties, shootoff reproductions, §7 formula checks (Worthington Whip $(S-1)/2S$, Speculative Tree $(1-\alpha^K)/(1-\alpha)$, turbulent multiplexing idle-fraction bounds), quantum-topology claims (Grover-style Δ_β scaling and Kronig-Penney band gaps as $\beta_2 > 0$), evidence-table deficits (e.g., T+2 settlement $\Delta_\beta = 2B$), self-hosted formal artifact parsing/round-trip validation with `aeon-logic`, a parser shootoff benchmark against Java SANY startup-parse baselines, mechanized TLA+ model checking across six formal modules (C1–C4, §7 formulas, protocol/settlement deficits, quantum deficit identity, band-gap void), and a Lean 4 theorem package with explicit-assumption theorems and theorem schemas for top-level claims verify the strongest operational claims section by section [21, 25, 26, 27].
- **Open-source flow + compression runtime** (168 passing tests): `@affectively/aeon` flow/compression tests verify 10-byte self-describing flow frames, UDP fragmentation/ACK behavior, frame reassembly, flow protocol semantics and topological compression properties [20].
- **Open-source topology engine** (114 passing tests): `@affectively/aeon-pipelines` tests cover fork/race/fold/vent primitives, fold strategies, Reynolds/backpressure/turbulent multiplexing, quantum modalities, flow-bridge wire compatibility, domain scenarios and microbenchmarks [2].
- **Open-source topology analyzer suite** (24 passing tests): `TopologyAnalyzer/TopologySampler` tests in `@affectively/aeon` validate Betti extraction, Δ_β diagnostics, β_2 void detection and executable protocol-topology contrasts [20].

Total validated tests referenced here: **495 passing tests**, plus parser-validated formal artifacts, mechanized Lean theorem builds and mechanized TLC runs across six TLA+ modules, with executable commands and source-visible assertions in the linked repositories [2, 20, 21, 25, 26, 27].

12. Limitations

Benchmark substrate. Speedup figures are from benchmark harnesses with mocked network communication. Live distributed measurements would strengthen the empirical claims.

Cross-shard cost. The Worthington Whip's correction phase scales with shard count. The crossover where correction overhead exceeds parallelism gains is not yet characterized.

Formal model scope. C1–C4, §7 formulas, protocol/settlement deficits, quantum deficit identity and band-gap void are mechanized in a two-layer stack: finite-state transition models in TLA+ (TLC), plus Lean theorems with explicit assumptions for quantitative identities and theorem schemas for global claims, all preflighted through the self-hosted `aeon-logic` parser [21, 25, 26, 27]. Extending these proofs to richer timing/service distributions, unbounded state spaces and full stochastic semantics remains future work.

Queueing theory subsumption scope. Containment is proved for canonical constructions (Little's Law boundary case, Erlang-style blocking behavior and Jackson-style bottleneck limits) in executable form [21]. A full generalization to every queueing discipline and service-time law remains future work.

13. Conclusion

I began with a child handing a ball to another child in a line. Four hundred handoffs. I ended with a topological framework that subsumes queueing theory, predicts biological mutation rates, explains why HTTP/2 has head-of-line blocking and runs on 10-byte UDP frames.

The path between those two points is fork/race/fold: four operations that express any directed acyclic computation graph.

1. **Fork** raises β_1, injects potential energy V – create parallel paths, store work.
2. **Race** traverses homotopy-equivalent paths, converts $V \to K$ – take the fastest.
3. **Fold** projects $\beta_1 \to 0$, extracts work W – merge results deterministically.
4. **Vent** releases excess paths, dissipates heat Q – propagate down, never across.

These operations are not new. DNA replication has used them for 4 billion years. Myelinated neurons pipeline action potentials through them at 100 m/s – a velocity my formula predicts exactly. Photosynthetic antenna complexes fork/race excitation energy at quantum scale with >95% efficiency. A brainless slime mold used them to recreate the Tokyo rail network in 26 hours.

The conveyor belt – Ford's line, TCP's stream, the hospital's referral chain – is the degenerate case. It works when the answer is known, resources are unlimited and a central clock exists. In every other case – which is every real case – the natural topology has $\beta_1 > 0$, and forcing it to zero is where latency hides.

The framework's language was not invented from scratch. It was borrowed – deliberately and entirely – from physical theories that already describe the phenomena I formalize. Quantum physics provided the lexicon: superposition,

tunneling, interference, entanglement, measurement, collapse. These are not metaphors but structural correspondences, validated by photosynthetic systems where the quantum mechanics is literal. The Feynman path integral is not analogous to fork/race/fold – it *is* fork/race/fold: all paths forked, phases raced, amplitudes folded by interference, non-classical paths vented by destructive cancellation (§6.11). The entire tower of classical physics is nested fold operations on the path integral: the Schrödinger equation is the differential form of race, the stationary phase approximation is maximal venting, and Newton's $F = ma$ is the fully folded result at $\beta_1 = 0$ (§6.11). Fluid dynamics provided the scaling intuition: the pipeline Reynolds number predicts phase transitions between sequential and multiplexed scheduling with the same precision that the physical Reynolds number predicts laminar-to-turbulent transitions. Thermodynamics provided the conservation laws: fork injects potential energy, race converts it to kinetic, fold extracts useful work, and venting dissipates waste heat. The First Law ($V = W + Q$) holds exactly – every byte forked is accounted for. The virial theorem gives the exact partition for self-gravitating systems: $W = Q = V/2$ (§6.11). Shannon entropy is the Carnot limit. Frame headers are ground-state energy. The waste heat from vented paths is the cost of certainty – you need the losers to prove the winner is optimal.

The pattern is not merely widespread – it is convergent. Attention in transformers *is* race: the QK^T dot product is the race, softmax is continuous venting, the V projection is fold (§6.10). Training loss *is* waste heat: gradient descent minimizes $\partial Q/\partial\theta$ (§6.10). Protein folding *is* fold: the energy landscape funnel is fork/race/fold from Levinthal's 10^{143} conformations to the native state [24]. The hylomorphism of category theory – unfold/fold – *is* fork/fold. The Carnot cycle's four strokes *are* the four primitives. These are not analogies. They are structural isomorphisms that emerge because any system that transforms input to output under conservation (First Law), irreversibility (Second Law) and minimum overhead (Third Law) will converge to fork/race/fold. There is no alternative.

What this paper claims and what it does not. The scope of the convergence is striking enough to warrant explicit boundaries. This paper claims that fork/race/fold is the *computational structure* of quantum field theory, not an analogy to it. The Feynman path integral is a fork/race/fold computation – mathematically, not metaphorically. The Schrödinger equation is the differential form of the race phase. Newton's laws are the fully folded result at $\beta_1 = 0$. The virial theorem gives the exact energy partition. The weak force is a vent operator. Symmetry breaking is a fold. These are structural identities, not poetic borrowings.

This paper does *not* claim to be a grand unified theory. It does not explain why there are four fundamental forces, or why gravity is weak, or what dark energy is. The gap between "structural pattern that recurs everywhere" and "theory that explains everything" is precisely the gap between topology and physics. Topology describes *shape* – it says that when you compute with gravity,

electromagnetism, the strong force, or the weak force, the computation has the shape of fork/race/fold. It does not say *why* these forces exist or *why* they have the coupling constants they have. It lives on the topology side of that boundary, deliberately.

This paper does not describe the grand unified theory of reality. But it does describe the shape of the box it is going to come in.

The claim is narrower and, I believe, more defensible: whatever the forces *are*, their computational structure is fork/race/fold. Not because fork/race/fold is special, but because it is *inevitable*. Any system that conserves energy, that moves irreversibly forward in time, and that has nonzero ground-state overhead will converge to this shape. Physics satisfies those three constraints. Biology satisfies them. Computation satisfies them. The convergence is bottom-up – independent discovery under shared constraints – not top-down unification from a single principle. The slime mold did not read Feynman. Feynman did not study slime molds. They arrived at the same algorithm because the constraints are the same.

And this convergence has a practical corollary: **to see fork/race/fold is to see a system that has found its shape** (§6.12). The topological deficit $\Delta_\beta = \beta_1^* - \beta_1$ quantifies how far an implementation deviates from its problem's natural topology. Every system in this paper where $\Delta_\beta = 0$ – photosynthesis, DNA replication, saltatory conduction, the path integral – operates at or near theoretical efficiency. Every system where $\Delta_\beta > 0$ – sequential healthcare, T+2 settlement, HTTP/2 over TCP – has measurable, quantifiable waste. The deficit is a diagnostic: find it, close it, and the performance follows.

Put plainly: Bules are where rigor and aesthetics meet. They quantify the gap between a system that merely runs and one that is beautifully shaped for its work.

And fluid dynamics revealed the technique's most counterintuitive property: its worst case is small data, not large. The speedup accelerates with scale, approaching $B \times N$ on large datasets. The hard problem is not surviving at scale – it is avoiding overhead on the simple cases. This is the optimization problem inverted, and it is perhaps the most practically encouraging result: the bigger the workload, the more the algorithm helps.

The children in the line never needed topology. They just needed to stop waiting. Every system that waits when it could fork, that sequences when it could race, that accumulates when it could fold – every such system is leaving performance on the table. Performance that evolution discovered, that slime molds exploit without neurons, that the Feynman path integral computes at every quantum interaction, that 10 bytes of frame header can unlock.

Fork/race/fold is all you need.

References

[1] A. Tero, S. Takagi, T. Saigusa, K. Ito, D. P. Bebber, M. D. Fricker, K. Yumiki, R. Kobayashi, T. Nakagaki, "Rules for Biologically Inspired Adaptive Network Design," *Science*, 327(5964):439–442, 2010.

[2] T. W. Buley, "Aeon Pipelines: A Computation Topology Engine," open-source implementation, 2026. https://github.com/affectively-ai/aeon/tree/main/open-source/aeon-pipelines

[3] D. Akita, I. Kunita, M. D. Fricker, S. Kuroda, K. Sato, T. Nakagaki, "Experimental Models for Murray's Law," *Journal of Vascular Research*, 52(2):75–88, 2016.

[4] T. Nakagaki, H. Yamada, Á. Tóth, "Maze-Solving by an Amoeboid Organism," *Nature*, 407(6803):470, 2000.

[5] J. Lobry, "Asymmetric Substitution Patterns in the Two DNA Strands of Bacteria," *Molecular Biology and Evolution*, 13(5), 1996.

[6] G. S. Engel, T. R. Calhoun, E. L. Read, T.-K. Ahn, T. Mančal, Y.-C. Cheng, R. E. Blankenship, G. R. Fleming, "Evidence for Wavelike Energy Transfer Through Quantum Coherence in Photosynthetic Systems," *Nature*, 446(7137):782–786, 2007.

[7] J. D. C. Little, "A Proof for the Queuing Formula: $L = \lambda W$," *Operations Research*, 9(3):383–387, 1961.

[8] J. R. Jackson, "Jobshop-Like Queueing Systems," *Management Science*, 10(1):131–142, 1963.

[9] Y. Huang et al., "GPipe: Efficient Training of Giant Neural Networks using Pipeline Parallelism," NeurIPS 2019.

[10] D. Narayanan et al., "PipeDream: Generalized Pipeline Parallelism for DNN Training," SOSP 2019.

[11] G. Yu et al., "Orca: A Distributed Serving System for Transformer-Based Generative Models," OSDI 2022.

[12] Y. Leviathan, M. Kalman, Y. Matias, "Fast Inference from Transformers via Speculative Decoding," ICML 2023.

[13] C. Chen et al., "Accelerating Large Language Model Decoding with Speculative Sampling," arXiv:2302.01318, 2023.

[14] M. Belshe, R. Peon, M. Thomson, "Hypertext Transfer Protocol Version 2 (HTTP/2)," RFC 7540, 2015.

[15] J. Iyengar, M. Thomson, "QUIC: A UDP-Based Multiplexed and Secure Transport," RFC 9000, 2021.

[16] M. Isard et al., "Dryad: Distributed Data-Parallel Programs from Sequential Building Blocks," EuroSys 2007.

[17] D. G. Murray et al., "Naiad: A Timely Dataflow System," SOSP 2013.

[18] D. Bray, "Wetware: A Computer in Every Living Cell," Yale University Press, 2009.

[19] A. M. Turing, "The Chemical Basis of Morphogenesis," *Philosophical Transactions of the Royal Society B*, 237(641), 1952.

[20] T. W. Buley, "Aeon Core Runtime (Flow + Compression) and Test Suite," open-source implementation, 2026. https://github.com/affectively-ai/aeon/tree/main/open-source/aeon

[21] T. W. Buley, "Fork/Race/Fold Companion Tests," reproducibility suite, 2026. https://github.com/affectively-ai/aeon/tree/main/docs/ebooks/145-log-rolling-pipelined-prefill/companion-tests

[22] R. P. Feynman, A. R. Hibbs, "Quantum Mechanics and Path Integrals," McGraw-Hill, 1965.

[23] S. Chandrasekhar, "Stellar Structure and Stellar Atmospheres," §IX (The Virial Theorem), *An Introduction to the Study of Stellar Structure*, University of Chicago Press, 1939.

[24] J. N. Bryngelson, J. D. Onuchic, N. D. Socci, P. G. Wolynes, "Funnels, Pathways, and the Energy Landscape of Protein Folding: A Synthesis," *Proteins*, 21(3):167–195, 1995.

[25] L. Lamport, *Specifying Systems: The TLA+ Language and Tools for Hardware and Software Engineers*, Addison-Wesley, 2002.

[26] T. W. Buley, "Aeon Logic: Fork/Race/Fold Temporal Logic Engine and TLC/TLA Compatibility Layer," open-source implementation, 2026. https://github.com/affectively-ai/aeon-logic

[27] Lean FRO Team, "The Lean Theorem Prover (Lean 4)," software and documentation, 2026. https://lean-lang.org

Reproducibility

Source code, test suites and protocol comparison benchmarks are available under open-source license [2, 20, 21, 26]. The scheduler, flow protocol, compression subsystem, computation topology engine and formal parser/tooling layer are independently testable. The validation totals reported in §11 are reproducible from the linked suites.