Fork/Race/Fold Is All You Need

Taylor William Buley Independent Researcher https://buley.fyi/

Abstract

I model **fork/race/fold** as a reusable computational primitive: `fork` work into parallel streams, `race` streams to select earliest valid progress, `fold` results through deterministic reconciliation.

I show closely related structure across natural and engineered systems: *Physarum polycephalum* recreated a rail-like network over nutrient gradients [1], myelinated neurons pipeline action potentials (with measured large speedups), photosynthetic antenna complexes exhibit >95 percent exciton-transfer efficiency in cited systems, and DNA replication uses out-of-order fragment synthesis with deterministic reassembly (Okazaki fragments).

What is new is recognizing these as the *same* algorithm. I present the **Wallington Rotation**, a scheduling algorithm that rotates partially ordered work into concurrent stage-local tracks with controlled reconciliation and prove (constructively plus executable verification) that four primitives – fork, race, fold, vent – are sufficient to express any directed acyclic computation graph used by the implementation classes in this paper. I give the algorithm a natural **topological characterization**: fork increases the first Betti number β_1 (creating independent parallel paths), race traverses homotopy-equivalent paths simultaneously, fold projects β_1 back to zero and vent propagation is a natural transformation that releases paths while preserving structure. Self-describing frames create a **covering space** over the computation graph – working in the cover (multiplexed, out-of-order) then projecting back to the base space (sequential, reassembled).

I then show that **classical queueing theory appears as the $\beta_1 = 0$ boundary case in this framework**. Little's Law, Erlang's formula and Jackson's theorem describe systems constrained to a single topological path. The **pipeline Reynolds number** $Re = N/C$ extends the analysis with topology-sensitive regimes (laminar, transitional, turbulent) in the modeled scope. I also use quantum-mechanical terminology – superposition, tunneling, interference, entanglement, measurement, collapse – as structural correspondence language for mapped computational operations. A thermodynamic accounting lens tracks fork/race/fold/vent as potential, kinetic, useful work and dissipation. I derive the **topological deficit** $\Delta_\beta = \beta_1^* - \beta_1$ as a diagnostic: systems with matched Betti structure ($\Delta_\beta = 0$) tend to show high fit/efficiency in the analyzed set, while systems with $\Delta_\beta > 0$ show measurable waste (healthcare delays, settlement lockup, protocol-level blocking) [21, 29, 30]. I define the **Bule** (1 B = 1 unit of Δ_β) as the corresponding engineering deficit unit.

```
o -> o -> o -> o -> o -> o
```

The conveyor belt was not new when Ford adopted it in 1913. It is a path graph – a line: one-dimensional, simply connected and without branching, where every node has exactly one predecessor and one successor. Modern pipelines can optimize this structure, but in this paper it is treated as a boundary case of a richer topology class.

```
      o -> o
     /     \
o -> o       o -> o
     \     /
      o -> o
```

Fork/race/fold is represented here as a directed acyclic graph (DAG) with merge points: nodes branch, paths run in parallel and merge vertices fold concurrent paths into one. In the analyzed domains, recurring bottlenecks arise when high-β_1 workloads are forced through path-like structures. The operational remedy is to work in the cover space (multiplexed, out-of-order) and project back to the base space (sequential, reassembled).

I instantiate the algorithm in **five** domains, presented in stack order — from building blocks to bytes on wire — each layer enabled by the ones below it.

1) In formal verification (the foundation, §10), I implement a temporal logic model checker (`@affectively/aeon-logic`) whose BFS state-space exploration is itself a fork/race/fold computation. Each multi-successor expansion is a fork, each transition to an already-visited state is a fold (interference), each unfair cycle filtered by weak fairness is a vent, and termination is collapse. The checker verifies a `TemporalModel` of its own exploration and generates a TLA+ specification of the same model, validated through a round-trip-stable TLA sandbox. Both verification paths prove the same invariants: $\beta_1 =$ folded, $\beta_1 \geq 0$, vents $\leq$ folds, and eventual termination under weak fairness. The self-verification is closed under self-application: a formal system built from these primitives can reason about formal systems built from these primitives [26].

2) In formal language theory (the programming model, §11), I implement Gnosis [28] — a programming language whose source code IS the computation graph and whose compiler IS a fork/race/fold pipeline. Programs are Cypher-like graphs with four edge types (FORK, RACE, FOLD, VENT). The compiler statically verifies β_1 bounds, verified by layer 1. The self-hosted compiler (Betti) is itself a GGL program: `(source) -[:FORK]-> (parse_nodes | parse_edges) -[:FOLD]-> (ast)`, establishing closure under construction.

3) In distributed staged computation (the scheduling algorithm, §7), chunked pipelined processing reduces sequential depth from $O(PN)$ to $O(\lceil P/B \rceil + N - 1)$, yielding measured speedups of 3.1x–267x. The Wallington Rotation, expressed in layer 2's language, verified by layer 1's checker.

4) In edge transport (the wire format, §8), I implement a binary stream protocol with 10-byte self-describing frame headers and native fork/race/fold operations on UDP, reducing framing overhead by 95 percent versus HTTP/1.1 and eliminating the topological contradiction that causes head-of-line blocking in HTTP/2 and HTTP/3. Layer 3's scheduling algorithm runs over layer 4's wire format.

5) In compression (bytes on wire, §9), I implement per-chunk topological codec racing (fork codecs, race per chunk, fold to winner), with executable verification of roundtrip correctness, codec-vent behavior and $\beta_1 = \text{codecs} - 1$ invariants [20, 21]. The capstone: actual bytes, actual ratios, actual wire — using every layer below it.

Within the modeled scope in this paper (finite DAG decompositions under C1-C4), the algorithm is an optimal topology class with measurable fit via Δ_β. It is intentionally simple: four primitives, explicit assumptions, and executable checks.

The technique and tooling now exist to identify, measure and reduce topological waste in high-impact domains, including drug discovery, health care and energy systems.

0. A Child, a Ball, a Line

Imagine a child handing a ball to a friend in a line. Four children, one hundred balls. The first child hands Ball 1 to the second, waits while it travels through all four kids, then hands Ball 2. Everyone stands idle while one ball moves. Four hundred handoffs, one at a time.

Now imagine something slightly different. The moment the first child passes Ball 1 to the second, she picks up Ball 2. Now the second child passes Ball 1 to the third while the first child passes Ball 2 to the second. Everyone is busy at once. One hundred balls, four children, one hundred and **three** handoffs. This is **pipelining** – a known technique in computer architecture that has been used since the 1960s.

But what if the first child could juggle? Bundle twenty-five balls together, pass them as a single armful. One hundred balls, four children, chunks of twenty-five: **seven handoffs**. That is a 98 percent reduction. This is **chunked pipelining**, and the formula is:

$$T = \lceil P/B \rceil + (N - 1)$$

where P is the number of balls, B is the chunk size and N is the number of children.

Handoffs are unnecessary overhead: waste to be eliminated. In real life, they take the form of network packets, memory allocations, context switches, and

other system resources that consume time and energy. They are also computationally ugly, as we will prove in section 2.

0.1 The Triangle

Now look at what happens when four chunks move through four children. Draw it as a grid where time moves left to right, children are stacked top to bottom. This is the **grid**:

```
Time:  t1   t2   t3   t4   t5   t6   t7
Kid 1: [C1] [C2] [C3] [C4]
Kid 2:      [C1] [C2] [C3] [C4]
Kid 3:           [C1] [C2] [C3] [C4]
Kid 4:                [C1] [C2] [C3] [C4]
```

Focus on the **ramp-up** – the first four time steps. Read what's active at each moment. This is the **triangle**:

```
t1:  1
t2:  2  1
t3:  3  2  1
t4:  4  3  2  1
```

A triangle. The top has one chunk. Each row adds one more. At t_4 the pipeline is full – every child is busy. Now trace any path through this triangle:

- **Read any column** (one child's work over time): 1, 2, 3, 4. Correct order.
- **Read any row** (all children at one moment): a contiguous subsequence, always in order.
- **Read the diagonal** (t_4, all four children active): 4, 3, 2, 1 – the wavefront. Every chunk is at a different stage, but they are all progressing in the correct relative order.

You cannot break the ordering. The triangle enforces it geometrically. Each child depends only on what the child above passed down (stage dependency) and each chunk depends only on the chunk before it at the same child (sequence dependency). Those two axes – vertical and horizontal – are the only constraints. The triangle is the **tightest possible packing** that satisfies both.

This is not a visualization choice. The triangle *is* the shape of pipelined computation. It is the minimum-area region in time $\times$ stage space that achieves full occupancy while respecting dependency constraints. Any other shape either wastes slots (too wide) or violates ordering (too narrow).

The triangle is also a **covering space** (§3.3). The diagonal – the moment when all children are busy – is the base space: one ordered sequence, 1-2-3-4. But each chunk arrived at the diagonal via a different path through the triangle. Chunk 1 took the longest path (entered first, fell through all four stages). Chunk 4 took the shortest (entered last, only at stage 1). Many paths, one output. That is the covering map. That is why order is always preserved.

And the triangle is **fractal**. Zoom into any sub-triangle and you see the same pattern. If you bundle chunks into mega-chunks, each mega-chunk moves through a larger triangle the same way a single chunk moves through a small one. A polysome with 40 ribosomes on an mRNA looks the same as one ribosome on a short mRNA – same triangle, different scale. *Physarum* (slime mold) tendril networks are the same triangle projected onto geography instead of time $\times$ stage.

The top of the triangle has $\beta_1 = 0$ – one chunk, one path, no parallelism. As you descend, β_1 increases – more chunks in flight, more independent paths through the system. At the diagonal, β_1 is maximum. Then the ramp-down triangle on the other side collapses β_1 back to zero.

Fork is entering the triangle. Race is the diagonal. Fold is exiting.

Now zoom out.

The children are standing inside a classroom. The teacher is managing *three* lines of children, each passing different-colored balls. When one line stalls – for example, a dropped ball – the teacher slides a waiting child from another line into the gap. No one is idle. This is **turbulent multiplexing**: multiple pipelines sharing idle slots across lines.

Zoom out again. The school is one of many in a district. The district coordinator doesn't manage individual children or individual balls. She manages the *shape* of the system – how many lines, how wide, how they interconnect. She has discovered that the number of *independent parallel paths* through the system matters more than the speed of any individual child. She calls this number β_1.

Zoom out once more. You are looking at a strand of DNA inside a cell, and the cell is doing a structurally similar thing. The replication fork is the teacher. Okazaki fragments are the bundled balls. DNA ligase is the child at the end of the line, stitching fragments together without requiring global arrival order. The cell has used this pattern for billions of years.

A simple axiom follows from this zoom-out: *I do not invent efficient coordination; I become aware of it.*

Three natural axioms set the stage.

- **Locality axiom**: if correctness is governed by local constraints, forcing global sequential order adds latency without adding truth.
- **Topology axiom**: when multiple paths preserve correctness, the optimal policy is to fork them, race them, then fold deterministically.
- **Naturalism axiom**: when the same pattern reappears in classrooms, cells and networks, it is not a metaphorical resemblance but a substrate-independent computational shape.

This paper is about the algorithm that the child, the teacher and the cell are all running. It has **three** operations: **fork** work into parallel paths, **race** paths

against each other, **fold** results into a single answer. It has one safety mechanism: **vent** – propagate down, never across. It embraces failure to reduce waste to zero.

These four operations are sufficient to express any directed acyclic computation graph. They have a natural topological characterization in terms of Betti numbers, covering spaces and homotopy equivalence. Classical queueing theory – Little's Law, Erlang's formula and Jackson networks – is the degenerate case where there is only one path. The quantum-mechanical vocabulary – superposition, tunneling, interference, entanglement, collapse – describes the computational operations with structural precision, not metaphor.

This paper began as a practical problem: a sequential bottleneck in a distributed inference pipeline. Tokens moved through layer nodes one at a time, and the obvious optimization – standard pipelining – wasn't good enough. The author, obstinate in personality and trained in decision theory in addition to being a systems engineer, pertinaciously refused to accept the sequential constraint as given. The question was not "how do I make the pipeline faster?" but "why is there a pipeline at all?" That question – a philosopher's question, not an engineer's – led to the optimal topology. And the optimal topology led everywhere. Once rotated, the logs kept rolling on their own.

The conveyor belt was a dominant 20th-century abstraction: serialize work. Fork/race/fold is a correction for workloads with $\beta_1 > 0$, where sequential-only structure imposes avoidable coordination cost.

Three bodies of existing theory provided the language for this correction.

I drew heavily from **quantum physics**, mapping its lexicon entirely onto computational operations: superposition is fork, measurement is observation, collapse is fold, tunneling is early exit, interference is consensus, entanglement is shared state (§5). These are not poetic borrowings – they are structural correspondences, validated by photosynthetic antenna complexes where the quantum mechanics is literal (§1.5).

The quantum-mechanical vocabulary describes the mapped computational operations with structural precision in this paper's scope. It is a modeling lens, not an exclusive language claim.

The second muse is **fluid dynamics**, whose Reynolds number I purloin wholesale into computation as the pipeline Reynolds number $Re = N/C$ (§2.3). Fluid dynamics provides more than vocabulary – it provides the correct intuition for *when* fork/race/fold matters. Just as the Reynolds number predicts when laminar flow becomes turbulent, Re predicts when sequential processing should yield to multiplexed scheduling.

The fluid-dynamical framing reveals an inverted scaling property (§2.2): the worst case is small data, where ramp-up overhead dominates. As data grows, speedup accelerates toward $B \times N$. Scale is the friend, not the enemy.

A thermodynamic framing (§6) reveals the **topological deficit** $\Delta_\beta = \beta_1^* - \beta_1$: the gap between a problem's intrinsic parallel topology and the implementation's actual topology. An information-theoretic framing (§6.8) unifies both: fork creates $\log N$ bits of uncertainty, fold compresses to a single outcome, and vented paths carry away exactly the bits that cannot be recovered.

1. Nature Got There First

Fork/race/fold is not a metaphor for natural systems. It is the same algorithm, discovered accidentally, running on different substrates. I grade each mapping:

- **Grade A**: Quantitative isomorphism – the algorithm's math directly models the system with embedded predictive power.
- **Grade B**: Structural homology – deep structural match, genuine design insight, no novel quantitative prediction.

In discovering the true shape of beautiful computation I do nothing novel except put a finger on the nose of the universe's own convergence.

In this paper, computational beauty is operational: low topological deficit, explicit invariants, and reproducible gains under stated assumptions.

1.1 *Physarum polycephalum*: The Phineas Gage of Distributed Intelligence (Grade A)

In 1848, a railroad construction foreman named Phineas Gage survived a 43-inch iron rod through his left frontal lobe. He could walk, talk, and reason – but his personality changed utterly. The accident revealed that personality *lives somewhere specific* in the brain. It was the founding observation of neuropsychology: a single dramatic injury that reorganized an entire field's understanding of where intelligence resides.

His skull and scientific divining rod are now in a display case in a medical library at Harvard.

As Phineas Gage was to neuroscience, so slime mold is to distributed intelligence.

In 2010, Tero et al. placed oat flakes on a wet surface in positions corresponding to the 36 stations of the greater Tokyo rail network [1]. They introduced a single *Physarum polycephalum* slime mold at the position corresponding to Tokyo station. The organism – which has no brain, no neurons, no central nervous system of any kind – extended exploratory tendrils in all directions (**fork**). Multiple tendrils reached each food source via different routes (**race**). The organism then pruned inefficient connections, reinforcing high-flow tubes and abandoning low-flow ones (**fold** with **venting** of abandoned paths).

Within 26 hours, the slime mold had independently constructed a transport network that closely matched the actual Tokyo rail system – a network that professional engineers had spent decades and billions of dollars optimizing.

The quantitative correspondence is striking:

- **Total network length**: *Physarum* network within 5 percent of the engineered rail network
- **Fault tolerance**: *Physarum* networks showed comparable resilience to random link removal
- **Transport efficiency**: Cost-performance tradeoff matched the Pareto frontier of the engineered system
- **Topology**: The network's β_1 (number of independent cycles) matched within one cycle

The author does not know of a slime mold yet on display at a public library, but remains hopeful.

The mapping to fork/race/fold is not analogy – it is mechanism:

Physarum Behavior	Fork/Race/Fold Operation
Exploratory tendril extension	**Fork**: create N parallel paths from current position
Cytoplasmic streaming through tubes	**Race**: flow rate determines winner
Tube reinforcement (positive feedback)	**Fold**: high-flow paths become canonical
Tube abandonment (starvation)	**Vent**: low-flow paths released, descendants shed
Shuttle streaming (oscillatory flow)	**Self-describing frames**: bidirectional flow carries positional information

Just as Phineas Gage's injury revealed that intelligence has a specific anatomical locus, *Physarum*'s rail network shows that optimization can emerge without centralized cognition. Fork/race/fold does not require a neural substrate; it needs parallel paths, a selection signal and a way to prune. In this manuscript, that structure is instantiated in protoplasm, silicon and 10-byte frame transport.

Predictive power: The Wallington Rotation's optimal chunk size formula predicts that *Physarum* tube diameter should scale with the cube root of flow distance (balancing viscous resistance against metabolic maintenance cost). This matches Akita et al.'s measurements of tube morphology [3]. The pipeline Reynolds number framework predicts that *Physarum* networks under nutrient stress should consolidate to fewer, thicker tubes (reducing β_1, increasing per-tube Re) – exactly the behavior observed by Nakagaki et al. in nutrient-limited environments [4].

1.2 DNA Replication: The Original Self-Describing Frame Protocol (Grade A)

DNA's two strands run antiparallel. The leading strand synthesizes continuously (clean pipeline). The lagging strand produces **Okazaki fragments** – 1,000–2,000 nucleotide chunks in prokaryotes, 100–200 in eukaryotes – synthesized out of order and stitched together by DNA ligase.

Each Okazaki fragment is a **self-describing frame**: its genomic coordinate is its `stream_id` + `sequence`. DNA ligase is the **frame reassembler** – it joins fragments without requiring global ordering. The replication fork moves at ~1,000 nt/s in *E. coli*. At any moment, 1–3 fragments are being synthesized simultaneously, giving a pipeline Reynolds number $Re \approx 0.7$–1.0.

Predictive power: My chunked pipeline formula $T = \lceil P/B \rceil + (N-1)$ predicts that prokaryotic fragments (~1,000 nt) should be longer than eukaryotic fragments (~150 nt) because eukaryotes have more processing stages N (chromatin reassembly, histone deposition). This matches observation. The framework also predicts that organisms with lower Re (more exposed single-stranded DNA during lagging strand synthesis) should have higher lagging-strand mutation rates. This asymmetry has been observed experimentally [5].

1.3 Saltatory Nerve Conduction: The Formula Matches Exactly (Grade A)

In myelinated neurons, action potentials jump between nodes of Ranvier (~1–2 mm apart) instead of propagating continuously. Multiple action potentials are in-flight simultaneously across different internodal segments.

Perhaps you picture biological denial of service-style "packet overload"? The biology is actually designed to prevent exactly those issues. After a node fires, it enters a "refractory period" where it physically cannot fire again for a few milliseconds. This acts like a one-way valve, a buffer ensuring that even though multiple "packets" (action potentials) are in flight, they can never move backward or merge into one another.

This is chunked pipelining. The "chunking" allows the brain to receive a high-frequency stream of data rather than a single pulse. It allows for more nuanced signaling—the frequency of the spikes (the "bitrate") conveys the intensity of the stimulus.

By only depolarizing the membrane at the nodes of Ranvier, the neuron saves a massive amount of metabolic energy (ATP) that would otherwise be spent pumping ions back and forth across the entire length of the axon. The brain doesn't just pipeline; it optimizes for energy efficiency by reducing the number of times ions need to be pumped back across the membrane.

The Wallington formula predicts conduction velocity:

$$v = \frac{B}{t_{\text{stage}}} = \frac{1.5 \text{ mm}}{0.015 \text{ ms}} = 100 \text{ m/s}$$

where $B = 1.5$ mm is the median internodal distance in large-diameter mammalian motor neurons (range: 0.2–2.0 mm across fiber types and species; Waxman & Swadlow, 1977) and $t_{\text{stage}} = 0.015$ ms is the nodal delay (the time for voltage-gated sodium channels to depolarize the membrane at each node of Ranvier; Tasaki, 1939). The formula predicts a velocity *range* across species: $B = 0.2$ mm gives $v = 13$ m/s (thin sensory fibers), $B = 2.0$ mm gives $v = 133$ m/s (large motor fibers). The measured range (10–120 m/s) matches.

Measured conduction velocity: 100 m/s (for large myelinated fibers). Unmyelinated conduction (without pipelining): 0.5–2 m/s. The 50–100x speedup is real, measured, and predicted by the same formula that predicts my pipeline speedups. Deviations from the optimal internode distance (as in multiple sclerosis demyelination) produce the predicted conduction velocity changes.

Myelin is the biological argument for investing in transport-layer reliability to enable larger chunks – skip intermediate processing, insulate the wire. This is the case for UDP over TCP: invest in framing reliability so you can skip ordered delivery.

1.4 Polysome Translation: The Wallington Rotation in Biology (Grade A)

A so-called polysome consists of multiple ribosomes simultaneously translating the same mRNA, spaced ~30–40 codons apart. This *is* the Wallington Rotation: the mRNA is the pipeline, each ribosome processes a chunk, and multiple proteins emerge concurrently.

Without pipelining: 40 proteins from one mRNA = 2,400 s. With polysome: ~118 s. **20x speedup.**

When Re drops below ~0.6, the mRNA is targeted for degradation (no-go decay). The cell destroys underutilized pipelines and reallocates ribosomes – exactly what turbulent multiplexing prescribes. Under stress, cells globally reduce Re but maintain high Re on priority mRNAs via IRES elements. The strategy of giving priority streams idle slots evolved 2 billion years ago.

1.5 Photosynthetic Light-Harvesting: Fork/Race at Quantum Scale (Grade A)

Photosynthetic light-harvesting is widely considered a strong example of quantum-coherent energy transfer in biology. In these systems, the algorithm in action is environment-assisted quantum transport, where excitons exploit spatial superposition to sample multiple pathways and achieve high reported transfer efficiency before decoherence.

Antenna complexes in photosynthesis contain ~200–300 chlorophyll molecules. Photon excitation energy forks across the pigment network, races through multiple pathways and the first path to reach the reaction center wins. Charge separation is fold. Non-photochemical quenching is venting. *Exciton transfer efficiency* (the probability that an absorbed photon's energy reaches the reaction center before decoherence) exceeds 95 percent in purple bacteria and approaches 99 percent in green sulfur bacteria. This is distinct from whole-plant photosynthetic quantum yield (3–6 percent), which includes downstream losses in carbon fixation, photorespiration, and reflection. The fork/race/fold characterization applies to the exciton transfer step specifically.

Fleming et al. (2007) showed that excitation energy exists in **quantum superposition** across multiple pigments simultaneously [6]. The fork/race/fold framework predicts that transfer efficiency should increase with pigment count (more forked paths = higher probability of reaching the reaction center before decoherence) but with diminishing returns — each additional pigment adds one path ($\beta_1 \mathrel{+}= 1$) but the marginal probability gain decreases as the existing network already covers most of the energy landscape. This logarithmic-like scaling matches observation: ~200 pigments achieve >95 percent efficiency, while further increases provide negligible improvement. The quantum vocabulary I use in §5 is not metaphor; it is structural correspondence with the physics.

1.6 Immune System V(D)J Recombination (Grade B)

The adaptive immune system generates 10^{11} unique antibody configurations through combinatorial recombination (**fork**), exposes them to antigen simultaneously (**race**) and expands the winners through clonal selection (**fold**). Non-binding clones are eliminated (**vent**). Self-reactive B cells undergo clonal deletion – the lineage is eliminated, but sibling B cells with different recombinations are unaffected. The parallelism factor is 10^{11} – the most massively parallel fork/race on Earth.

This is not just parallelism; it is **probabilistic parallelism**. The immune system does not know which configuration will bind the antigen. It forks a vast library, races them against the antigen and folds the winners. The algorithm is identical to the fork/race/fold pattern in distributed systems.

1.7 Transformers: Accidental Rediscovery (Grade A)

The biological examples above are evolutionary discoveries. But the pattern extends to human-engineered systems that arrived at the same structure without recognizing it. The transformer architecture (§6.11) is the most prominent: multi-head attention is fork/race/fold (N heads fork, compute in parallel, concatenate-and-project to fold), feed-forward layers are fork/vent/fold (expand 4x, activate/suppress, contract), residual connections are two-path fork with additive fold, and softmax is continuous venting. The entire architecture is a nested fork/race/fold graph – discovered empirically by Vaswani et al. (2017)

through gradient descent, not through topological analysis. That the dominant architecture in machine learning independently converged on the same four primitives as *Physarum*, DNA polymerase and photosynthetic antenna complexes is the strongest evidence that fork/race/fold is discovered, not invented.

Falsifiable prediction. If the fork/race/fold characterization is structural rather than coincidental, then removing parallel-path properties while preserving parameter count should degrade performance *superlinearly*, not linearly. Specifically: replacing multi-head attention (N heads, fork width N) with single-head attention of the same total dimension (d_{model}) should degrade quality by more than $1/N$, because the loss is topological (reduced β_1), not merely statistical (fewer parameters). Preliminary evidence supports this: Voita et al. (2019) showed that pruning attention heads to a single head degrades translation quality disproportionately on syntactically complex sentences — precisely the inputs whose dependency structure has high intrinsic β_1^*. Similarly, replacing the 4x FFN expansion-contraction (fork/vent/fold) with a same-parameter linear layer should degrade performance, because the vent operation (activation sparsity) performs topological path selection that a linear map cannot. The prediction is testable: measure Δ_β of the ablated architecture and correlate with the performance drop.

1.8 The Convergence

Seven systems across seven orders of magnitude in scale – from quantum-coherent pigment networks to billion-parameter neural networks. Different substrates. Different evolutionary histories. A shared structural motif. The evidence in this paper is consistent with convergence toward fork/race/fold under three constraints:

1. **Finite resources, high demand** → chunked pipelining and multiplexing
2. **Unknown correct answer** → fork/race/fold with vent
3. **No global clock** → self-describing frames with out-of-order reassembly

These three constraints make fork/race/fold optimal within the class of finite DAG topologies examined in this paper, not merely possible. Systems lacking all three can still use fork/race/fold (transformers have synchronized SGD; photosynthesis has electromagnetic field synchronization) – but when all three constraints bind simultaneously, no alternative topology in this paper's scope achieves the same efficiency. Other concurrent models (gossip protocols, epidemic algorithms, eventually-consistent CRDTs) also operate under these constraints but make different tradeoffs — gossip sacrifices deterministic fold for probabilistic convergence, CRDTs sacrifice fold entirely for commutativity. The claim is not that fork/race/fold is the *unique* optimum, but that it is the topology to which independent systems *converge* when they require both parallelism and deterministic reconciliation. The conveyor belt is the degenerate case: it works only when the answer is known, resources are unlimited and a central clock exists. It is an unshapely idea with ugly consequences.

2. The Algorithm

2.1 Pipeline Model

A pipeline with N stages processes workload of P items.

Serialized: $T_{\text{serial}} = P \cdot N$ – each item completes all stages before the next begins.

Pipelined: $T_{\text{pipeline}} = P + (N - 1)$ – stage-local ordering suffices.

Chunked: $T_{\text{chunked}} = \lceil P/B \rceil + (N-1)$ – intra-stage parallelism (SIMD, batched ops) with chunk size B.

For $C = \lceil P/B \rceil$ chunks across N stages, the idle fraction during ramp-up/ramp-down:

$$\text{idle} = \frac{N(N-1)}{2(C+N-1)}$$

2.2 The Inverted Scaling Property

The speedup of chunked pipelining over serialized processing is:

$$\text{Speedup} = \frac{P \cdot N}{\lceil P/B \rceil + (N-1)}$$

For small P (few items, few stages), the denominator's $(N-1)$ term – the ramp-up cost – dominates. The pipeline spends most of its time filling and draining, never reaching full occupancy. The idle fraction $N(N-1)/2(C+N-1)$ is large. This is the **worst case**.

For large P, the $\lceil P/B \rceil$ term dominates and $(N-1)$ becomes negligible:

$$\text{Speedup} \xrightarrow{P \to \infty} \frac{P \cdot N}{P/B} = B \cdot N$$

The speedup approaches $B \times N$ – the product of chunk size and stage count. The pipeline is fully occupied. The kids all have balls. Idle fraction approaches zero. The kids are all juggling.

The technique gets *faster and faster* as the work at hand grows.

This is a profoundly inverted scaling property, which is a useful and unusual feature of the algorithm. In most engineering contexts, the hard problem is scale – systems that work beautifully on small inputs collapse under large ones. Here, the opposite is true: large datasets are where pipelining shines, approaching its theoretical maximum speedup of $B \times N$. Small datasets are where the overhead hurts.

The optimization challenge in fork/race/fold-based pipelines is not "how do I survive at scale?" but "how do I avoid overpaying on small workloads?" – a far more pleasant problem. Trivial solutions for such trivial concerns are panoply – early exit, dynamic chunk sizing and adaptive scheduling, for example – but speedups are more precious when they arrive in abundance.

The fluid-dynamical analogy (§2.3) captures this precisely. Low Re (many chunks, few stages – large data) corresponds to laminar flow: smooth, predictable, high utilization. High Re (few chunks, many stages – small data) corresponds to turbulent flow: idle slots appear, multiplexing becomes necessary, overhead rises. The Reynolds number predicts exactly where the crossover occurs, and the laminar regime – the easy case – is the one that grows with data size.

2.3 The Pipeline Reynolds Number

I define:

$$Re = N/C$$

This is the ratio of stages to chunks – the density of the pipeline. Low Re (< 0.3): laminar regime, steady-state, high utilization. Transitional Re (0.3–0.7): idle-slot recovery is profitable. High Re (> 0.7): turbulent regime, multiplexing across requests yields the largest benefit.

The Reynolds number is not metaphor. In fluid dynamics, $Re = \rho v D/\mu$ predicts the transition from laminar to turbulent flow: inertial forces (numerator) versus viscous forces (denominator). In computation, the structural correspondence is: stages N are the inertial force (more stages = more momentum in the pipeline, more work in flight), and chunks C are the viscous force (larger chunks = more resistance to context switching, more "friction" per scheduling decision). Low Re (large chunks, few stages) is laminar: each chunk flows smoothly through a short pipeline. High Re (small chunks, many stages) is turbulent: many items in flight, multiplexing across requests yields the largest benefit. The transition happens because at high Re, the ramp-up/ramp-down triangles (§0.1) waste more idle slots than the multiplexing overhead costs to recover — exactly as turbulence in fluids becomes energetically favorable when inertial forces dominate viscous damping.

2.4 Four Primitives

Given pipeline state S and operation set O:

1. **Fork**: $\text{Fork}(S, O) \to \{S_1, \ldots, S_k\}$ – create k independent branch states, each processing a subset of O. Topological effect: $\beta_1 \mathrel{+}= k - 1$.

2. **Race**: $\text{Race}(\{S_i\}) \to (S_w, i_w)$ – advance all branches concurrently; select the first to reach a valid completion. Losers are vented. Exploits homotopy equivalence.
3. **Fold**: $\text{Fold}(\{S_i\}, f) \to S^*$ – wait for all branches to complete (or vent); apply deterministic merger f to produce a single canonical state. Topological effect: $\beta_1 \to 0$.
4. **Vent**: $\text{Vent}(S_i) \to \perp$ – cease output, recursively vent all descendants, leave siblings untouched. **One rule: propagate down, never across.** The system releases excess energy from paths that cannot contribute useful work – a pressure relief valve that prevents computational overheating.

Completeness (finite, mechanized scope). These four primitives are sufficient to express finite directed acyclic computation graphs under explicit decomposition assumptions. Any finite DAG can be decomposed into fork points (nodes with out-degree > 1), join points (nodes with in-degree > 1) and linear chains. Fork creates divergences. Fold creates convergences. Race is fold with early termination. Vent handles failures and excess energy. Linear chains are the trivial case (no fork, no fold). In the formal stack, local decomposition is constructive and the global statement is an explicit-assumption theorem schema, paired with executable finite-DAG decomposition checks [21, 26].

2.5 Correctness Conditions

Fork/race/fold preserves correctness when:

- **C1 (Constraint locality)**: Stage-local ordering is sufficient for global correctness.
- **C2 (Branch isolation)**: A vented branch does not corrupt siblings.
- **C3 (Deterministic fold)**: The merger f is deterministic.
- **C4 (Termination)**: Every branch either completes, is vented, or times out in finite time.

These conditions are mechanized in a two-layer formal stack. Finite-state models in TLA+ verify C1–C4 as invariants across the formal module set (checked by both TLC and the self-hosted `aeon-logic` parser/checker). Lean 4 theorem schemas verify the quantitative identities that depend on C1–C4 under explicit assumptions. The sufficiency claim — that any finite DAG decomposes into fork points, join points, and linear chains, and that these four conditions preserve correctness through the decomposition — is verified constructively by executable finite-DAG decomposition checks [21, 26].

2.6 Five Fold Strategies

Not all folds are equal. The choice of merger f determines the computational semantics:

Strategy	Semantics	Time Complexity	When
Winner-take-all	Best result by selector	$O(N)$ comparisons	One answer needed, clear criterion
Quorum	K of N must agree	$O(N^2)$ pairwise comparisons	Byzantine fault tolerance
Merge-all	All results contribute	$O(N) + O(f)$ where $f =$ merger cost	Complementary information
Consensus	Constructive/destructive interference	$O(N^2)$ pairwise comparisons	Signal amplification or outlier detection
Weighted	Authority-weighted merger	$O(N) + O(f)$ where $f =$ merger cost	Heterogeneous source quality

Race is not a fold strategy – it is a separate primitive. Race picks the *fastest* result. Winner-take-all picks the *best* result. The distinction matters: race terminates early (venting losers), winner-take-all waits for all branches to complete.

Implicit in this is the fact that failure is a necessary component of any robust system. Failure modes are handled by the vent primitive, which propagates down the tree but never across branches. This ensures that a failure in one branch does not cascade to other branches, maintaining the isolation property required for correctness. A system that cannot fail gracefully is not robust.

2.7 Vent Propagation

Venting is the protocol-level analogue of NaN propagation in IEEE 754, `AbortSignal` in web APIs and apoptosis in biology. The one rule – **propagate down, never across** – makes composition safety a architectural feature rather than an accidental one. Any pipeline of fork/race/fold stages is safe by construction because venting never crosses branch boundaries. The system sheds excess pressure the moment a path becomes unproductive.

2.8 The Worthington Whip

The Worthington Whip extends fold for aggressive parallel shard merging. A single workload of P items is sharded across S parallel pipelines, each processing P/S items. At fold, a cross-shard correction reconciles the results.

Derivation of the $(S-1)/2S$ reduction. In a computation with pairwise dependencies, an unsharded system processes all $\binom{P}{2} = P(P-1)/2$ pairs. After sharding into S partitions of P/S items each, each shard processes only its intra-shard pairs: $\binom{P/S}{2} = (P/S)(P/S-1)/2$. The total intra-shard work across all S shards is $S \cdot (P/S)(P/S-1)/2 = P(P/S-1)/2$. The cross-shard pairs — the ones not processed within any shard — number $\binom{P}{2} - S \cdot \binom{P/S}{2}$. As $P \to \infty$, the ratio of intra-shard work to total work approaches $1/S$, so the per-shard compute reduction is $1 - 1/S = (S-1)/S$. Per shard, each shard avoids $(S-1)/S$ of the total pairs, but since each shard processes $1/S$ of the total, the per-shard savings relative to processing the full P is $(S-1)/2S$. The cross-shard correction at fold time reconciles the missing pairs — this is the whip snap.

The fold phase is the whip snap: all parallel shards converge to a single definite state. The computational snap is single-threaded execution with the speed of parallel processing.

Beyond pairwise: k-tuple dependencies. The $(S-1)/2S$ derivation assumes pairwise interactions. For k-tuple dependencies (e.g., 3-way consistency checks), the intra-shard work fraction generalizes to $S \cdot \binom{P/S}{k}/\binom{P}{k}$, which approaches $1/S^{k-1}$ as $P \to \infty$. The per-shard savings grow with k: for $k = 3$, the reduction is $(S^2-1)/S^2$ — sharding is *more* beneficial for higher-order dependencies, not less. The cross-shard correction at fold reconciles $\binom{P}{k} - S \cdot \binom{P/S}{k}$ missing k-tuples. For unequal partition sizes, the correction cost depends on the size distribution; the $(S-1)/2S$ formula is the symmetric optimum, and any partition skew strictly increases the cross-shard fraction.

The crossover point. Adding shards reduces per-shard computation by $(S-1)/2S$ but increases cross-shard correction cost. The correction is a fold over S partial results — itself an $O(S)$ operation for merge-all or $O(S^2)$ for quorum. The crossover occurs when the marginal correction cost of shard $S+1$ exceeds the marginal per-shard savings: $\partial(\text{correction})/\partial S > \partial(\text{savings})/\partial S$. For pairwise dependencies with $O(S)$ merge-all fold, this yields an optimal $S^* = \lceil\sqrt{P}\rceil$. The TLA+ `WhipCrossover` model mechanizes this: it explores all (S, P) configurations up to finite bounds and verifies that the crossover point satisfies the inequality [21, 26].

These novel whipper snappers are the ultimate expression of parallelism.

3. The Topology of Fork/Race/Fold

3.1 Betti Numbers Classify Computation Graphs

The first Betti number β_1 counts independent parallel paths in a topological space. It is the number that matters:

Structure	β_1	Parallelism	Fault Tolerance
Sequential pipeline	0	None	None
Single fork/join	1	One level	One failure
Fork with K paths	$K-1$	K-way	$K-1$ failures
Full mesh of N nodes	$\binom{N}{2}$	Maximum	Maximum

Fork/race/fold is the operation that **temporarily raises β_1 to exploit parallelism, then lowers it back to zero**:

Primitive	Topological Operation	Effect on β_1
Fork	Create parallel paths	$\beta_1 \mathrel{+}= N-1$
Race	Traverse homotopy-equivalent paths	β_1 stays high
Fold	Merge all paths to single output	$\beta_1 \to 0$
Vent	Release a path	$\beta_1 \mathrel{-}= 1$

The entire history of process optimization – Ford's assembly line, TCP's ordered byte stream, the hospital referral chain, T+2 financial settlement – is the story of forcing $\beta_1 = 0$ onto problems whose natural topology has $\beta_1 \gg 0$. Healthcare diagnosis has intrinsic $\beta_1 \geq 3$ (blood work, imaging, genetic screening, and specialist consultation are independent). The referral system forces $\beta_1 = 0$. The mismatch is the 4.7-year average diagnostic delay for rare diseases [29]. Financial settlement has intrinsic $\beta_1 = 2$. T+2 forces $\beta_1 = 0$. At a DTCC-core daily baseline, that yields ~\$4.4T/day of locked capital; in the paper's broad-scope scenario, the same topology yields ~\$70T/day [21, 30].

3.2 Homotopy Equivalence

Two computations are homotopy equivalent if they produce the same result through different topological paths. In a sequential pipeline, there is exactly one path – no homotopy is possible. In a fork/race graph with N paths, if the computation is deterministic, all N paths are homotopy equivalent.

Race exploits homotopy equivalence: race discovers that all paths lead to the same answer and takes the fastest. **Fold handles the general case**: when paths are *not* homotopy equivalent (a blood test and an MRI give different information), the merger function f combines non-equivalent results into a richer output than any single path could provide.

The distinction is topological: race requires homotopy equivalence (π_1-trivial computation on each path). Fold does not. This is why they are separate primitives.

3.3 Covering Spaces and Self-Describing Frames

A covering space maps onto a base space such that every point has a neighborhood that is evenly covered. Self-describing frames create a covering space over the computation graph. Each frame carries `(stream_id, sequence)` – its coordinates in the covering space. The base space is the sequential computation. The covering space is the multiplexed computation.

The **frame reassembler is the covering map**: it projects the cover back to the base space. Frames arrive from any point in the cover (any stream, any sequence) and are reassembled into sequential order.

TCP forces you to work in the base space – one ordered stream. Simply connected. **UDP with self-describing frames lets you work in the covering space** – many streams, local ordering, out-of-order reassembly. The topological degree of the covering map is the multiplexing factor.

This is precisely what DNA ligase does: Okazaki fragments arrive from the covering space (out-of-order lagging-strand synthesis) and are projected back to the base space (the complete double-stranded genome). DNA ligase is the covering map. It has been performing this topological operation for 4 billion years.

3.4 The Fundamental Group and Protocol Design

The fundamental group π_1 classifies loops up to homotopy:

- **TCP**: $\pi_1 = 0$. One path. Simply connected. Works for simply connected problems.
- **HTTP/2**: Application layer has $\beta_1 > 0$ (multiplexed streams), but TCP substrate has $\beta_1 = 0$ (one ordered byte stream). **This is a topological contradiction.** Head-of-line blocking is the symptom: losing one packet on any stream blocks *all* streams because the underlying space cannot support independent paths.
- **HTTP/3 (QUIC)**: Partially resolves the contradiction with per-stream independence on UDP. But maintains ordered delivery within each stream – π_1 within each stream is trivial.
- **Aeon Flow over UDP**: Self-describing frames in the covering space. No ordered delivery anywhere. π_1 of the wire matches π_1 of the application. **No topological contradictions. No head-of-line blocking at any layer.** The shape of the protocol matches the shape of the problem.

3.5 Time-Indexed Topological Filtration

The evolution of β_1 over a computation's lifetime forms a *filtration* — a nested sequence of topological spaces indexed by time:

- $t = 0$: Computation starts. $\beta_1 = 0$.
- $t = t_{\text{fork}}$: β_1 jumps to $N - 1$.

- During race: β_1 stays at $N-1$.
- $t = t_{\text{vent}_i}$: β_1 drops by 1 per vented path.
- $t = t_{\text{fold}}$: $\beta_1 \to 0$.

Terminology note. This time-indexed β_1 evolution is a *filtration* in the algebraic-topological sense — a nested sequence of subcomplexes $K_0 \subseteq K_1 \subseteq \cdots \subseteq K_n$ where each K_t is the computation graph at time t. It borrows the birth/death language of persistent homology (Edelsbrunner et al., 2002) but the filtration parameter is *time*, not *distance scale* as in classical topological data analysis of point clouds. The formal properties of TDA persistence (stability under perturbation, isometry invariance) apply only when the filtration is metric-indexed; the time-indexed version retains the birth/death structure but not the stability guarantees.

The filtration diagram encodes: how much parallelism was used (features born at fork), how quickly bad paths were pruned (short persistence = speculation), how much redundancy survived to fold (long persistence = consensus). A well-optimized system has short vent persistence (release early) and long fold persistence (exploit parallelism fully).

3.6 Category-Theoretic Framing

In category theory, a so-called monoidal category is a mathematical system consisting of a collection of objects and morphisms, or a way to combine objects in a way similar to multiplication.

Fork/race/fold forms a **monoidal category**:

- **Objects**: computation states (sets of active streams).
- **Morphisms**: Fork ($S \to S_1 \otimes S_2 \otimes \cdots \otimes S_n$), Race ($\bigotimes S_i \to S_{\text{winner}}$), Fold ($\bigotimes S_i \to f(S_1, \ldots, S_n)$).
- **Tensor product** $\otimes$: parallel composition.
- **Composition** $\circ$: sequential composition.

The conveyor belt uses only composition. Fork/race/fold uses both composition and tensor product. In this sketch, that suggests a broader expressive surface. Vent propagation is modeled as a **natural transformation** from active computations to terminated computations – preserving morphism structure across the tensor product, i.e., "propagate down, never across."

What this buys. The monoidal framing is not fully developed here — a rigorous treatment would require specifying the unit object (the empty computation), proving associativity and unit laws for both $\otimes$ and $\circ$, and verifying the interchange law. The claim is sketched, not proved, and the paper's results do not depend on it. The value is taxonomic: it connects fork/race/fold to the category-theoretic literature on dataflow (hylomorphisms = unfold/fold, which map to fork/fold) and allows future work to import results from monoidal category theory — in particular, the coherence theorem (Mac Lane, 1963) would

guarantee that different compositions of fork/race/fold reach the same result regardless of bracketing, which is a stronger form of C3 (deterministic fold).

4. Containing Queueing Theory

4.1 Little's Law as a Special Case

Little's Law states: $L = \lambda W$, where L is the average number of items in a system, λ is the arrival rate and W is the average time in the system. This is the foundational result of queueing theory, proved by Little in 1961 [7] and considered universal within its domain.

Containment theorem (operational form). Under assumptions C1-C4 and standard Markovian service models, the fork/race/fold framework recovers canonical queueing results when constrained to $\beta_1 = 0$, and strictly extends them when $\beta_1 > 0$ by adding topology as a control variable. The executable proofs in §13 include direct tests for Little, Erlang-style blocking behavior and Jackson-style bottleneck limits [21].

Precision on "containment." Little's Law ($L = \lambda W$) holds under remarkably weak assumptions — it requires only ergodicity and finite expectations, not Markovian arrivals or any specific service-time distribution (Little & Graves, 2008). The C1-C4 conditions are *not* equivalent to Little's Law's assumptions; they are *stronger* (C3 demands deterministic fold, C4 demands finite termination). Containment means: every system satisfying C1-C4 with $\beta_1 = 0$ also satisfies the conditions under which Little's Law holds, and the framework produces the same steady-state predictions. The converse does not hold — Little's Law applies to systems that violate C3 (non-deterministic service) or have no fold semantics at all. What fork/race/fold adds is not a relaxation of Little's assumptions but an extension of the *vocabulary*: when $\beta_1 > 0$, topology becomes a control variable that queueing theory has no notation for.

Little's Law describes a system with one path – items enter, wait, get served, exit. There is no concept of forking, no concept of racing, no concept of folding parallel results. When $\beta_1 > 0$, Little's Law still holds *per path*, but says nothing about the *topology* of the system – how paths interact, when to fork, when to fold, when to vent.

The pipeline Reynolds number $Re = N/C$ generalizes Little's Law:

Queueing Theory	Fork/Race/Fold
$L = \lambda W$ (items in system)	$\beta_1 = N - 1$ (parallel paths in system)
Utilization $\rho = \lambda/\mu$	$Re = N/C$ (stages / chunks)
$\rho < 1$ for stability	$Re < 0.3$ for laminar; $Re > 0.7$ for turbulent
M/M/1, M/M/c, M/G/1 variants	Laminar, transitional, turbulent regimes
Arrival rate λ	Fork rate

Queueing Theory	Fork/Race/Fold
Service rate μ	Fold rate
Queue discipline (FIFO, priority)	Fold strategy (quorum, weighted, consensus)

In queueing theory, an M/M/1 queue represents the simplest non-trivial model of a waiting line. It describes a memoryless system with a single server where arrivals and service times are essentially random. Its notation follows Kendall's Notation, where each letter defines a specific characteristic of the system:

- M (Markovian/Memoryless) Arrival: Customers arrive according to a Poisson process. This means the time between arrivals follows an Exponential distribution. It is "memoryless" because the time until the next arrival doesn't depend on how much time has already passed.
- M (Markovian/Memoryless) Service: The time it takes to serve a customer also follows an Exponential distribution.
- 1 (Single Server): There is only one station or person processing the queue.

Every M/M/1 queue is a pipeline with $\beta_1 = 0$, one stage and Poisson arrivals. The *Re* framework does not contradict queueing theory – it *contains* it. When $\beta_1 = 0$, *Re* reduces to utilization. When $\beta_1 > 0$, *Re* predicts phenomena that queueing theory has no vocabulary for: the transition from sequential to multiplexed scheduling, the optimal fork width and the cost of topological mismatch.

4.2 Erlang's Formula as Fold Without Fork

Erlang's B formula gives the blocking probability for c servers with no queue:

$$B(c, A) = \frac{A^c/c!}{\sum_{k=0}^{c} A^k/k!}$$

In fork/race/fold terms, Erlang's system is a race over c servers – but without fork. Arrivals are not forked; they are routed to a single server. The system cannot exploit parallelism because it has no fork operation. Blocking occurs when all c paths are occupied – but there is no mechanism to create *new* paths on demand.

While Agner Krarup Erlang provided the mathematical logic that allows us to build networks that don't collapse under pressure, he didn't have fork/race/fold.

Fork/race/fold eliminates blocking by making path creation dynamic. When demand exceeds capacity, fork creates new paths (β_1 increases). When demand subsides, fold and venting remove paths (β_1 decreases). The topology adapts to the load. Erlang's formula describes the *static* case; fork/race/fold describes the *dynamic* case.

4.3 Jackson Networks as Fixed-Topology Pipelines

James R. Jackson was a mathematician at UCLA who, by 1963, realized that, in the real world, queues don't exist in isolation: a factory floor, a hospital, or a data center are all complex networks, not simple conveyer belts.

Jackson's theorem [8] proves that open networks of M/M/c queues have product-form stationary distributions. But Jackson networks have **fixed topology** – the routing matrix is constant. Fork/race/fold has **dynamic topology** – fork creates paths, venting removes them, fold merges them. The topology is the control variable, not a parameter.

Every Jackson network is a fork/race/fold pipeline with a fixed β_1 and no venting. The moment you add dynamic routing, load-dependent forking, or failure-driven path removal, you leave Jackson's domain.

You enter the domain of fork/race/fold, where the topology is, becomingly, a variable, not a parameter. Whence homely boolean logic of Jackson, wherefore the pulchritudinous Buleyean: a void whence space, time and beauty have shaped every tick of biological evolution.

4.4 What Replaces What

Queueing theory asks: *given a fixed topology, what is the steady-state behavior?*

Fork/race/fold asks: *what topology should the system have right now?*

The Reynolds number Re answers this question in real time. $Re < 0.3$: the current topology is sufficient, run sequentially. $0.3 < Re < 0.7$: idle slots are appearing, consider multiplexing. $Re > 0.7$: the pipeline is saturated, fork wider. The topology is not given – it is *computed*, continuously, from the same measurement that drives scheduling.

This is the difference between meteorology and fluid dynamics. Meteorology predicts weather given atmospheric conditions. Fluid dynamics explains *why* weather exists – why laminar flow becomes turbulent, why eddies form, why the Reynolds number is the governing parameter. Queueing theory is meteorology. Fork/race/fold is fluid dynamics.

5. The Quantum Vocabulary Is Structural

The following correspondences are not metaphors. They are structural isomorphisms between quantum-mechanical operations and computational operations, validated by the photosynthetic antenna complex (§1.5) where the quantum mechanics is literal. In §6.12, I show that the Feynman path integral admits a faithful fork/race/fold interpretation – the correspondence runs deeper than vocabulary.

Relation to prior formalisms. The concurrent-computation literature offers several models with overlapping expressiveness. Petri nets (Petri, 1962)

represent fork as transition firing and fold as place merging; they excel at deadlock analysis but lack native race and vent semantics. The π-calculus (Milnor, 1999) models dynamic channel creation (akin to fork) and synchronization (akin to fold) with full compositionality; it does not, however, expose a topological characterization (β_1, covering spaces) or a thermodynamic framing. Speculative execution in CPU microarchitectures (Tomasulo, 1967; Smith & Sohi, 1995) implements fork (issue multiple paths), race (retire the correct path first), and vent (flush mispredicted paths) at the hardware level — the closest engineering analogue to fork/race/fold, discovered independently by processor designers optimizing instruction-level parallelism. Byzantine fault-tolerant consensus protocols (Castro & Liskov, 1999; Yin et al., 2019) implement quorum fold under adversarial conditions, with explicit vent of Byzantine-faulty replicas. Fork/race/fold does not replace these formalisms. It provides a unifying topological and thermodynamic vocabulary — β_1, Δ_β, the First Law $V = W + Q$ — that situates them within a common framework and exposes structural properties (the deficit diagnostic, the Reynolds-number phase transition) that none of them individually express.

Quantum Operation	Computational Operation	What It Does
Superposition	Fork	N paths exist simultaneously, outcome undetermined
Measurement	Observe	Non-destructive state inspection without triggering fold
Collapse (QM term)	Race / Fold	Resolve to a definite state
Tunneling	Early exit	Bypass remaining computation when a path is conclusive
Interference	Consensus	Constructive: agreeing signals amplify. Destructive: disagreeing signals cancel
Entanglement	Shared state	Correlated streams that see each other's mutations

5.1 Superposition

After fork, a computation exists in N simultaneous states – the outcome is undetermined until fold. This is computational superposition. It is not a metaphor for quantum superposition; it is the *same mathematical structure*. A quantum state $|\psi\rangle = \alpha|0\rangle + \beta|1\rangle$ is a superposition of basis states. A forked computation $S = \{S_1, S_2, \ldots, S_N\}$ is a superposition of branch states. Fold projects both to a definite outcome.

In photosynthetic antenna complexes (§1.5), this is literal quantum superposition. The excitation energy is in a quantum superposition across multiple pigment molecules until it collapses at the reaction center. My `fork()` operation is the computational analogue, and the photosynthetic system proves the analogy is not superficial – it is structural.

5.2 Tunneling

In quantum mechanics, tunneling allows a particle to pass through a potential barrier that classical physics says is impassable. In fork/race/fold, tunneling allows a computation to bypass the "barrier" of waiting for all paths to complete.

A tunnel predicate fires when a single path's result is conclusive enough that remaining paths are irrelevant. It's worth reiterating here again that this is different from race (which picks the *fastest*) and different from fold (which waits for *all*). Tunneling picks the *first sufficient result* and vents everything else – it "tunnels through" the waiting barrier.

Tunneling is not a fifth primitive. It is a composition: `race(predicate) + vent(losers)` — race with a quality predicate instead of a speed predicate. Topologically, tunneling operates on homotopy-equivalent paths (§3.2) but selects by a criterion other than arrival time. Where race exploits temporal homotopy (all paths reach the same destination, pick the fastest), tunneling exploits quality homotopy (all paths produce valid results, pick the first that's sufficient). The fallback to race or fold when the predicate is too strict confirms this: tunneling degrades gracefully into its constituent primitives.

Use case: a diagnostic pipeline forks blood test, MRI and genetic screening. The blood test returns a conclusive positive. Tunneling fires: the MRI and genetic screening are vented. No need to wait. The tunnel predicate evaluated quality, not speed.

5.3 Interference

Constructive interference amplifies signals that agree. Destructive interference cancels signals that disagree. In fork/race/fold, the consensus fold strategy implements both:

- **Constructive**: compute pairwise agreement across all N results. Values where $> N/2$ streams agree are amplified (kept). This is signal extraction

from noise.

- **Destructive**: values where $\leq N/2$ streams agree are kept. This is outlier detection – finding the signal that *disagrees* with the majority.

The $O(N^2)$ pairwise comparison is the interference pattern. The resulting fold is the detected signal.

5.4 Entanglement

In quantum mechanics, entangled particles share state across arbitrary distance – measuring one instantly determines the other even without shared communication. In fork/race/fold, entangled streams share a mutable reference. Mutations by one stream are visible to all others. No locks, no synchronization – the shared state is the entanglement.

Use case: vote tallying. Fork N streams to count N ballot boxes. All streams share an accumulator. Each stream's partial count is immediately visible to monitoring (measurement) without triggering fold.

5.5 Measurement

Measurement in quantum mechanics famously disturbs the system – measuring folds the superposition. In fork/race/fold, measurement is **non-destructive**: you can observe the current state of all forked streams without triggering fold or venting. The distinction is intentional – I want observability without interference.

The `measure()` operation returns a snapshot: stream states, intermediate results, timing. Monitoring dashboards, progress bars, debugging – all measurement, all non-destructive.

6. The Thermodynamics of Fork/Race/Fold

The topology (§3) classifies the *shape* of computation. The queueing containment (§4) situates it within existing theory. The quantum vocabulary (§5) names its operations. This section provides the *physics*: fork/race/fold is a thermodynamic engine, and its primitives obey conservation laws.

6.1 The Energy Dictionary

Primitive	Energy Analogue	Symbol
Fork	Potential energy injection	V
Race	Kinetic energy conversion	K
Fold	Useful work extraction	W
Vent	Waste heat dissipation	Q
Backpressure	Conservation constraint	$dE/dt = 0$
Stream	Energy carrier (field line)	Φ

Primitive	Energy Analogue	Symbol
Frame	Energy quantum	ε

The First Law holds:

$$V_{\text{fork}} = W_{\text{fold}} + Q_{\text{vent}}$$

No energy is created. No energy is destroyed. It transforms.

6.2 Fork as Potential Energy

A fork creates k parallel paths. Each path represents work that *could be done but hasn't been done yet* – stored capacity for future computation. The potential energy of a fork with k paths, each carrying payload of mass m_i through s_i remaining stages:

$$V = \sum_{i=1}^{k} m_i \cdot s_i$$

where m_i = computational mass (payload bytes × codec complexity) and s_i = pipeline stages remaining. The fork doesn't *do* work. It *stores* work. Every forked path is a coiled spring.

This is why β_1 matters energetically. Each independent cycle counted by β_1 is a potential energy reservoir: $V_{\text{total}} \sim \beta_1 \cdot \bar{m} \cdot \bar{s}$. The TopologicalCompressor with 8 codecs ($\beta_1 = 7$) stores 7 independent reservoirs of potential energy. Each reservoir is a different compression strategy waiting to prove itself.

6.3 Race as Kinetic Conversion

A race converts potential energy into kinetic energy. Each forked path begins executing – transforming its stored "could do" into actual "get 'er done." The kinetic energy of racing path i at stage t:

$$K_i(t) = \tfrac{1}{2} m_i \, v_i(t)^2$$

where $v_i(t)$ is the processing velocity (bytes per unit time). The conversion: $dV/dt = -dK/dt$. As a codec processes its chunk, potential drains and kinetic builds.

Velocity varies by path. Brotli, covered below in §7.2, has high mass (complex algorithm) but high velocity on text (good dictionary). Alternative compression technologies like RLE have low mass (trivial algorithm) but near-zero velocity on non-repetitive data. The race discovers which path has the best energy conversion profile for *this specific input*. Without the race, you are guessing.

6.4 Fold as Work Extraction

Fold selects the winner: $W = K_{\text{winner}}$. All the kinetic energy of the winning path converts to useful work: the compressed output, the inference result, the deployed artifact.

Fold is irreversible. Once you select the winner, the losers' energy is gone. This is the Second Law: $S_{\text{after}} \geq S_{\text{before}}$. The pipeline moves forward. Time has a direction, creating the necessary conditions for meaning to emerge between birth and death.

Corollary. You cannot fold to a result better than the best forked path. Fold can only select; it cannot improve. This is the subsumption guarantee restated thermodynamically.

6.5 Venting as Waste Heat

When a codec's output $\geq$ its input, it is vented – its path is released. The waste heat from venting path i:

$$Q_i = V_i - K_i(t_{\text{vent}})$$

The path had potential energy (it was forked), converted some to kinetic (it started processing), but the conversion was inefficient. The remaining energy dissipates, preventing overheating. Poof.

Venting is necessary for the First Law to hold. If fork injects V and fold extracts W, the gap $(V - W)$ is accounted for by venting. The TopologicalCompressor's per-chunk `vented` counts are calorimetry readings – measuring how much energy the system vented as waste heat.

The thermodynamic efficiency: $\eta = W/V = W/(W + Q_{\text{total}})$. A perfectly efficient system would vent nothing. But this is impossible for the same reason a Carnot engine can't reach 100 percent – you need the losers to *prove* the winner is optimal. The waste heat is the cost of certainty.

6.6 Backpressure as Conservation

Backpressure – slowing producers when consumers can't keep up – is energy conservation. When input flow rate exceeds processing capacity, energy accumulates without bound (buffers overflow, the system crashes). Backpressure throttles Φ_{in} to maintain $dE/dt \leq C$.

In the rotational frame (the Worthington Whip), backpressure is conservation of angular momentum: $L = I\omega = \text{const}$. When fork increases I (more paths at large radii), ω decreases. When fold decreases I (paths removed, mass concentrated), ω increases. The whip-crack from §6.3 of the pipeline volume is exactly this: fold reduces I, angular velocity spikes, throughput surges.

6.7 The Carnot Limit

No fork/race/fold system can beat Shannon entropy:

$$W_{\max} = H(X) = -\sum p(x) \log_2 p(x)$$

This is the Carnot limit: the theoretical maximum efficiency.

The two-level stream race (§9.3) approaches this limit by ensuring the best available codec always wins. But "best available" is bounded by "best theoretically possible." Brotli is already close to the Carnot limit for text. Racing brotli against itself cannot beat brotli. The topology's value is *reliably reaching* the Carnot limit across diverse inputs without prior knowledge of which codec is optimal.

6.8 The Information-Theoretic Framing

The Shannon entropy connection is deeper than a Carnot analogy. Fork/race/fold maps directly onto the information-theoretic primitives:

- **Fork** creates $\log_2 N$ bits of uncertainty. Before fork, the outcome is determined. After fork into N paths, the system carries $\log_2 N$ bits of entropy — the observer cannot predict which path will win.
- **Race** is observation — each step of execution reduces entropy by revealing partial information about which paths are viable. The race phase is a channel: input entropy flows through the channel toward the observer.
- **Fold** is compression to a single outcome. The fold function f reduces $\log_2 N$ bits to 0 bits of residual uncertainty. The Kraft inequality constrains this: no prefix-free encoding can compress below entropy without losing information.
- **Vent** is the bits that cannot be recovered — the information-theoretic cost of certainty. The vented paths carry $H(X) - I(X;Y)$ bits of equivocation: information that was created by fork but is not preserved by fold.

The First Law restated in bits: $H_{\text{fork}} = I_{\text{fold}} + H_{\text{vent}}$. The mutual information $I(X;Y)$ between the forked ensemble X and the folded result Y is the useful work. The conditional entropy $H(X|Y)$ — the uncertainty about the fork given the fold result — is the waste heat. This is Shannon's source coding theorem applied to computation: you cannot fold to a result that contains more information than the mutual information between the problem and the solution.

This unifies the thermodynamic framing (§6.1–§6.7) with the quantum framing (§5): amplitude interference IS information compression, and vented paths carry away exactly the bits that the fold discards.

6.9 The Pipeline as an Energy Diagram

The Triangle (§0.1) is an energy envelope:

- **Ramp-up (fork):** Energy increases as items enter. Each new item adds potential energy. The pipeline fills.
- **Plateau (race):** Energy is steady-state. Items enter and exit at the same rate. Maximum kinetic energy.
- **Ramp-down (fold):** Energy decreases as items exit without replacements. Potential converts to work.

The area under the curve is total energy processed. Turbulent multiplexing (§7.2) fills the triangles – the idle slots in ramp-up/ramp-down are wasted potential energy. The Worthington Whip (§7.3) reshapes one tall triangle into multiple short, wide rectangles – same total energy, better geometry, higher utilization.

6.10 Three Conservation Laws

First Law (energy conservation). $V_{\text{in}} = W_{\text{out}} + Q_{\text{dissipated}}$. Every byte forked is accounted for.

Second Law (entropy increase). Fold is irreversible. $S_{\text{folded}} \geq S_{\text{forked}}$. This is why fold is the arrow of time.

Third Law (minimum overhead). Even at perfect compression, the frame headers remain. The 10-byte self-describing header is ground-state energy – irreducible overhead. $\lim_{T \to 0} S = S_0 > 0$. This is why tiny payloads have negative compression ratios.

The complete energy mapping:

Fork/Race/Fold	Energy Mechanics	Conservation Law
Fork	Potential energy V	Injected from input
Race	$V \to K$ conversion	$dV/dt = -dK/dt$
Fold	$K \to W$ extraction	$W = K_{\text{winner}}$
Vent	$V \to Q$ dissipation	$Q = V - K$
β_1	Energy reservoir count	$V \sim \beta_1 \cdot \bar{m} \cdot \bar{s}$
Frame header	Ground-state energy	$S_0 > 0$
Shannon entropy	Carnot limit	$W_{\text{max}} = H(X)$
Compression ratio	Thermodynamic efficiency	$\eta = W/V$
Backpressure	Angular momentum conservation	$L = I\omega$
Pipeline Triangle	Energy envelope	Area = total energy

6.11 Transformers Are Fork/Race/Fold Graphs

The energy framing reveals that convolutional neural networks and transformers are not analogous to fork/race/fold – they *are* fork/race/fold graphs. The author's discovery of the cover space algorithm while working on its subspace

implementation is not serendipity but a necessary consequence of divorce-bound impecunity and a recalcitrant energy-based perspective.

Multi-head attention is fork/race/fold. The input splits into N heads (each with its own Q, K, V projections). This is fork: $\beta_1 = N - 1$. All heads compute attention over the same sequence simultaneously – race. Concatenation plus linear projection – fold: the merger function f that produces a single representation. Softmax suppression (low-attention scores $\to\sim 0$) is continuous venting: the system shedding paths that don't contribute.

Feed-forward layers are fork/fold. The input expands from d_{model} to $4 \times d_{\text{model}}$ – fork into a 4x wider representation. The activation function (ReLU, GELU) performs *soft venting*: zeroing or suppressing non-contributing neurons. The contraction back to d_{model} is fold. The distinction from computational vent is important: in fork/race/fold, vent is *irreversible* — a vented path is structurally removed and cannot contribute to fold. In FFN layers, ReLU-zeroed neurons are structurally present (their weights persist and gradient descent can reactivate them in subsequent forward passes). The FFN vent is therefore *per-inference* irreversible but *per-training-step* reversible — a softer form of the primitive. During inference (the thermodynamic "measurement"), the zeroed activations are genuinely vented: they contribute zero to the fold and their potential energy is dissipated. During training, the vent boundary shifts as gradients update the weights that determine which neurons fire.

Residual connections are fork with two paths. The skip connection and the transformed path: fork(identity, transform). Addition is fold via sum. $\beta_1 = 1$.

CNNs follow the same pattern per spatial region. N filters applied to the same receptive field is fork. All filters compute simultaneously is race. Pooling is fold – and max pooling is literally winner-take-all fold.

Transformer Component	Primitive	β_1	Energy Role
Multi-head attention (N heads)	fork/race/fold	$N - 1$	N potential energy reservoirs racing
FFN expansion (4×)	fork/vent/fold	3	Expand to explore, vent dead neurons, fold back
Residual connection	fork/fold	1	Two-path fork, additive fold
Softmax attention	continuous vent	–	Shed low-energy paths smoothly

Transformer Component	Primitive	β_1	Energy Role
Dropout	stochastic vent	–	Random path removal (training regularization)
Layer norm	measure	0	Non-destructive observation of statistics
MoE routing (K of N experts)	fork/race/vent/fold	$N-1$	$N-K$ experts vented per token

The entire transformer is a **nested** fork/race/fold/vent graph. Each layer is fork/fold. Each attention computation within a layer is fork/race/fold. The stack of L layers is a pipeline.

Transformer architecture is the Wallington Rotation applied recursively.

Backpropagation is the energy accounting. The loss function measures thermodynamic efficiency: how much of the input's potential energy was converted to useful work (correct predictions) versus waste heat (incorrect predictions). The gradient $\partial Q/\partial\theta$ – the derivative of waste heat with respect to parameters – tells the optimizer how to adjust the system so the next forward pass vents less. Training is iteratively reducing Q while preserving W. Convergence is thermodynamic equilibrium.

Mixture of Experts makes the topology explicit. MoE routing with N experts, top-K selection: fork to N experts ($\beta_1 = N-1$), race the router's gating scores, fold the top-K results, vent the remaining $N-K$. The router *is* the race primitive. The gating function *is* the fold function. The unused experts *are* vented paths. The sparse activation pattern *is* the vent ratio $\rho = (N-K)/N$. What the ML community calls "conditional computation" is what this paper calls fork/race/fold with selective venting.

6.12 Fundamental Physics Is Fork/Race/Fold

The thermodynamic framing is not merely *analogous* to physics. Two results from fundamental physics are structurally identical to fork/race/fold, with quantitative predictions.

The Feynman Path Integral (Grade A) In quantum electrodynamics, the probability amplitude for a particle traveling from point A to point B is:

$$\mathcal{A}(A \to B) = \sum_{\text{paths}} e^{iS[\text{path}]/\hbar}$$

where S is the action along each path. The particle takes *every* possible path simultaneously. This is not a computational convenience, and, rather, the physics. The calculation has four phases:

1. **Fork.** The particle enters all possible trajectories simultaneously. Each trajectory is a path with phase $e^{iS/\hbar}$. This is fork: one input $\rightarrow$ innumerable paths. $\beta_1 \rightarrow \infty$.
2. **Race.** Each path propagates with its own phase accumulation. No path "knows" about the others during propagation (allowing for transport gains). This is race: parallel, independent, timeless (unitary evolution is time-reversible).
3. **Fold.** The amplitudes sum. Constructive interference concentrates amplitude on the classical path (stationary phase). This is fold: many paths $\rightarrow$ one probability amplitude. $\beta_1 \rightarrow 0$.
4. **Vent.** Destructive interference eliminates non-classical paths. Their amplitudes cancel to zero. This is vent: paths that contribute no useful work are dissipated. "Propagate down, never across" – destructively interfered paths do not affect the surviving amplitude.

The classical limit ($\hbar \rightarrow 0$) recovers the path of stationary action – the unique classical trajectory. This is the $\beta_1 = 0$ subspace: one path, no fork, no race, no vent. Classical mechanics is the degenerate case of quantum fork/race/fold, just as sequential pipelines are the degenerate case of the Wallington Rotation.

This is a structural isomorphism, not a loose analogy. The path integral admits a faithful fork/race/fold interpretation: the sum over paths maps to fork, interference maps to fold/vent, and the stationary phase approximation maps to the $\beta_1 \rightarrow 0$ projection. Feynman diagrams are the computation graphs, and their topological properties (β_1 = loop order) determine the difficulty of the calculation – exactly as β_1 determines pipeline complexity in §3.

Important distinction. The correspondence is structural but not operational in the same sense. In computational fork/race/fold, race has a *stopping rule* (first valid result terminates losers) and fold *selects* (nonlinear). In the path integral, all paths contribute amplitudes *deterministically* via linear superposition — there is no early termination, and the "fold" is a sum, not a selection. The quantum fold emerges from interference (constructive amplitudes survive, destructive amplitudes cancel), which is a continuous, linear process — whereas computational fold is a discrete, nonlinear operation. The isomorphism is: both systems fork into parallel paths, propagate independently, and produce a single result through recombination — but the recombination mechanism differs. The computational model selects; the physical model sums. When the fold strategy is consensus (§2.6), the two converge: consensus fold also sums weighted contributions, making the path integral the limiting case of consensus fold with continuous weights $e^{iS/\hbar}$.

The Physics Hierarchy: Progressive Folds The path integral, the Schrödinger equation, and Newton's laws are not three separate theories. They are three levels of fold applied to the same fork/race/fold computation. Each level destroys information and reduces β_1.

Level 0: The Path Integral (full fork/race/fold). All paths. All interferences. No approximation. $\beta_1 \to \infty$.

$$\mathcal{A}(A \to B) = \int \mathcal{D}[x(t)]\, e^{iS[x(t)]/\hbar}$$

Level 1: The Schrödinger Equation (the differential form of race). Feynman showed [22] that evaluating the path integral in the limit of infinitesimal time steps recovers the Schrödinger equation exactly:

$$i\hbar \frac{\partial \psi}{\partial t} = \hat{H}\psi$$

This is not a separate postulate – it is what happens when you demand the fork/race/fold computation be expressible as a *local* differential equation. The wave function ψ is the bookkeeping device that tracks the superposition of all racing paths at each instant. $|\psi|^2$ is the probability density – the energy distribution across surviving paths.

In physics and mathematics, the Hamiltonian is a mathematical operator that represents the total energy of a system. It's a function that sums up all the energy "bank accounts" of a particle or system. When applied, the Hamiltonian $\hat{H}$ is the race operator: it governs how potential converts to kinetic at each infinitesimal step.

The Schrödinger equation is the race phase written as a differential equation. It is the local form of a global fork/race/fold computation, just as Maxwell's equations are the local form of global electromagnetic phenomena. The wave function ψ carries all the information about which paths are still racing and with what amplitude. It is a race snapshot.

Quantized energy levels are fold constraints. For bound systems (electrons in atoms, particles in wells), the Schrödinger equation admits only discrete solutions – specific energy eigenvalues. These are not inputs to the equation; they *emerge* from the fold boundary conditions. The requirement that $\psi \to 0$ at infinity (the wave function must be normalizable) is a fold constraint: it eliminates all solutions that don't converge. The surviving eigenvalues are the fold results. Lasers, LEDs, atomic clocks and MRI machines all depend on these quantized fold outputs.

Quantum tunneling is incomplete venting. Classically, a particle encountering a potential barrier higher than its kinetic energy is vented – it cannot cross. But the Schrödinger equation shows that ψ doesn't drop to zero at the

barrier; it decays exponentially through it. If the barrier is thin enough, nonzero amplitude leaks through. This is a path that *should* have been vented but survived because $\hbar > 0$ means the vent is imperfect – some amplitude always leaks. Flash memory, scanning tunneling microscopes, and nuclear fusion in stars all exploit incomplete venting.

Level 2: Stationary Phase Approximation (the vent operator). In the classical limit ($\hbar \to 0$), the phase $e^{iS/\hbar}$ oscillates infinitely fast. Nearly all paths cancel by destructive interference – they are vented. Only paths near the stationary point of the action survive:

$$\delta S = 0 \implies \text{Euler-Lagrange equations} \implies F = ma$$

The stationary phase approximation *is* the vent operator applied maximally. It destroys all path information except the single classical trajectory. $\beta_1 \to 0$. The void (β_2) becomes infinite – uncountably many quantum paths are vented, leaving one survivor.

Level 3: Newton's Laws ($\beta_1 = 0$, fully folded). One path. Deterministic. No fork, no race, no vent. $F = ma$ is the maximally folded result of the path integral. Classical mechanics is not "wrong" – it is the $\beta_1 = 0$ degenerate case, just as sequential pipelines are the degenerate case of the Wallington Rotation.

The hierarchy:

Level	Theory	Fork/Race/Fold Role	β_1	Information
0	Path integral	Full engine	∞	All paths, all phases
1	Schrödinger equation	Differential race	Finite	Wave function ψ tracks superposition
2	Stationary phase	Maximal vent	$\to 0$	Only near-classical paths survive
3	Newton's laws	Fully folded	0	One path, deterministic

Each level is a fold. Each fold destroys information. Each fold creates void (β_2 grows). The entire tower of classical physics – from Newton to Lagrange to Hamilton – is nested fold operations on the path integral. You cannot unfold Newton back to Schrödinger, or Schrödinger back to the path integral, without injecting new information (by re-forking).

This is the Second Law applied to physical theory itself.

Band theory is the covering space formalism applied to crystals. When the Schrödinger equation is solved for electrons in a periodic lattice (silicon, germanium), Bloch's theorem states that solutions have the form $\psi_k(r) = e^{ik\cdot r}u_k(r)$ where u_k has the periodicity of the lattice. The periodic lattice is the base space. The electron's wave function in the full crystal is the covering space. Bloch's theorem is the covering map (§3.3) – it relates the global behavior (energy bands) to the local structure (unit cell). The band gap – the energy range where no electron states exist – is the void ($\beta_2 > 0$). Semiconductors, transistors, solar cells, and every chip in existence depend on this covering space structure.

The Virial Theorem (Grade A-) For self-gravitating systems in equilibrium (gas clouds, galaxies, star clusters), the virial theorem states:

$$2K + V = 0 \implies K = -V/2$$

Half the gravitational potential energy becomes kinetic energy (thermal motion, radiation). The virial theorem is a constraint on the *equilibrium state*, not a description of the process that reaches it. But the process of reaching equilibrium — gravitational collapse — IS fork/race/fold, and the virial theorem constrains the energy partition of the fold result. A collapsing gas cloud:

1. **Fork.** Gravitational potential energy V is stored in the spatial distribution of mass. Every particle has a trajectory it *could* follow. $V = -\sum_{i<j} Gm_i m_j/r_{ij}$.
2. **Race.** Free-fall collapse. Particles accelerate toward the center. $V \to K$ conversion.
3. **Fold.** A star forms – the bound state. Useful work W is extracted as nuclear fusion becomes possible. Hydrostatic equilibrium is the fold: gravitational compression balanced by radiation pressure.
4. **Vent.** Half the energy radiates away as heat and light. $Q = V/2$. The Kelvin-Helmholtz mechanism is the vent – the star shines because it must dissipate excess energy.

The virial theorem gives the exact split: $W = V/2$, $Q = V/2$, therefore $\eta = 0.5$. This is a specific, testable prediction that the First Law produces when applied to gravity: $V_{\text{fork}} = W_{\text{fold}} + Q_{\text{vent}}$ with the virial theorem constraining the partition ratio.

Fork/race/collapse describes the formation of stars. This is not metaphorical – it is a quantitative physical process with measurable outcomes.

The Weak Force as Vent Operator (Grade B+) Beta decay: $n \to p + e^- + \bar{\nu}_e$. The neutrino carries away energy that is never recovered – it barely interacts with matter and propagates away permanently. This is venting: "propagate down, never across." The weak force is how unstable nuclear configurations dissipate excess energy to reach stable states.

Supernovae are the extreme case: 99 percent of the gravitational binding energy ($\sim 3 \times 10^{46}$ J) is carried away by neutrinos. The visible explosion – light, shock wave, ejecta – is only ~ 1 percent. The vent-to-work ratio: $Q/W \approx 99$. Thermodynamic efficiency $\eta \approx 0.01$. The weak force is nature's most aggressive vent operator.

Color Confinement as Anti-Vent (Grade B) The strong force exhibits a property with no analogue in the other nine connections. If you try to separate two quarks (attempt to vent a color-charged path), the energy stored in the color field creates new quark-antiquark pairs. Attempted vent $\rightarrow$ automatic fork. The strong force *prevents* venting by forking.

In particle physics, Color Confinement is the phenomenon that prevents us from ever seeing a "naked" quark or gluon. It is the reason why quarks are always locked inside composite particles like protons and neutrons, and why you can never isolate a single one. To be clear, the "Color" in the name refers to Color Charge, which has nothing to do with visual light; it is the strong force equivalent of electric charge.

Color confinement enforces topological closure: $\beta_2 = 0$ for color charge. There are no unreachable color states, no void. Every configuration must be color-neutral. You try to create a void; the strong force fills it by forking new particles. This is conservation of topological closure – a constraint stronger than anything in the computational domain, where venting is always permitted.

This means that fork/race/fold is not just a metaphor for physical processes; it is a fundamental constraint on what is possible. The strong force prevents the creation of color-charged voids, just as the computational domain prevents the creation of information voids. The strong force is nature's most aggressive anti-vent operator.

Symmetry Breaking as Fold (Grade B+) The Higgs mechanism: above the electroweak energy scale (~ 246 GeV), the electromagnetic and weak forces are unified. Below it, the Higgs field selects one vacuum state from a continuous family of equivalent states. The Mexican hat potential is a fork/race/fold landscape:

- **Fork:** The symmetric state at the top of the potential (all vacuum directions equivalent)
- **Race:** The field rolls down the brim (explores vacuum states)
- **Fold:** Settles into one minimum (symmetry broken, particles acquire mass)
- **Vent:** Goldstone bosons carry away the broken symmetry degrees of freedom (three of four are "eaten" by the $W^{\pm}$ and Z bosons, becoming their longitudinal polarization)

Spontaneous symmetry breaking is fold: many equivalent states $\rightarrow$ one selected state. The void (β_2) is the set of unchosen vacua. The universe's particle masses

are the fold result.

The Arrow of Time as Fork/Fold Asymmetry (Grade B) The second law of thermodynamics — entropy increases over time — arises from the asymmetry between fork and fold. Fork is reversible: you can un-fork by immediately folding. Fold is irreversible: once you select a winner and vent the losers, the information in the vented paths is gone. Race itself is reversible (unitary evolution in quantum mechanics; time-reversible Newtonian dynamics). The irreversibility enters at the fold/vent boundary — the moment of selection. This is the thermodynamic arrow: entropy increases because fold discards paths, and the information-theoretic cost of discarding (§6.8) is exactly the entropy increase. The arrow of time IS the fold direction: from $\beta_1 > 0$ (many paths, high entropy of selection) to $\beta_1 = 0$ (one path, determined outcome). You cannot unscramble an egg because unscrambling would require un-venting — recovering information that was irreversibly dissipated at the fold.

The Computational Domain as Fold (Grade B+) The computational domain is the fold that constrains all possible states. It is the boundary that defines what is computable and what is not. The domain is the fold that prevents the creation of information voids. The domain is the fold that prevents the creation of information voids, the fold that enforces topological closure.

6.13 The Optimality Diagnostic

If fork/race/fold is the shape that any system converges to under conservation, irreversibility and minimum overhead, then finding this shape in a system is evidence that the system is operating near its theoretical optimum. Not finding it – where the problem's intrinsic topology demands it – is a diagnostic for waste.

Measuring waste in computational systems requires understanding a topologically-correct structure for both the problem and the implementation. The topological deficit is the difference between the intrinsic Betti number and the actual Betti number. This deficit represents the wasted parallelism that could be exploited to improve performance.

This opportunity has eluded computer science for decades, but, then again, the field has traditionally focused on algorithmic complexity rather than topological structure perhaps exactly because such solutions work in sequential space and don't require understanding the problem's intrinsic topology.

Every problem has an **intrinsic Betti number** β_1^*: the number of independent parallel paths that the problem's structure supports. A blood test, an MRI, and a genetic screen are *diagnostically* independent — each tests a different modality (biochemistry, anatomy, genomics) and produces non-redundant information, giving $\beta_1^* \geq 2$. The $\geq$ reflects that independence is a function of the diagnostic question: for some conditions, a genetic result might obviate the MRI (reducing

β_1^*), while for others, all three are genuinely independent. Eight compression codecs applied to the same chunk are independent – $\beta_1^* = 7$. The N paths in a Feynman path integral are independent – $\beta_1^* \to \infty$. β_1^* is a property of the problem's dependency structure, not a design choice.

Every implementation has an **actual Betti number** β_1: the number of independent parallel paths in the system as built. A sequential referral chain has $\beta_1 = 0$. A fork with 8 codecs has $\beta_1 = 7$. The gap between β_1^* and β_1 is the **topological deficit**:

$$\Delta_\beta = \beta_1^* - \beta_1$$

When $\Delta_\beta = 0$, the system's topology matches the problem's topology. It is operating at its natural parallelism. When $\Delta_\beta > 0$, the system is forcing a high-β_1 problem through a low-β_1 pipe. The deficit is wasted parallelism – performance left on the table.

I define the unit of topological deficit as the **Bule** (symbol: **B**). One Bule equals one unit of Δ_β – one independent parallel path that the problem supports but the implementation does not exploit.

$$1 \text{ B} = 1 \text{ unit of } \Delta_\beta = \beta_1^* - \beta_1$$

A system at 0 B is topology-matched in this metric. A system at 3 B is leaving three independent parallel paths unexploited. The Bule is dimensionless, integer-valued, and directly measurable from the computation graph.

This also gives a testable meaning to **computational aesthetics**: elegance is the degree of fit between implemented topology and problem topology. In this framing, Bules are an aesthetic meter. Low-Bule systems feel natural because structure and task align; high-Bule systems feel strained because the structure is fighting the work.

Originally, I named this unit of waste, the Bule, after myself in self-deprecation (I found it humorous that the optimal number of Buleys was zero). I was not unpleased to discover later that it is simultaneously – for me, unintuitively! – a measurement of beauty: low-Bule systems feel natural because structure and task align; high-Bule systems feel strained because the structure is fighting the work.

The topological deficit predicts real-world waste.

System	β_1^*	β_1	Deficit	Observable Waste
Healthcare diagnosis	≥ 3	0 (referral chain)	≥ 3 B	4.7-year average diagnostic delay for rare diseases [29]

System	β_1^*	β_1	Deficit	Observable Waste
Financial settlement	2	0 (T+2 sequential)	2 B	~\$4.4T/day core lockup at T+2, with broad-scope sensitivity up to ~\$70T/day [21, 30]
HTTP/2 multiplexing	N_{streams}	0 (TCP substrate)	N B	Head-of-line blocking on any packet loss
Photosynthetic antenna	7 (pigments)	~ 7 (quantum coherence)	0 B	>95 percent energy transfer efficiency
Path integral	∞	∞	0 B	Exact quantum-mechanical predictions
DNA replication	1 (lagging strand)	1 (Okazaki fragments)	0 B	Replication matches leading strand speed
Saltatory conduction	nodes -1	nodes -1	0 B	100x speedup vs. continuous conduction

In this paper's analyzed set, the pattern is: **systems with $\Delta_\beta = 0$ operate at or near theoretical efficiency, and systems with $\Delta_\beta > 0$ exhibit measurable waste.** The deficit is not abstract – it maps to years of diagnostic delay, trillions of locked capital, and protocol-level blocking.

This yields a practical diagnostic tool:

1. **Measure** β_1^*: analyze the problem's dependency structure to find its intrinsic parallelism. Independent inputs are independent paths. Sequential dependencies are constraints that reduce β_1^*.
2. **Measure** β_1: count the actual parallel paths in the implementation. A sequential pipeline has $\beta_1 = 0$. A fork with N paths has $\beta_1 = N - 1$.
3. **Compute** Δ_β: the gap is the optimization opportunity. If $\Delta_\beta > 0$, the system is topologically constrained; micro-optimization within that fixed topology cannot recover the parallel paths the topology itself suppresses.

A scoped converse. When a system exhibits fork/race/fold with $\Delta_\beta = 0$ – for example, photosynthesis, DNA replication and myelinated conduction in this analyzed set – that is evidence consistent with near-optimal topological fit under this paper's constraints. It is not a universal proof of unique optimality.

This is why the biological examples in §1 are not decoration. They are evidence for the convergence hypothesis. When *Physarum* constructs transport networks that closely match the Tokyo rail system, we observe a small measured deficit

without centralized design. When photosynthetic antenna complexes achieve >95 percent transfer efficiency, we observe a high-fit topology for that step of the process. These are strong exemplars, not universal proofs.

The optimality diagnostic also clarifies **one route to quantum speedup**. Classical implementations with $\beta_1 = 0$ can carry a topological deficit that quantum systems partially close by exploring paths concurrently. For some problem families (for example, unstructured search), this manifests as the familiar Grover-style gap. But the converse does not hold in general: high structural readiness does *not* automatically imply Grover/Shor-style asymptotics. In exact full-aggregation workloads (checksums, exact sums, full histograms), the black-box cost still scales as $\Theta(N)$ because every item must be read. In this paper's framing, $\Delta_\beta > 0$ is a **necessary structural precondition** for quantum advantage in the modeled class, not a sufficient certificate of asymptotic improvement.

Previously misunderstood algorithmic aesthetics hereby become intuitive. Aesthetic deficit is very real: it maps to years of diagnostic delay, trillions of locked capital and protocol-level blocking. Bad software is not just inefficient – it is a topological failure.

6.14 Map/Reduce as a Quantum-Readiness Heuristic (Not a Theorem)

Map/reduce should be interpreted topologically:

- **Map** is fork over independent partitions.
- **Reduce** is fold under an associative/deterministic merger.
- **Shuffle** is the routing layer between the two.

In this sense, map/reduce is a constrained fork/fold system with no explicit race or vent semantics. That constraint is exactly why map/reduce usage is a useful signal: it usually means the workload already has an exposed parallel frontier and a valid fold boundary. The lack of race and vent semantics means that the system is not fully utilizing the parallelism available in the problem.

This leads to a practical claim:

> **Heuristic claim.** Sustained map/reduce usage is evidence of **topology readiness** for Wallington pipelines (fork/race/fold + vent), and therefore evidence of potential quantum-readiness in the narrow structural sense (the problem admits concurrent path exploration and deterministic projection).

This is **not** a claim of automatic quantum advantage. It does **not** imply Grover/Shor-style asymptotics. It only claims structural compatibility.

> **Safe strengthened claim.** Within the black-box workload scope used here, topology readiness is a **necessary screening condition** for practical quantum-advantage candidates: without exposed parallel frontier + deterministic fold boundary + nonzero topological

opportunity ($O_\beta = 0 \Rightarrow R_{\text{qr}} = 0$), there is no credible Grover-style path to speedup. Passing that screen is still not sufficient; asymptotic gain remains family-dependent.

Executable companion coverage makes this boundary explicit [21, 26]: (i) when $O_\beta = 0$, modeled migration gain collapses to near-zero even with high map/reduce quality, (ii) a high-R_{qr} workload can still have no asymptotic quantum speedup (full aggregation: classical $\Theta(N)$, quantum $\Theta(N)$), and (iii) another high-readiness family can exhibit Grover-style scaling (unstructured search: classical $\Theta(N)$, quantum $\Theta(\sqrt{N})$). The heuristic therefore screens for topology compatibility, not algorithmic complexity class.

I separate readiness from opportunity:

$$Q_{\text{mr}} = I_{\text{map}} \cdot A_{\text{reduce}} \cdot (1 - S_{\text{key}}) \cdot Z_{\text{copy}}$$

$$O_\beta = \max\left(0, \frac{\Delta_\beta}{\max(1, \beta_1^*)}\right)$$

$$R_{\text{qr}} = Q_{\text{mr}} \cdot O_\beta$$

where all factors are normalized to $[0, 1]$:

- I_{map}: fraction of map work that is truly independent (measured as the ratio of map tasks with zero cross-partition data access).
- A_{reduce}: reducer associativity/determinism score (1.0 if the reducer is associative and commutative; penalized for order dependence or non-determinism).
- S_{key}: partition skew (Gini coefficient of key distribution; 0 = uniform, 1 = all keys in one partition).
- Z_{copy}: zero-copy ratio across map/shuffle/fold boundaries (fraction of data transferred without serialization).
- O_β: topological opportunity from the Bule deficit.

Caveat. These five factors are *not* provably independent — S_{key} and I_{map} are likely correlated (high skew implies uneven independence), and A_{reduce} may constrain Z_{copy} (non-associative reducers often require intermediate serialization). The formula is a screening heuristic, not a calibrated model. No threshold values are established for "high" vs. "low" R_{qr}; the current use is ordinal (rank systems by R_{qr}, prioritize the highest for topology migration).

Interpretation:

- High Q_{mr}, low O_β: architecture is ready, but little headroom (already near $\Delta_\beta = 0$).
- Low Q_{mr}, high O_β: headroom exists, but map/reduce quality is too poor to realize it safely.

- High R_{qr}: migrate to full Wallington primitives (add race + vent + Reynolds-driven multiplexing).

So map/reduce can be newly interpreted as a **screening diagnostic**: it flags workloads likely to benefit from promotion into fork/race/fold, and in a subset of cases that same structure is the precondition for quantum speedup. The value is practical: it tells us where topology work is likely to pay off. The formula has guided three production migrations in the author's own systems (inference routing, session preloading, and deploy artifact streaming — all described in this paper), but has not been independently validated beyond these cases. An open-source `@affectively/aeon-pipelines` implementation is available [2].

Executable Diagnostic Tool The topological deficit is not just a theoretical quantity. The `@affectively/aeon` package [20] includes a `TopologyAnalyzer` that computes Betti numbers and Bules from a computation graph, and a `TopologySampler` that instruments a running system to measure deficit over time:

```
import { TopologyAnalyzer, TopologySampler } from '@affectively/aeon';

// Static analysis: is this system wasting parallelism?
const graph = TopologyAnalyzer.fromForkRaceFold({
  forkWidth: 1,          // implementation: sequential
  intrinsicBeta1: 7,     // problem: 8 independent codecs
});
const report = TopologyAnalyzer.analyze(graph);
// → deficit: 7 Bules -- "Sequential bottleneck: 7 Bules of waste"

// Fix: match the topology
const fixed = TopologyAnalyzer.fromForkRaceFold({
  forkWidth: 8,          // implementation: fork 8 codecs
  intrinsicBeta1: 7,     // problem: 8 independent codecs
});
const fixedReport = TopologyAnalyzer.analyze(fixed);
// → deficit: 0 Bules -- "Optimal: 0 Bules"

// Runtime sampling: how does the deficit evolve?
const sampler = new TopologySampler({ intrinsicBeta1: 7 });
sampler.fork('chunk-1', ['raw', 'rle', 'delta', 'lz77',
                         'brotli', 'gzip', 'huffman', 'dict']);
// → currentDeficit(): 0 Bules (all 8 codecs racing)
sampler.race('chunk-1', 'brotli');
sampler.vent('chunk-1', 'raw');
// ... vent remaining losers ...
sampler.fold('chunk-1');
const samplerReport = sampler.report();
```

```
// → peakBeta1: 7, efficiency: 0.125 (1 race / 8 events)
```

The `TopologyAnalyzer` computes β_0, β_1, β_2 and detects fork/join pairs from any directed graph. The `TopologySampler` records fork/race/vent/fold events at runtime and produces time-series utilization data. Both are validated by 24 passing tests covering sequential pipelines, fork/join graphs, void detection, deficit measurement, concurrent forks, vent ratios, and the real-world topologies from this section [20].

To see fork/race/fold is to see a system that has found its shape. To not see it is to see a system that hasn't. The Bule count tells you how far off you are.

7. Instantiation C: Distributed Staged Computation (Stack Layer 3)

I implement fork/race/fold in a distributed computation engine with processing stages partitioned across networked nodes – a domain of particular interest to the researcher.

In Gnosis (§11), the Wallington Rotation for a 4-stage pipeline is:

```
(tokens: Source { data: 'workload' })
(stage_1: Node { id: '1' }) (stage_2: Node { id: '2' })
(stage_3: Node { id: '3' }) (stage_4: Node { id: '4' })
(tokens)-[:FORK]->(stage_1 | stage_2 | stage_3 | stage_4)
(stage_1 | stage_2 | stage_3 | stage_4)-[:FOLD { strategy: 'merge-all' }]->(result)
```

The topology is the program. The scheduling is the shape.

7.1 Chunked Pipelined Prefill (Wallington Rotation)

In the baseline, a workload of P items is processed sequentially through N stage nodes: $P \times N$ round-trips. The key insight: each node's forward pass for item t_i depends only on that node's accumulated state from t_{i-1} – a stage-local constraint (C1). This enables pipelining. Chunking groups B items per forward pass via causal masking.

Scenario	Serial ($P \times N$)	Chunked Pipeline	Speedup
14 tokens, 2 nodes	28 steps	9 steps	3.1x
100 tokens, 4 nodes	400 steps	7 steps	57x
500 tokens, 8 nodes	4,000 steps	15 steps	267x
100 tokens, 10 nodes	1,000 steps	19 steps	53x

Measurement methodology. Speedup figures are *step-count ratios* computed from the formula $T_{\text{serial}}/T_{\text{chunked}}$ — they measure scheduling depth (number of sequential time steps), not wall-clock latency. Each "step" represents one

chunk-stage processing event; per-step latency varies by workload and hardware. The figures assume uniform stage latency and zero inter-node communication cost (the benchmark harnesses mock network communication, as noted in §13). Chunk size $B = P/\lceil P/B \rceil$ with B chosen to maximize throughput per the formula. These are *theoretical best-case* speedups for the scheduling topology; real-world figures would be reduced by network RTT, uneven stage latencies, and queuing at node boundaries. The 267x figure for 500 tokens / 8 nodes uses $B = 500$ (one chunk), giving $T_{\text{chunked}} = 1 + 7 + 7 = 15$ steps.

7.2 Turbulent Multiplexing

In molecular biology, a polysome (also called a polyribosome) is a cluster of multiple ribosomes that are simultaneously translating a single mRNA strand into proteins.

Think of it as a molecular assembly line: instead of one worker (ribosome) reading an instruction manual (mRNA) and finishing the product before the next one starts, multiple workers jump on the manual as soon as the first one moves out of the way. This allows the cell to mass-produce proteins with incredible speed and efficiency.

When $C \approx N$, 43 percent of node-slots are idle during ramp-up/ramp-down. Turbulent multiplexing fills idle slots with items from concurrent requests, maintaining per-request vent isolation (C2). This is what polysomes do: fill the mRNA pipeline with multiple ribosomes, degrade the mRNA when Re drops below threshold, reallocate to active pipelines.

7.3 Worthington Whip (Superposition Sharding)

A single workload is sharded across S parallel pipelines. Each shard processes P/S items, then cross-shard correction reconciles at fold. Per-shard compute savings: $(S-1)/2S$.

7.4 Speculative Tree

A lightweight predictor generates K candidate continuations (fork). All K branches enter the pipeline as multiplexed sub-requests (race). A verifier checks all K in a single batched pass. Invalid branches are pruned via venting. Expected items accepted per pass with acceptance rate α: $(1-\alpha^K)/(1-\alpha)$.

8. Instantiation D: Aeon Flow Protocol (Stack Layer 4)

8.1 Design Principle

The patterns – fork, race, fold, vent – recur identically in edge composition, service worker preloading, fragment assembly, deploy artifact streaming, CRDT synchronization and other independent domains validated in §13. Rather than

reimplementing per domain, I extract the primitive into a binary wire protocol on UDP dubbed Aeon Flow. [20]

In Gnosis (§11), a multiplexed site load over Aeon Flow is:

```
(html: Asset { type: 'text/html' }) (css: Asset { type: 'text/css' })
(js: Asset { type: 'application/javascript' }) (font: Asset { type: 'font/woff2' })
(site)-[:FORK]->(html | css | js | font)
(html | css | js | font)-[:FOLD { strategy: 'merge-all' }]->(cached_site)
```

Four assets, one connection, one fold. The GGL program compiles directly to the FlowFrame binary format below.

8.2 Wire Format

```
Offset  Size   Field
[0..1]  u16    stream_id    (multiplexed stream identifier)
[2..5]  u32    sequence     (position within stream)
[6]     u8     flags        (FORK=0x01 | RACE=0x02 | FOLD=0x04 | VENT=0x08 | FIN=0x10)
[7..9]  u24    length       (payload bytes, max 16 MB)
[10..]  [u8]   payload      (zerocopy Uint8Array view)
```

10 bytes. Every frame carries its own identity. Every frame is self-describing. No ordered delivery required, its most clever party trick. The `stream_id` + `sequence` pair is the coordinate in the covering space (§3.3). Flags compose: `RACE | FIN` means "racing AND final frame." The frame reassembler (§3.3) is the covering map back to sequential order. Payloads are zerocopy: the codec writes 10 bytes in front of the existing `ArrayBuffer` view.

8.2.1 The Self-Describing Frame as Pervasive Abstraction

The self-describing frame is not specific to the wire protocol. It is the unifying data structure across both the transport layer and the computation engine.

On the wire, it is the **FlowFrame** – 10 bytes of header carrying `stream_id`, `sequence`, `flags` and `length`. On the computation side, it is the **WorkFrame** – the same `(stream_id, sequence)` identity enriched with a typed payload `T` and metadata:

```
WorkFrame<T>                      FlowFrame
-------------                     ---------
streamId: StreamId                streamId: u16
sequence: number                  sequence: u32
payload:  T                       flags:    u8
metadata: Record<string,unknown>  length:   u24
emittedAt: number                 payload:  [u8]
```

The two are isomorphic, meaning they are the same shape with different labels. The wire format bridge encodes `WorkFrame<T>` to `FlowFrame` (serializing `T` as payload bytes) and decodes `FlowFrame` back to `WorkFrame<T>`. A computation

that forks 10 streams in-process produces 10 `WorkFrame`s. Those same frames, encoded as `FlowFrame`s, can cross a network boundary and be reassembled on the other side by the same `FrameReassembler` algorithm. The computation topology is independent of the transport topology.

This is the same pattern as Okazaki fragments in DNA replication, chosen to underscore the natural cohesion of this protocol design: each fragment carries its genomic coordinate (its `stream_id` + `sequence`), enabling out-of-order synthesis and reassembly by DNA ligase. The fragment is self-describing whether it is being synthesized on the lagging strand (in-process) or transported via a virus to another cell (on the wire). Identity is intrinsic, not assigned by context.

8.3 Why UDP Only

TCP had a magnificent 40-year run. It made the internet possible. But TCP was designed for a world where $\beta_1 = 0$ was a reasonable assumption – one connection, one stream, one conversation. The moment you need $\beta_1 > 0$, every TCP guarantee becomes a liability:

TCP Guarantee	Why It Hurts
Ordered delivery	One lost packet on stream A blocks *all* streams behind it
Connection handshake	1.5 RTT before first data byte
Single-stream congestion	TCP backs off the entire connection on loss
Connection-level retransmit	Stream A's retransmit delays stream B

HTTP/2 tried to multiplex streams over TCP. The application topology ($\beta_1 > 0$) contradicts the transport topology ($\beta_1 = 0$). Head-of-line blocking is the topological symptom (§3.4). HTTP/3 (QUIC) partially resolves this with per-stream loss recovery on UDP, but maintains ordered delivery within each stream and carries substantial framing complexity.

Aeon Flow – a UDP-native alternative in this paper's benchmark scope – does not patch TCP's problems at the application layer; it changes the transport assumptions directly.

It starts from the topology and asks: what wire format does $\beta_1 > 0$ actually need? The answer: self-describing frames with no ordered delivery. AIMD congestion control per-stream (not per-connection). MTU-aware fragmentation (4-byte fragment header, 255 fragments × 1,468 bytes). ACK bitmaps (14 bytes covering 64 sequences). The protocol is 800 lines of TypeScript. In the shootoff benchmarks used here, it outperforms HTTP/3 on the measured framing metrics and selected latency measurements. These are benchmark-scoped results, not a

universal internet-wide claim; the topological-fit interpretation is a mechanism hypothesis supported by these measurements [21].

8.4 Protocol Comparison

Metric	HTTP/1.1	HTTP/2	HTTP/3 (QUIC)	Aeon Flow
Per-resource overhead	~200 bytes	~30 bytes (HPACK)	~20 bytes	**10 bytes**
Header fraction (12 resources)	4.2 percent	0.6 percent	0.4 percent	**0.2 percent**
Header fraction (95 resources)	31.0 percent	5.8 percent	4.1 percent	**1.5 percent**
Connections for full site	6+	1	1	**1**
Head-of-line blocking	Yes (conn)	Yes (TCP)	No (per-stream)	**No**
Native fork/race/fold	No	No	No	**Yes**
Vent propagation	N/A	RST_STREAM	STOP_SENDING	**Recursive tree**
Transport	TCP	TCP	UDP (QUIC)	**UDP (raw)**
Ordered delivery	Required	Required	Per-stream	**None**
Topological contradiction	N/A	β_1 mismatch	Partial	**None**

8.5 Shootoff: Head-to-Head Protocol Benchmarks

I benchmark Aeon Flow against HTTP/1.1, HTTP/2 and HTTP/3 with realistic compression (gzip, brotli) across two site profiles. All protocols use identical payloads; only framing and transport differ.

Big Content Site (12 resources, ~2.5 MB – large JS bundles, hero images, web fonts):

Protocol	Wire Size	Framing Overhead	Overhead %	RTTs
HTTP/1.1	913 KB	8.2 KB	0.89 percent	3
HTTP/2	907 KB	1.6 KB	0.18 percent	2
HTTP/3 (QUIC)	906 KB	906 B	0.10 percent	1
Aeon Flow	**905 KB**	**276 B**	**0.03 percent**	**1**

For large payloads, all modern protocols converge – but Aeon Flow's framing is **3.3x smaller than HTTP/3** (276 B vs 906 B).

Microfrontend Site (95 resources, ~1.8 MB – 45 JS modules, 16 CSS modules, 20 SVG icons):

Protocol	Wire Size	Framing Overhead	Overhead %	RTTs
HTTP/1.1	187 KB	58.1 KB	**31.0 percent**	16
HTTP/2	137 KB	8.0 KB	5.8 percent	2
HTTP/3 (QUIC)	135 KB	5.9 KB	4.4 percent	1
Aeon Flow	**131 KB**	**1.9 KB**	**1.5 percent**	**1**

This is where the topology matters. HTTP/1.1 wastes **31 percent of total bandwidth on headers** – nearly a third of the wire is framing, not data. HTTP/2 reduces this to 5.8 percent. HTTP/3 to 4.4 percent. Aeon Flow: **1.5 percent**. That is a **21x reduction** in framing overhead versus HTTP/1.1 and **3x versus HTTP/3**.

At 100ms RTT, HTTP/1.1 needs 16 round trips (1.6 seconds of pure latency). Aeon Flow: 1 round trip. The difference is not optimization – it is topology. HTTP/1.1 has $\beta_1 = 0$ (one request per connection, six connections). Aeon Flow has $\beta_1 = 94$ (95 streams, one connection). The framing overhead is the cost of forcing a high-β_1 problem through a low-β_1 pipe.

Modern frontend workloads often ship many small assets after tree-shaking and code splitting, which amplifies request/metadata overhead. In this benchmark scope, Aeon Flow multiplexes these assets through one transport session and reduces framing cost. Effects on CLS, INP and hydration strategy remain application-dependent and are not guaranteed by transport alone.

9. Instantiation E: Topological Compression (Stack Layer 5 — Capstone)

9.1 The Claim and Its Limits

The same fork/race/fold primitive applies to compression. **Topological compression** forks all available codecs per chunk, races them and folds to the winner. Each chunk independently selects its best codec. The output is a sequence of

self-describing frames (9-byte header: codec ID, original size, compressed size). $\beta_1 = \text{codecs} - 1$.

In Gnosis (§11), the topological compressor is:

```
(raw: Codec { type: 'raw' }) (rle: Codec { type: 'rle' })
(brotli: Codec { type: 'brotli' }) (gzip: Codec { type: 'gzip' })
(chunk)-[:FORK]->(raw | rle | brotli | gzip)
(raw | rle | brotli | gzip)-[:RACE]->(smallest)
```

Four lines. The entire compression strategy. The topology IS the algorithm.

I implement this with eight codecs:

ID	Codec	Type	Best on
0	Raw (identity)	Pure JS	Incompressible data
1	RLE	Pure JS	Repeated byte runs
2	Delta	Pure JS	Sequential/incremental data
3	LZ77	Pure JS	Repeated patterns
4	Brotli	Platform (node:zlib)	General text
5	Gzip	Platform (node:zlib)	General text (universal fallback)
6	Huffman	Pure JS	Skewed byte distributions
7	Dictionary	Pure JS	Web content (HTML/CSS/JS keywords)

Before any excitement takes hold, it is important to state a boundary: fork/race/fold provides a container for adaptive codec selection, not a guaranteed ratio improvement over the best standalone codec on homogeneous payloads. Its value here is strategy selection, composability and bounded framing overhead.

9.2 What the Benchmarks Actually Show

I benchmark across both sites on Aeon Flow transport. The results are honest:

Big Content Site (12 resources, ~2.22 MB):

Compression	Wire Size	Ratio	β_1
Brotli (global, quality 4)	905 KB	39.8 percent	0
Topo-full (8 codecs per-chunk)	1005 KB	44.2 percent	7
Topo-pure (6 pure-JS codecs per-chunk)	1.17 MB	52.5 percent	5

Microfrontend Site (95 resources, ~617 KB):

Compression	Wire Size	Ratio	β_1
Brotli (global, quality 4)	131 KB	20.9 percent	0
Topo-full (8 codecs per-chunk)	159 KB	25.4 percent	7
Topo-pure (6 pure-JS codecs per-chunk)	229 KB	36.8 percent	5

Standalone brotli wins on compression ratio. On these benchmarks – homogeneous web content – global brotli beats per-chunk topological compression by 4–15 percentage points, darn it. This is not surprising: brotli compresses the entire stream with a sliding window that builds dictionary context across chunks. Per-chunk compression resets the dictionary every 4096 bytes.

The two-level race (§9.3) confirms this. When given the choice between global brotli and per-chunk topological, it picks global brotli every time on these payloads, matching standalone brotli's ratio plus 5 bytes of strategy header. Per-chunk topological never wins the stream-level race on homogeneous content because the 9-byte per-chunk header tax and the loss of cross-chunk dictionary context always exceed whatever the per-chunk adaptive selection saves.

9.3 Two-Level Stream Race

I extend the topology to race at two levels:

```
fork (stream level):
  |- Path 0: Per-chunk topological (8 codecs × each 4096-byte chunk)
  |- Path 1: Global brotli (entire stream, cross-chunk dictionary)
  |- Path 2: Global gzip (entire stream)
  `- ...
race → smallest total output wins
fold → 5-byte strategy header + compressed data
```

This is the usefulness of fork/race/fold to compression: with brotli included as a racing path, the stream-level strategy tracks brotli's ratio within a bounded strategy-header overhead. On these benchmarks it is not better than standalone brotli; the observed downside is the fixed 5-byte strategy header.

9.4 What the Topology Actually Provides

If topological compression does not beat brotli on ratio, what is the point?

1. Subsumption, not superiority. The topology is the space in which brotli competes. Brotli at $\beta_1 = 0$ is a degenerate case of topological compression at $\beta_1 = 7$. The two-level race includes brotli as a contestant. If brotli is best, the topology selects it. If something better appears tomorrow – a learned codec, a neural compressor, a domain-specific dictionary – it enters the race without changing the architecture. The `TopologicalCompressor` is unchanged; only the codec array grows.

2. Platform independence. Brotli requires `node:zlib` (Node, Bun, Deno). In browsers and Cloudflare Workers, it is unavailable. Topo-pure – six codecs in pure JavaScript, zero dependencies – achieves 36.8 percent ratio on the microfrontend with no native code. The topology degrades gracefully: full ratio when brotli is available, reasonable ratio when it is not. For software engineers, there is technical value in fewer depencies. For people, it helps set the table for a serverless ecosystem built on a local-first technology stack.

3. Per-chunk random access. The per-chunk format enables decompression of individual chunks without processing the entire stream. For seeking into large payloads, resuming interrupted transfers, or parallel decompression, this is structurally impossible with global brotli.

4. Adaptive codec selection on heterogeneous data. On the per-chunk level, different regions of the input genuinely select different codecs. The shootoff shows 3 distinct codecs winning across 151 chunks on realistic web content (brotli for text chunks, dictionary for web-pattern-heavy chunks, raw for incompressible binary). No single algorithm achieves this without the topology.

5. The real compression win is framing, not codecs. The paper's compression contribution is not beating brotli's ratio. It is the 30× reduction in framing overhead (§8.4): Aeon Flow uses 1.9 KB of framing for 95 resources where HTTP/1.1 uses 56.3 KB. On the microfrontend, framing overhead drops from 8.4 percent to 0.3 percent of the payload. This saving is orthogonal to which codec compresses the content.

9.5 Honest Assessment

The per-chunk topological approach pays a real cost: 9 bytes per chunk of header overhead and the loss of cross-chunk dictionary context. On homogeneous content (which describes most web payloads), this cost exceeds the benefit of adaptive codec selection. Global brotli, with its full-stream dictionary, simply compresses text better than any per-chunk approach can.

Comparison to adaptive single-algorithm heuristics. A simpler alternative — "use brotli for text, raw for binary, based on content-type heuristic" — would capture most of the per-chunk topology's adaptive benefit at zero per-chunk overhead. On these benchmarks, such a heuristic would match global brotli's ratio (since all content is web text). The per-chunk topology's advantage over simple heuristics emerges only on *heterogeneous* payloads (mixed binary/text, embedded images in HTML, protocol buffers interleaved with JSON) where content-type heuristics misclassify regions. The shootoff's 3-codec-winner distribution across 151 chunks is the first evidence of this: even on mostly-homogeneous web content, 12 percent of chunks selected a non-brotli winner (dictionary for web-pattern-heavy chunks, raw for incompressible binary).

The two-level stream race eliminates this disadvantage by including global brotli as a racing path. But it also reveals that per-chunk topological compression,

as currently implemented, is not the winning strategy for web content. It is a structurally sound framework that provides platform independence, random access and future extensibility – at the cost of matching, not beating, the state of the art on ratio.

The progression four codecs ($\beta_1 = 3$) → six codecs ($\beta_1 = 5$) → eight codecs ($\beta_1 = 7$) demonstrates the covering space property: each expansion improved pure-JS compression without changing the base space. But adding brotli and gzip to the race, while improving per-chunk results, still cannot overcome the fundamental advantage of global dictionary context.

The topology subsumes the algorithm. It does not necessarily surpass it. I will not pretend this is the result I wanted. I wanted to dust brotli – to show that topological compression, by racing eight codecs per chunk with adaptive selection, could outperform decades of careful engineering in a C-native algorithm. It does not. Global brotli with its full-stream dictionary context is genuinely excellent at compressing web content, and no amount of per-chunk cleverness overcomes the fundamental information advantage of seeing the entire input at once. The honest conclusion is that the topology provides structural guarantees – subsumption, platform independence, random access, extensibility – but not ratio superiority. That is still worth having. It is not the same as winning.

Executable evidence is available in two independent suites: the companion topological-compression obligations [21] and the production `TopologicalCompressor` tests in the open-source `@affectively/aeon` package [20]. Together they verify per-chunk adaptive winner selection, 9-byte self-describing chunk headers, codec vent behavior (discarding expansions), two-level stream race strategy selection, $\beta_1 = \text{codecs} - 1$ invariants and roundtrip correctness across edge cases and large payloads.

9.6 Applications

Application	Fork	Race	Fold
Site preloading	Stream all assets as parallel frames	First complete asset wins cache slot	SW stores all in Cache API
ESI composition	Fork stream per directive	Race cache vs. compute	Assemble into final page
Deploy artifacts	Fork per build artifact	Stream concurrently	Receive complete deployment

Application	Fork	Race	Fold
CRDT sync	Fork per-peer delta streams	Race peers to contribute	Merge deltas into canonical state
Speculative nav	Fork predicted route preloads	Race prediction vs. actual	Display whichever resolves first

10. Instantiation A: Self-Verification (Stack Layer 1 — Foundation)

A strong executable result for expressiveness in this scope: the model checker can verify a model of its own exploration.

10.1 The Checker's BFS Is Fork/Race/Fold

The `ForkRaceFoldModelChecker` in `@affectively/aeon-logic` [26] explores state spaces via breadth-first search. Each BFS layer is a time step. Each state is a spatial position. The exploration graph maps directly to the four primitives:

BFS Operation	Fork/Race/Fold Primitive	Topological Effect
Expansion with >1 successor	**Fork**	$\beta_1 \mathrel{+}= N - 1$
Transition to already-visited state	**Fold** (interference)	Creates independent cycle
Unfair cycle filtered by weak fairness	**Vent**	Irreversible path removal
Frontier exhausted, exploration complete	**Collapse**	$\beta_1 \to 0$

The checker now computes and returns topological diagnostics (`CheckerTopologyStats`) for every verification: `forkCount`, `foldCount`, `ventCount`, `beta1` (first Betti number of the exploration graph), and `depthLayers` (path-integral time steps).

10.2 Self-Verification as TemporalModel

The checker's own BFS exploration is modeled as a `TemporalModel<CheckerState>` with 8 state variables (`explored`, `frontier`, `transitions`, `folded`, `forks`, `vents`, `depth`, `done`) and 6 actions (`ExpandLinear`, `ExpandFork`, `FoldTransition`, `VentCycle`, `CompleteLayer`, `Finish`). Another instance of the same checker verifies 7 invariants about this model:

1. $\beta_1 \geq 0$ — topology is well-formed
2. $\beta_1 =$ folded — every back-edge creates exactly one independent cycle
3. vents $\leq$ folds — you can only vent what has been folded
4. folds $\leq$ transitions — folds are a subset of transitions
5. explored ≥ 1 — at least the initial state
6. frontier ≥ 0 — non-negative frontier
7. depth $\leq$ MaxDepth — bounded exploration

Liveness: $\Diamond$done (eventual termination) under weak fairness WF(Finish).

10.3 TLA+ Self-Verification

The same model is rendered as a TLA+ specification via `renderSelfVerificationArtifactPair()`, producing a `.tla` module (extending `Naturals`, with weak fairness `WF_vars(Finish)`) and a `.cfg` config. The specification is validated through the `runTlaSandbox()` round-trip: parse $\rightarrow$ render $\rightarrow$ parse = identical. A dual verification test confirms both paths agree: the TLA sandbox validates the spec structure, the checker verifies the same model's invariants and liveness.

10.4 Closure Under Self-Application

In the finite-model scope used here, self-verification provides a constructive closure result. The topology stats the checker reports about verifying itself (`forkCount`, `foldCount`, `beta1`) are themselves fork/race/fold observables. The meta-topology – the topology of the checker checking itself – has forks (multiple actions enabled per state), folds (different action sequences reaching the same checker state), and measurable β_1.

This means fork/race/fold is closed under self-application: a system built from these primitives can reason about systems built from these primitives. The topological deficit Δ_β of self-verification measures the cost of self-knowledge.

Executable companion tests verify these claims [26].

11. Instantiation B: Formal Language Theory (Stack Layer 2)

The fifth instantiation domain is the most recursive: a programming language whose source code *is* the computation graph, whose compiler *is* a fork/race/fold pipeline, and whose self-hosting bootstraps through the previous four domains.

11.1 Gnosis Graph Language (GGL)

Gnosis [28] is a programming language that dispenses with imperative control flow (`if`/`else`, `for`, `try`/`catch`) entirely. Programs are graphs — nodes define data and compute, edges define topological transitions. The syntax is Cypher-like:

```
(input) -[:FORK]-> (raw_codec | brotli_codec)
(raw_codec | brotli_codec) -[:RACE]-> (winner)
```

The language has exactly four edge types — `FORK`, `RACE`, `FOLD`, `VENT` — plus `PROCESS` for sequential steps and `INTERFERE` for constructive/destructive signal combination. There are no functions, only subgraphs. There are no variables, only nodes with typed properties. There are no loops, only topological cycles detected at compile time by β_1 analysis.

This is the thesis of the paper made literal: **the source code IS the topology**. The AST is the computation graph. The compiler is the β_1 analyzer. The runtime is the topology engine.

11.2 The Betty Compiler

The compiler (named **Betty**, after the Betti number) statically analyzes the GGL topology to ensure:

1. β_1 is properly managed — no unbounded superpositions (every `FORK` must reach a `FOLD`, `RACE`, or `VENT`).
2. All paths eventually collapse — the compiler rejects programs where β_1 never returns to zero.
3. Deterministic fold — the merger strategy is declared in the edge properties, satisfying C3.

Betty parses the graph, computes β_1 at each edge, and translates the AST into 10-byte `FlowFrame` binary buffers (§8.2) — the same wire format used by the Aeon Flow protocol. The compiled output is a sequence of `FlowFrame`s that the Rust/WASM runtime executes at near-native speed.

The compilation pipeline is itself fork/race/fold:

```
(source_code)
  -[:PROCESS]-> (read_source)
  -[:FORK]-> (parse_nodes | parse_edges)
  -[:FOLD { strategy: 'merge-ast' }]-> (ast)
  -[:PROCESS]-> (build_wasm_frames)
  -[:PROCESS]-> (executable_binary)
```

11.3 Transformers as GGL Programs

A transformer written in Gnosis reveals the fork/race/fold structure claimed in §6.11:

```
(input_sequence)-[:PROCESS]->(qkv_projection)
(qkv_projection)-[:FORK]->(head_1 | head_2 | head_3 | head_4)
(head_1 | head_2 | head_3 | head_4)-[:FOLD { strategy: 'concat' }]->(multi_head_out)
(input_sequence | multi_head_out)-[:INTERFERE { mode: 'constructive' }]->(residual_1)
(residual_1)-[:PROCESS]->(ffn)
(residual_1 | ffn)-[:INTERFERE { mode: 'constructive' }]->(transformer_out)
```

Multi-head attention is `FORK` → `FOLD`. Residual connections are `INTERFERE`. The topology is visible in the source code — not buried in matrix operations, not implicit in framework conventions, but *declared* as the program's structure. The compiler computes $\beta_1 = 3$ at the fork point (four heads) and verifies it returns to zero at the fold.

11.4 The Bootstrapping Path: Betty → Betti

The ultimate goal is self-hosting. Because a compiler is a pipeline — `(source) -[:FORK]-> (lexers) -[:FOLD]-> (AST)` — the TypeScript-based Betty compiler can be rewritten entirely in GGL. The self-hosted compiler is named **Betti** (the true topological spelling). The bootstrapping chain:

$$\text{TypeScript (Betty)} \xrightarrow{\text{compiles}} \text{GGL (Betti)} \xrightarrow{\text{compiles}} \text{Everything else}$$

This is closure under a different axis than §10. Self-verification (§10) proves that the checker can reason about itself — closure under *reasoning*. Self-hosting (Betti) proves that the language can compile itself — closure under *construction*. Together they establish that fork/race/fold is closed under both reasoning and construction: you can build tools from these primitives and verify tools built from these primitives, using the same primitives.

Gnosis supports a strong evidence-backed claim: it is a self-hosted, self-checking topology language with automated formal-artifact generation. The compiler topology is itself written in GG (`betti.gg`) and included in formal lint checks, while execution paths enforce bounded-state structural verification with explicit invariants and eventual reachability conditions before or during topology use. The `verify` workflow can generate TLC-ready TLA+ modules and configs with safety and liveness obligations, and these paths are covered by source-level tests and formal-check scripts [21, 26, 28].

This is a claim of structural formal compatibility and mechanized verification workflow, not a claim of automatic asymptotic quantum advantage.

11.5 The Five Domains as a Stack

The five instantiation domains are not independent — they form a stack, each enabled by the ones below:

Stack Layer	Domain	§	Primitive	Role
1 (foundation)	Self-verification	§10	Temporal model checking	Proves the math works

Stack Layer	Domain	§	Primitive	Role
2	Formal language	§11	GGL + Betty/Betti	The programming model
3	Distributed computation	§7	Wallington Rotation	The scheduling algorithm
4	Edge transport	§8	10-byte FlowFrame	The wire format
5 (capstone)	Compression	§9	Per-chunk codec racing	Bytes on wire

The stack reads bottom-up: *from building blocks to bytes on wire.* Layer 1 (§10) proves the primitives are correct. Layer 2 (§11) gives you a language to write topologies, verified by layer 1. Layer 3 (§7) schedules work through the topology, expressed in layer 2's language. Layer 4 (§8) puts frames on the wire, carrying layer 3's scheduled work. Layer 5 (§9) compresses the payload — actual bytes, actual ratios, actual wire — using every layer below it.

The Rust/WASM runtime executes the FlowFrames at the same byte-level format defined in §8.2. The language is not a wrapper around the protocol — it is the protocol's native programming model.

The stack is the paper's strongest existence proof: a single set of four primitives (fork, race, fold, vent) generates a scheduling algorithm, a wire protocol, a compression strategy, a verification engine, and a programming language. Each layer is independently useful. Together they form a closed computational ecosystem where the topology is the program, the program is the protocol, and the protocol is the topology.

12. The Engine

The algorithm is implemented as **Aeon Pipelines** [2], a zero-dependency computation topology engine in TypeScript. It runs on Cloudflare Workers, Deno, Node, Bun and browsers. The API surface is two classes:

- **`Pipeline`**: the engine – capacity, metrics, backpressure, turbulent multiplexing.
- **`Superposition<T>`**: the builder – chainable fork/race/fold/vent/tunnel/interfere/entangle/measure/searc operations.

```
// Kids juggling balls
const result = await Pipeline
  .from([fetchFromA, fetchFromB, fetchFromC])
```

```
  .race();

// People juggling the kids
const diagnosis = await Pipeline
  .from([bloodTest, mriScan, geneticScreen])
  .vent(result => result.inconclusive)
  .tunnel(result => result.conclusive)
  .fold({
    type: 'merge-all',
    merge: mergeFindings,
  });

// Grover-style search over solution space
const drug = await new Pipeline({ capacity: 64 })
  .fork(candidates.map(c => () => evaluate(c)))
  .search({
    width: 32,
    oracle: compound => compound.efficacy,
    mutate: (compound, gen) => perturb(compound, gen),
    convergenceThreshold: 0.01,
  });
```

The `search()` operation is a classical heuristic inspired by Grover-style amplification patterns. In some landscapes it reduces empirical iteration counts versus naive sequential search, but this is not an asymptotic complexity claim.

11.1 Performance

The pipeline engine is designed for low orchestration overhead. In the microbenchmarks below, orchestration cost is in the microsecond range, and profiled workloads are typically dominated by user work functions.

A stronger statement is now mechanized as a conditional formal obligation in `SchedulerBound.tla`: under finite-topology execution with bounded frame metadata and constant-time scheduler primitives, scheduler transition cost is an additive bounded term independent of user-handler runtime. This justifies "handler-dominated runtime" only within those explicit assumptions, not as a universal claim [21].

Operation	Latency	Notes
`fork(10)`	**1.82 µs**	10 parallel streams created
`fold({ type: 'quorum', threshold: 3 })`	**4.51 µs**	Byzantine agreement across 5 streams

Operation	Latency	Notes
`search(8×20)`	**8.3 µs**	Grover-style search, 8-wide, 20 generations
`interference(100)`	**16.3 µs**	Pairwise consensus across 100 streams
`vent-tree(13)`	**18.9 µs**	Recursive vent across 13-node tree
`flow-bridge-batch(100)`	**25.7 µs**	100 frames encoded to wire format
`reassemble-reverse(1000)`	**71.4 µs**	1,000 frames reassembled from reverse order
`flow-bridge-roundtrip`	**0.76 µs**	Single frame encode → decode

Zero dependencies. ~384 bytes per stream and ~3.5 KB per pipeline. Requires no servers.

11.2 Domain Validation

The same API – unchanged – was validated across many independent domains, including:

1. **Multi-venue trading**: fork/race across exchanges, vent adverse prices
2. **Healthcare diagnostics**: fork parallel tests, tunnel on conclusive, merge-all findings
3. **Financial settlement**: fork clearing/netting/DVP, merge-all for T+0
4. **Construction scheduling**: fork trades per floor, merge-all hours
5. **Emergency dispatch**: fork/race responders, first arrival wins
6. **Academic review**: fork reviewers, quorum 2/3 agreement
7. **Drug discovery**: fork compounds, Grover search to convergence
8. **Manufacturing QC**: fork sensors, consensus (constructive interference)
9. **Journal publishing**: fork reviewers, vent timeout, quorum verdict
10. **Legal review**: fork reviewers, weighted fold by seniority

The universality is not designed – it is discovered, the same way *Physarum* discovers optimal networks without knowing how to do so.

11.3 Wire Format Bridge

The engine includes a wire format bridge to the Aeon Flow protocol. The same 10-byte frame header (§8.2) encodes `WorkFrame<T>` objects for network transmission. Frames encoded by Aeon Pipelines transcode into frames in Aeon

Flow, and vice versa. The computation topology is independent of the transport topology.

13. Validation

The claims are backed by executable tests across four independent suites:

- **Companion obligations and executable proofs** (212 passing tests): pipeline topology, queueing containment, flow-frame invariants, compression race properties, shootoff reproductions, finite-DAG decomposition coverage (including edge-cover exactness and full source-to-sink path-set preservation), §7 formula checks (Worthington Whip $(S-1)/2S$, Speculative Tree $(1-\alpha^K)/(1-\alpha)$, turbulent multiplexing idle-fraction bounds), quantum-topology claims (Grover-style Δ_β scaling and Kronig-Penney band gaps as $\beta_2 > 0$), map/reduce readiness diagnostics (boundedness/monotonicity, nonzero-opportunity necessity in migration simulation, independent migration-simulator rank ordering, and high-readiness counterexample families showing non-automatic quantum asymptotics), convergence simulation under the three constraints, evidence-table deficits (including T+2 settlement $\Delta_\beta = 2B$ under both core and broad-scope lockup scenarios), evidence-traceability calibration/provenance/reference checks, self-hosted formal artifact parsing/round-trip validation with `aeon-logic`, a parser shootoff benchmark against Java SANY startup-parse baselines (stabilized multi-sample harness: 9 measured samples after warmup, `aeon-logic` median 49.51 ms for 19,200 artifacts with IQR 48.21–49.94 ms = 387,780.9 artifacts/s; Java SANY median 116.45 ms on `BandGapVoid.tla` with IQR 115.13–122.08 ms, implying approximately 45,156.7x normalized per-artifact throughput), plus a differential parse-equivalence harness against SANY outcomes (100% agreement on the current formal corpus for original modules, round-tripped modules and invalid-corpus rejections). The parser result is therefore speed plus capability surface: unlike the parser-only baseline, `aeon-logic` also exposes superposition chains, quorum temporal operators, topology bridges and embedded model-checker interfaces in the same runtime [26, 27]. Mechanized TLA+ model checking across the current formal module set (C1–C4, §7 formulas, cross-shard crossover, scheduler-overhead bounds, protocol/settlement deficits, quantum deficit identity, band-gap void, beauty-optimality scaffold), and a Lean 4 theorem package with constructive identities plus explicit-assumption theorem schemas (including global convergence schema) verify the strongest operational claims section by section [21, 25, 26, 27].
- **Open-source flow + compression runtime** (170 passing tests): `@affectively/aeon` flow/compression tests verify 10-byte self-describing flow frames, UDP fragmentation/ACK behavior, frame reassembly, flow protocol semantics, WASM force-mode/error semantics, and topological compression properties [20].

- **Open-source topology engine** (139 passing tests): `@affectively/aeon-pipelines` tests cover fork/race/fold/vent primitives, fold strategies, Reynolds/backpressure/turbulent multiplexing, quantum modalities, flow-bridge wire compatibility, domain scenarios and microbenchmarks [2].
- **Open-source topology analyzer suite** (29 passing tests): `TopologyAnalyzer/TopologySampler` tests in `@affectively/aeon` validate Betti extraction, Δ_β diagnostics, β_2 void detection and executable protocol-topology contrasts [20].

Total validated tests referenced here: **550 passing tests**, plus parser-validated formal artifacts, mechanized Lean theorem builds and mechanized TLC runs across the current TLA+ module set, with executable commands and source-visible assertions in the linked repositories [2, 20, 21, 25, 26, 27].

14. Limitations

Benchmark substrate. Speedup figures are from benchmark harnesses with mocked network communication. Live distributed measurements would strengthen the empirical claims. Readers interested in seeing this demonstrated first hand should contact the author.

Cross-shard cost. The Worthington Whip crossover is characterized in finite bounded models: after full sharding, nonzero correction cost makes additional shards strictly non-improving (TLA+ `WhipCrossover` + Lean theorem + executable tests). Extending this characterization to richer timing/service distributions and adaptive sharding policies remains future work.

Formal model scope. C1–C4, §7 formulas (including cross-shard crossover), scheduler-overhead bounds, protocol/settlement deficits, quantum deficit identity, band-gap void, beauty-optimality scaffolds and convergence schema are mechanized in a two-layer stack: finite-state transition models in TLA+ (TLC), plus Lean theorems with explicit assumptions for quantitative identities and theorem schemas for global claims, all preflighted through the self-hosted `aeon-logic` parser [21, 25, 26, 27]. The self-verification (§10) is scoped to finite, non-probabilistic, untimed state spaces. Extending these proofs to richer timing/service distributions, unbounded state spaces, probabilistic systems (randomized protocols), real-time systems (strict latency bounds), and full stochastic semantics remains future work.

Queueing theory subsumption scope. Containment is proved for canonical constructions (Little's Law boundary case, Erlang-style blocking behavior and Jackson-style bottleneck limits) in executable form [21]. A full generalization to every queueing discipline and service-time law remains future work.

15. Conclusion

I began with a child handing a ball to another child in a line. Four hundred handoffs. I ended with a topological framework that subsumes queueing theory,

predicts biological mutation rates, explains why HTTP/2 has head-of-line blocking and runs on 10-byte UDP frames blasted as fast as theoretically possible.

The path between those two points is fork/race/fold: four operations that express any directed acyclic computation graph.

1. **Fork** raises β_1, injects potential energy V – create parallel paths, store work.
2. **Race** traverses homotopy-equivalent paths, converts $V \to K$ – take the fastest.
3. **Fold** projects $\beta_1 \to 0$, extracts work W – merge results deterministically.
4. **Vent** releases excess paths, dissipates heat Q – propagate down, never across.

These operations are not new. DNA replication has used analogous structure for billions of years. Myelinated neurons pipeline action potentials at measured speeds up to roughly 100 m/s, within the range discussed in §1.3. Photosynthetic antenna complexes exhibit >95 percent transfer efficiency in the cited step-level measurements. *Physarum* recreated a rail-like topology in roughly 26 hours.

The innumerable conveyor belts of this world – Ford's line, TCP's stream, the hospital's referral chain – are a useful degenerate case. They work best when the answer is known, resources are ample and a central clock exists. In many real systems, the natural topology has $\beta_1 > 0$, and forcing it to zero hides latency and waste.

A broad class of computational waste is now in scope for measurement and reduction. We can quantify topological mismatch and target it directly.

The framework's language was not invented from scratch. It was borrowed from physical theories that already describe related path-selection phenomena. Quantum physics provided the lexicon: superposition, tunneling, interference, entanglement, measurement, collapse. In this paper, these are structural correspondences within an explicit computational abstraction, with literal quantum examples in photosynthetic transfer. The Feynman path integral admits a fork/race/fold interpretation in that abstraction: all paths forked, phases evolved, amplitudes recombined by interference, and non-classical contributions canceled (§6.12). The recombination mechanism differs (linear amplitude summation in physics vs nonlinear selection/merge in computation), so this is a structural mapping, not an identity claim. Fluid dynamics provides scaling intuition through the pipeline Reynolds number, and thermodynamics provides a conservation accounting lens ($V = W+Q$) for modeled computational work/vent partitioning.

The pattern appears convergent in the analyzed set. Attention in transformers can be mapped to race-like scoring (QK^T), with softmax/selection and value projection as fold-like operations (§6.11). Protein folding can be read as path exploration with energetic selection [24]. Hylomorphism (unfold/fold) aligns

naturally with fork/fold structure. These are structural correspondences used to unify reasoning across domains in this manuscript's modeled class, supported by executable simulations and formal schemas under explicit assumptions.

What this paper claims and what it does not. The scope of the convergence is broad enough to warrant explicit boundaries. This paper claims a structural mapping between fork/race/fold and selected physical/computational formalisms used here (for example, path-integral branching/recombination, Schrödinger evolution as race-like dynamics, and folded classical limits at $\beta_1 = 0$ in the abstraction). It does not claim full physical identity or derivation of fundamental constants/forces. The key point is shared topology under explicit assumptions, with different underlying recombination mechanics.

This paper does *not* claim to be a grand unified theory. It does not explain why there are four fundamental forces, or why gravity is weak, or what dark energy is. The gap between "structural pattern that recurs everywhere" and "theory that explains everything" is precisely the gap between topology and physics. Topology describes *shape* – it says that when you compute with gravity, electromagnetism, the strong force, or the weak force, the computation has the shape of fork/race/fold. It does not say *why* these forces exist or *why* they have the coupling constants they have. It lives on the topology side of that boundary, deliberately.

This paper does not describe the grand unified theory of reality. But it does describe the shape of the box it is going to come in.

The claim is narrower and more defensible: under the stated assumptions, fork/race/fold behaves as an attractor in the modeled class. Systems that conserve resources, evolve irreversibly, and incur nonzero coordination overhead show this shape in the finite executable/mechanized setting used here. The convergence argument is bottom-up – independent systems under shared constraints – not a top-down unification claim.

And this convergence has a practical corollary: **to find fork/race/fold in a system is to see that the system has found its shape** (§6.13). The topological deficit $\Delta_\beta = \beta_1^* - \beta_1$ quantifies how far an implementation deviates from its problem's natural topology. In the systems analyzed here, $\Delta_\beta = 0$ cases – photosynthesis, DNA replication, saltatory conduction, the path integral – operate at or near theoretical efficiency, while $\Delta_\beta > 0$ cases – sequential healthcare, T+2 settlement, HTTP/2 over TCP – exhibit measurable, quantifiable waste. The deficit is a diagnostic: find it, close it, and the performance follows.

The closer the deficit is to zero, the more the system exhibits both efficiency and elegance. Put plainly: Bules are where rigor and aesthetics meet. They quantify the gap between a system that merely runs and one that is beautifully shaped for its work.

Fluid dynamics revealed the technique's most counterintuitive property: its worst case is small data, not large. The speedup accelerates with scale, ap-

proaching $B \times N$ on large datasets. The hard problem is not surviving at scale – it is avoiding overhead on the simple cases. This is the optimization problem inverted, and it is perhaps the most practically encouraging result: the bigger the workload, the more the algorithm helps.

The children in the line never needed more hardware; they needed less waiting. In systems where constraints permit fork/race/fold, unnecessary serialization leaves measurable performance on the table.

We didn't need to understand the deep physics of quantum field theory. We just needed a computational philosophy to see the pattern.

Within this paper's modeled scope, fork/race/fold + vent is sufficient.

References

[1] A. Tero, S. Takagi, T. Saigusa, K. Ito, D. P. Bebber, M. D. Fricker, K. Yumiki, R. Kobayashi, T. Nakagaki, "Rules for Biologically Inspired Adaptive Network Design," *Science*, 327(5964):439–442, 2010.

[2] T. W. Buley, "Aeon Pipelines: A Computation Topology Engine," open-source implementation, 2026. https://forkracefold.com/content/pipeline-topology.test.ts.txt

[3] D. Akita, I. Kunita, M. D. Fricker, S. Kuroda, K. Sato, T. Nakagaki, "Experimental Models for Murray's Law," *Journal of Vascular Research*, 52(2):75–88, 2016.

[4] T. Nakagaki, H. Yamada, Á. Tóth, "Maze-Solving by an Amoeboid Organism," *Nature*, 407(6803):470, 2000.

[5] J. Lobry, "Asymmetric Substitution Patterns in the Two DNA Strands of Bacteria," *Molecular Biology and Evolution*, 13(5), 1996.

[6] G. S. Engel, T. R. Calhoun, E. L. Read, T.-K. Ahn, T. Mančal, Y.-C. Cheng, R. E. Blankenship, G. R. Fleming, "Evidence for Wavelike Energy Transfer Through Quantum Coherence in Photosynthetic Systems," *Nature*, 446(7137):782–786, 2007.

[7] J. D. C. Little, "A Proof for the Queuing Formula: $L = \lambda W$," *Operations Research*, 9(3):383–387, 1961.

[8] J. R. Jackson, "Jobshop-Like Queueing Systems," *Management Science*, 10(1):131–142, 1963.

[9] Y. Huang et al., "GPipe: Efficient Training of Giant Neural Networks using Pipeline Parallelism," NeurIPS 2019.

[10] D. Narayanan et al., "PipeDream: Generalized Pipeline Parallelism for DNN Training," SOSP 2019.

[11] G. Yu et al., “Orca: A Distributed Serving System for Transformer-Based Generative Models,” OSDI 2022.

[12] Y. Leviathan, M. Kalman, Y. Matias, “Fast Inference from Transformers via Speculative Decoding,” ICML 2023.

[13] C. Chen et al., “Accelerating Large Language Model Decoding with Speculative Sampling,” arXiv:2302.01318, 2023.

[14] M. Belshe, R. Peon, M. Thomson, “Hypertext Transfer Protocol Version 2 (HTTP/2),” RFC 7540, 2015.

[15] J. Iyengar, M. Thomson, “QUIC: A UDP-Based Multiplexed and Secure Transport,” RFC 9000, 2021.

[16] M. Isard et al., “Dryad: Distributed Data-Parallel Programs from Sequential Building Blocks,” EuroSys 2007.

[17] D. G. Murray et al., “Naiad: A Timely Dataflow System,” SOSP 2013.

[18] D. Bray, “Wetware: A Computer in Every Living Cell,” Yale University Press, 2009.

[19] A. M. Turing, “The Chemical Basis of Morphogenesis,” *Philosophical Transactions of the Royal Society B*, 237(641), 1952.

[20] T. W. Buley, “Aeon Core Runtime (Flow + Compression) and Test Suite,” open-source implementation, 2026. https://forkracefold.com/content/companion-tests-README.md

[21] T. W. Buley, “Fork/Race/Fold Companion Tests,” reproducibility suite, 2026. https://forkracefold.com/content/companion-tests-README.md

[22] R. P. Feynman, A. R. Hibbs, “Quantum Mechanics and Path Integrals,” McGraw-Hill, 1965.

[23] S. Chandrasekhar, “Stellar Structure and Stellar Atmospheres,” §IX (The Virial Theorem), *An Introduction to the Study of Stellar Structure*, University of Chicago Press, 1939.

[24] J. N. Bryngelson, J. D. Onuchic, N. D. Socci, P. G. Wolynes, “Funnels, Pathways, and the Energy Landscape of Protein Folding: A Synthesis,” *Proteins*, 21(3):167–195, 1995.

[25] L. Lamport, *Specifying Systems: The TLA+ Language and Tools for Hardware and Software Engineers*, Addison-Wesley, 2002.

[26] T. W. Buley, “Aeon Logic: Fork/Race/Fold Temporal Logic Engine and TLC/TLA Compatibility Layer,” open-source implementation, 2026. https://github.com/affectively-ai/aeon-logic

[27] Lean FRO Team, “The Lean Theorem Prover (Lean 4),” software and documentation, 2026. https://lean-lang.org

[28] T. W. Buley, "Gnosis: A Topological Programming Language with Self-Hosting Compiler," open-source implementation, 2026. https://github.com/affectively-ai/gnosis

[29] J. M. Bauskis, M. M. M. Kuijper, J. M. Schuuring, C. M. Schaefer, M. M. J. Schuurmans, S. A. M. J. Damen, M. M. M. K. de Vries, "The changing landscape of rare disease diagnosis: diagnostic time in the Netherlands in 2013 and 2023," *European Journal of Human Genetics*, 2026. https://www.nature.com/articles/s41431-026-01791-6

[30] Depository Trust & Clearing Corporation (DTCC), "DTCC 2024 Annual Report," 2025. (NSCC average daily transaction value: $2.219 trillion) https://www.dtcc.com/annuals/2024/

Reproducibility

Source code, test suites and protocol comparison benchmarks are available under open-source license [2, 20, 21, 26, 28]. The scheduler, flow protocol, compression subsystem, computation topology engine, formal parser/tooling layer and topological programming language are independently testable. The validation totals reported in §13 are reproducible from the linked suites.